



DEEP
LEARNING
INSTITUTE



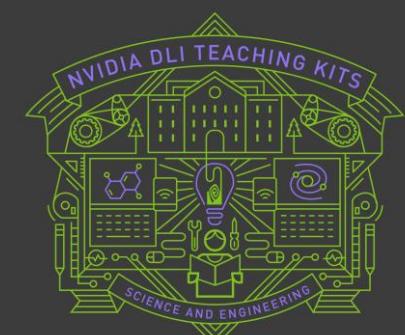
Deep Learning for Science and Engineering Teaching Kit

Deep Learning for Scientists and Engineers

Lecture 9: Physics-Informed Neural Networks (PINNs)

Instructors: George Em Karniadakis, Khemraj Shukla, Lu Lu

Teaching Assistants: Vivek Oommen and Aniruddha Bora





The Deep Learning for Science and Engineering Teaching Kit is licensed by NVIDIA and Brown University under the
[Creative Commons Attribution-NonCommercial 4.0 International License](#).

Course Roadmap

Module-1 (Basics)

- Lecture 1: Introduction
- Lecture 2: A primer on Python, NumPy, SciPy and *jupyter* notebooks
- Lecture 3: Deep Learning Networks
- Lecture 4: A primer on TensorFlow and PyTorch
- Lecture 5: Training and Optimization
- Lecture 6: Neural Network Architectures

Module-2 (PDEs and Operators)

- Lecture 7: Physics-Informed Neural Networks (PINNs)
- Lecture 8: PINN Extensions
- Lecture 9: Neural Operators

Module-3 (Codes & Scalability)

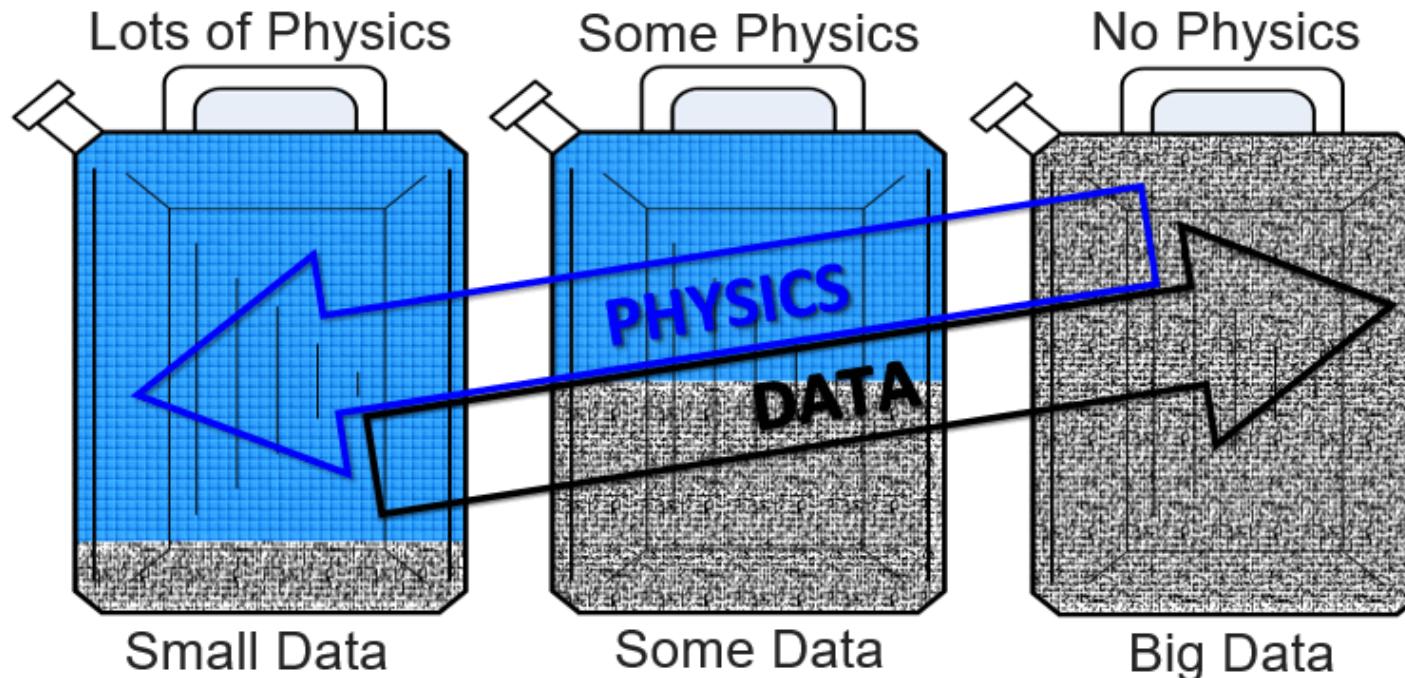
- Lecture 10: Multi-GPU Scientific Machine Learning

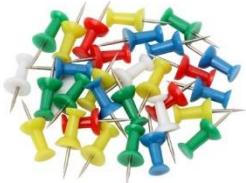
Contents

- ❑ Data + Physical Laws
- ❑ Data + Physical Laws + Neural Networks
- ❑ What is a PINN and Why PINNs
- ❑ PINN for Burgers Equation
- ❑ PINN for Boundary Value Problems
- ❑ Soft Constraints and Weights
- ❑ Hard Constraints: Boundary Conditions
- ❑ Linearly Constrained Neural Networks
- ❑ Hard Constraints: Design and Optimization
- ❑ Weighted Residual Methods
- ❑ hp-VPINNs: Domain Decomposition
- ❑ Variational Neural Networks
- ❑ Convergence Theory of PINNs
- ❑ Convergence Theory of hp-VPINNs
- ❑ Error Decomposition
- ❑ Error estimates of PINNs based on Quadrature
- ❑ PINNs vs DRM (Deep Ritz Method)
- ❑ Summary
- ❑ References

Data + Physical Laws

Three scenarios of
Physics-Informed Learning Machines





Data + Neural Networks + Physical Laws

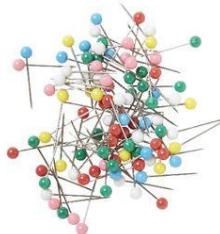
PINNs



Physics-informed neural networks: A deep learning framework
for solving forward and inverse problems involving nonlinear
partial differential equations

M Raissi, P Perdikaris, GE Karniadakis
Journal of Computational Physics 378, 686-707, **2019**

arXiv:1711.10561; arXiv:1711.10566 (2017)



PINNs: Physics-Informed Neural Networks

AI FOR SCIENCE

“... All of the **knowledge** are led to **physics equations** and now **embedded into the neural networks**. And it gives a gigantic head start, **physics informed neural networks...**”

Design Space Exploration
ICF + MERLIN – Fusion

Inverse Problems
LIGO – Gravitational Waves

Faster Prediction
ANI + MD - Chemistry

Real-time Steering
ITER – Fusion Energy

EXPERIMENTATION

NEURAL ESTIMATION

SIMULATION

DATA

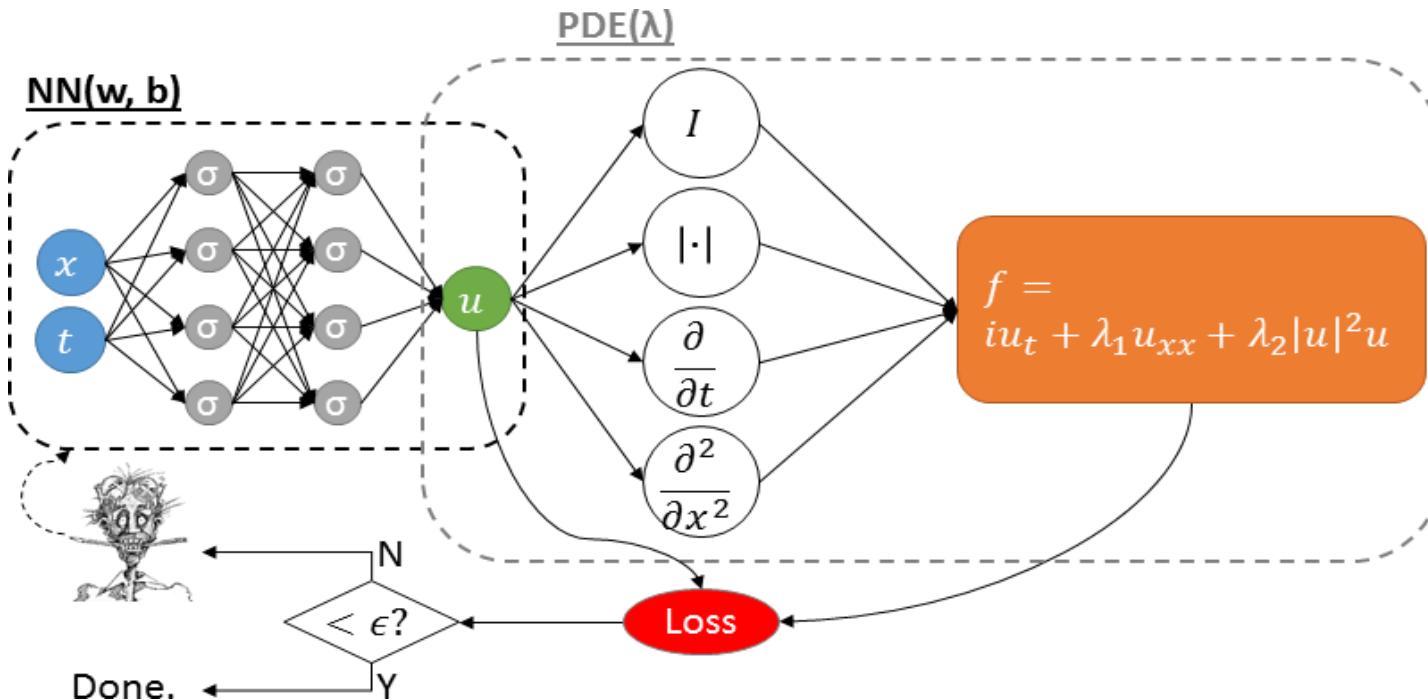
Fast Approximation

Real-time Steering

NVIDIA CEO J. Huang at SC19,
the annual supercomputing conference, Nov 19

What is a PINN? Physics-Informed Neural Network

We employ two (or more) NNs that share the same parameters



➤ Minimize:

$$MSE_u = \frac{1}{N_u} \sum_{i=1}^{N_u} |u(t_u^i, x_u^i) - u^i|^2,$$

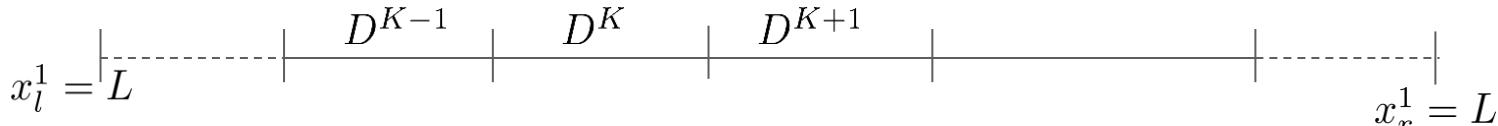
$$MSE = MSE_u + MSE_f,$$

$$MSE_f = \frac{1}{N_f} \sum_{i=1}^{N_f} |f(t_f^i, x_f^i)|^2.$$

Why PINNs: PINN vs (DG) FEM

- Simple Linear Wave Equation with $a = 30$

$$\frac{\partial u}{\partial t} + a \frac{\partial u}{\partial x} = 0, \quad x \in [-1, 1]$$



with initial condition $u(x, 0) = -\sin(\pi x)$ and periodic boundary conditions.

- The exact solution reads

$$u(x, t) = -\sin(\pi x - at)$$

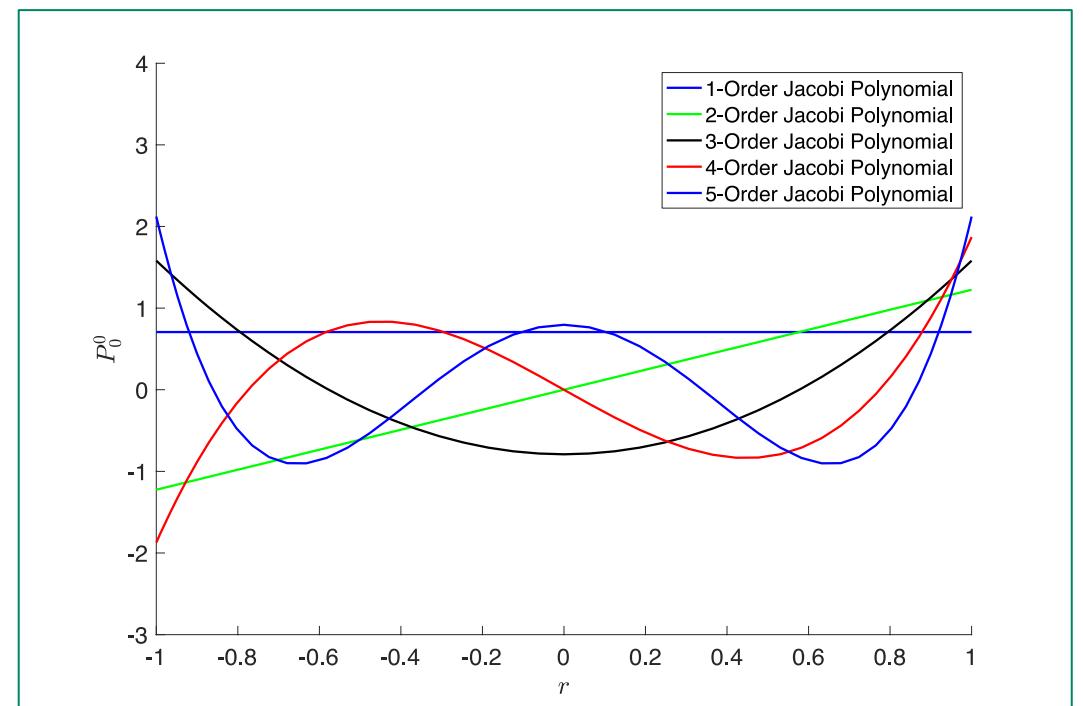
- The local solution is well approximated as

$$u_h^k(x, t) = \sum_{i=1}^{N_p} u_h^k(x_i^k, t) l_i^k(x)$$

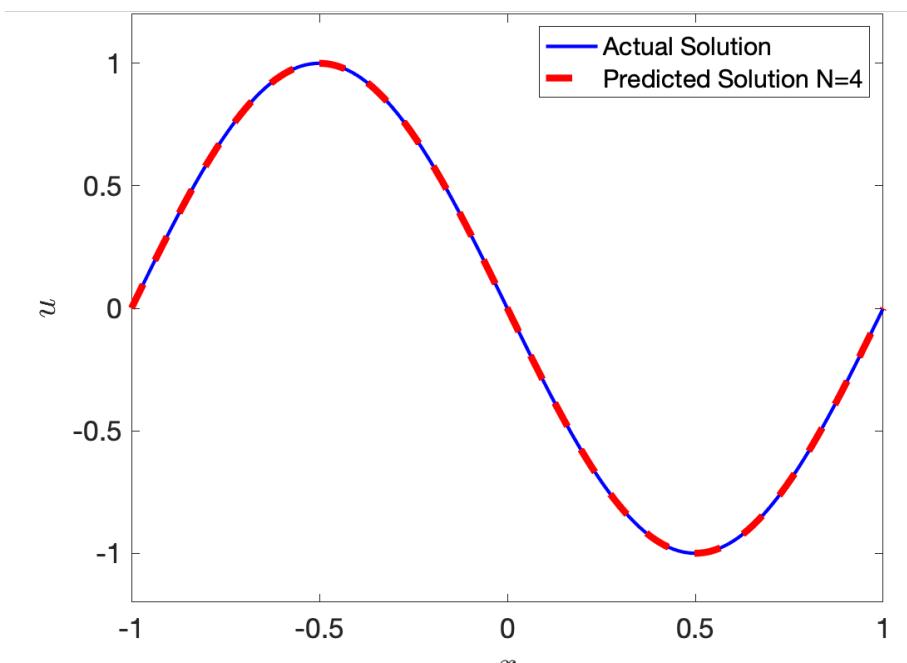
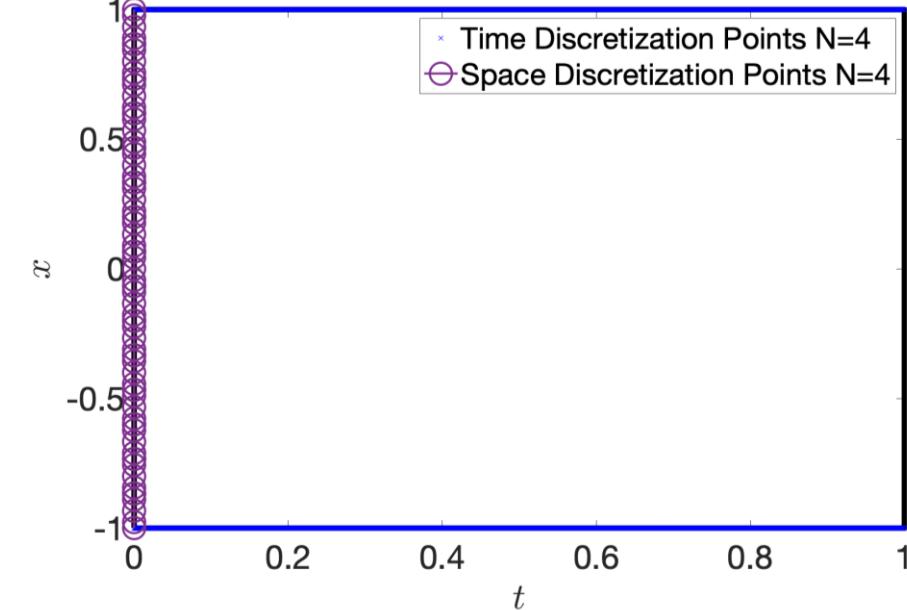
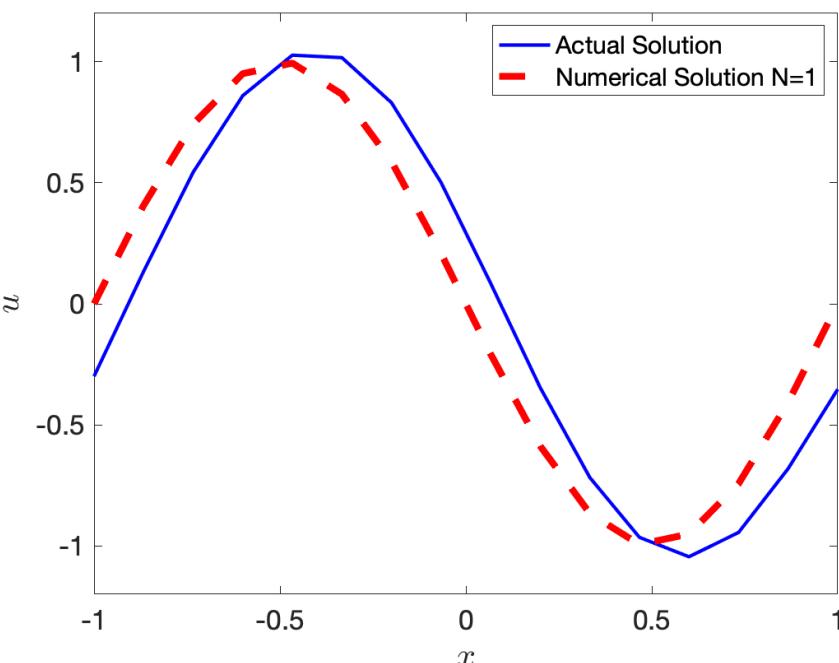
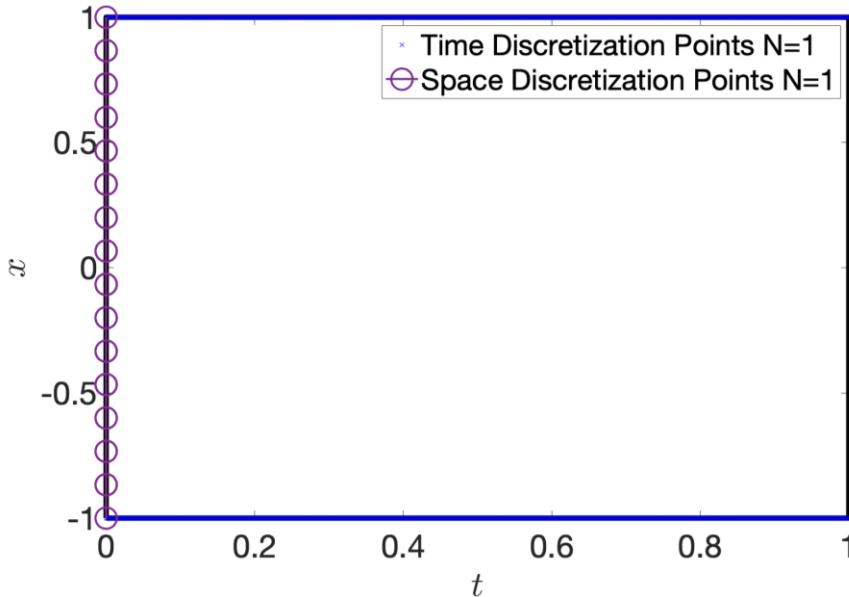
The local semi-discrete schemes look like

$$\mathcal{M}^k \frac{d\mathbf{u}_h^k}{dt} + 2\pi \mathcal{S} \mathbf{u}_h^k = [l^k(x)(2\pi u_h^k - (2\pi u_h)^*)],$$

where $(2\pi u_h)^*$ is numerical flux.

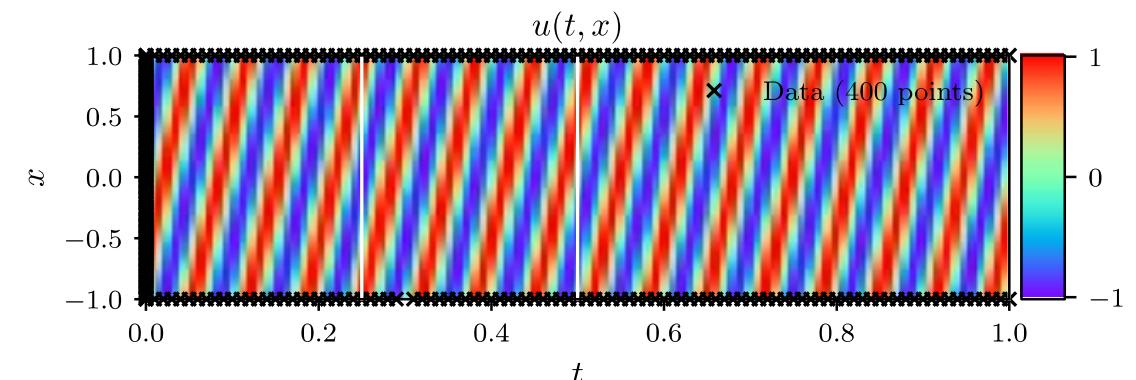
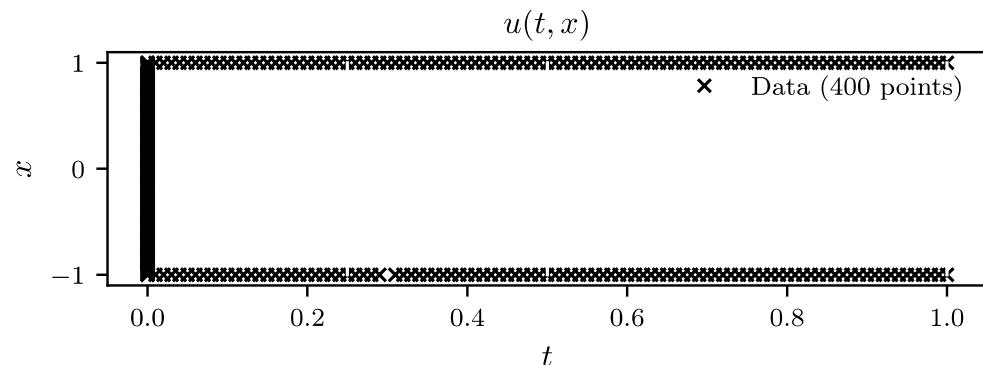


Why PINNs: PINN vs **(DG) FEM**

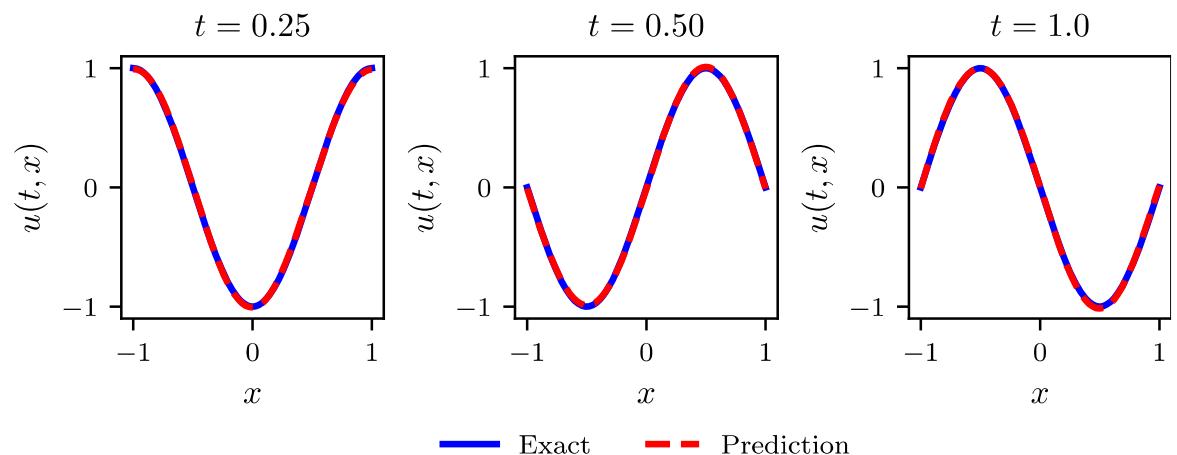
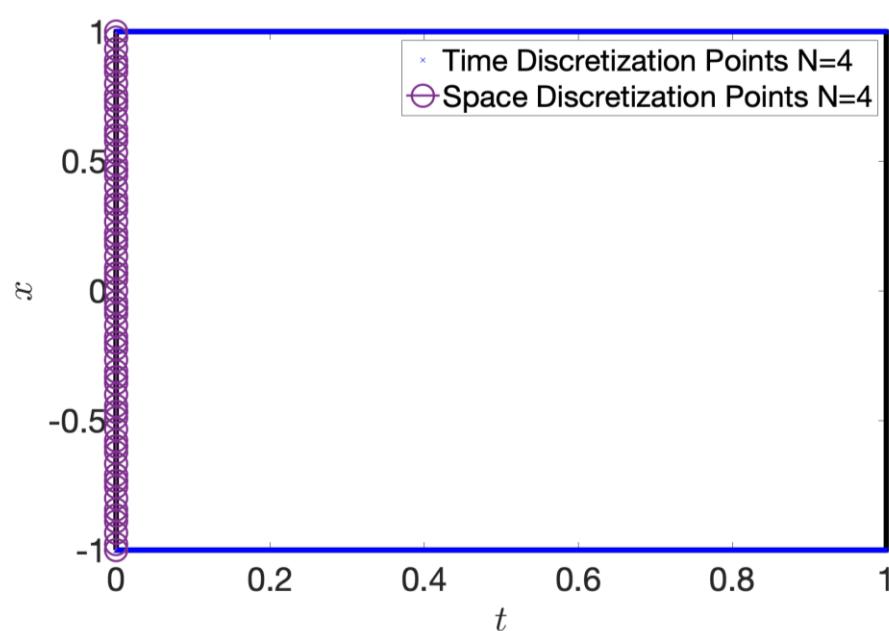


Why PINNs: PINN vs (DG) FEM

Data Distribution PINN



Data Distribution PINN



PINN for Burger's Equation: Implementation

Burger's Equation is defined as

$$u_t + uu_z - (0.01/\pi)u_{xx} = 0$$

1. Import Modules and Deep Learning Framework

```
import sys
sys.path.insert(0, 'Utilities/')
import os

from scipy.interpolate import griddata
from pyDOE import lhs
from plotting import newfig, savefig
from mpl_toolkits.mplot3d import Axes3D
import matplotlib.gridspec as gridspec
from mpl_toolkits.axes_grid1 import make_axes_locatable
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import time
import scipy.io
```

2a. Initialize Layers and NN Parameters

```
layers = [2, 20, 20, 20, 20, 20, 20, 20, 20, 20, 1]
L = len(layers)
W = [hyper_initial([layers[l-1], layers[l]]) for l in range(1, L)]
b = [tf.Variable(tf.zeros([1, layers[l]])) for l in range(1, L)]
```

2b. Glorot Normal Initialization

```
def hyper_initial(size):
    in_dim = size[0]
    out_dim = size[1]
    std = np.sqrt(2.0/(in_dim + out_dim))
    return tf.Variable(tf.random.truncated_normal(shape=size, stddev = std))
```

PINN for Burger's Equation: Implementation

3a. Forward Pass

```
def net_u(x, t, w, b):
    u = DNN(tf.concat([x,t],1), w, b)
    return u
```

3b. DNN Function:

```
# Neural Network
def DNN(X, W, b):
    A = X
    L = len(W)
    for i in range(L-1):
        A = tf.tanh(tf.add(tf.matmul(A, W[i]), b[i]))
    Y = tf.add(tf.matmul(A, W[-1]), b[-1])
    return Y
```

4. Residual Computation

```
def net_f(x,t,W, b, nu):
    with tf.GradientTape(persistent=True) as tape1:
        tape1.watch([x, t])
    with tf.GradientTape(persistent=True) as tape2:
        tape2.watch([x, t])
        u=net_u(x,t, W, b)
        u_t = tape2.gradient(u, t)
        u_x = tape2.gradient(u, x)
    u_xx = tape1.gradient(u_x, x)
    del tape1
    f = u_t + u*u_x - nu*u_xx
    return f
```

PINN for Burger's Equation: Implementation

5. Backward Propagation

```
def train_step(W, b, X_u_train_tf, u_train_tf, X_f_train_tf, opt, nu):
    x_u = X_u_train_tf[:,0:1]
    t_u = X_u_train_tf[:,1:2]
    x_f = X_f_train_tf[:,0:1]
    t_f = X_f_train_tf[:,1:2]
    with tf.GradientTape() as tape:
        tape.watch([W,b])
        u_nn = net_u(x_u, t_u, W, b)
        f_nn = net_f(x_f,t_f, W, b, nu)
        loss = tf.reduce_mean(tf.square(u_nn - u_train_tf)) + tf.reduce_mean(tf.square(f_nn))
    grads = tape.gradient(loss, train_vars(W,b))
    opt.apply_gradients(zip(grads, train_vars(W,b)))
    return loss
```

6. Predict

```
def predict(X_star_tf, w, b):
    x_star = X_star_tf[:,0:1]
    t_star = X_star_tf[:,1:2]
    u_pred = net_u(x_star, t_star, w, b)
    return u_pred
```

PINN for Burger's Equation: Implementation

7. Data Preparation and Training

```
nu = 0.01/np.pi
noise = 0.0
N_u = 100
N_f = 10000
Nmax=10000

layers = [2, 20, 20, 20, 20, 20, 20, 20, 20, 20, 1]
L = len(layers)
W = [hyper_initial([layers[l-1], layers[l]]) for l in range(1, L)]
b = [tf.Variable(tf.zeros([1, layers[l]])) for l in range(1, L)]

data = scipy.io.loadmat('./Data/burgers_shock.mat')

t = data['t'].flatten()[:,None]
x = data['x'].flatten()[:,None]
Exact = np.real(data['usol']).T
X, T = np.meshgrid(x,t)
X_star = np.hstack((X.flatten()[:,None], T.flatten()[:,None]))
u_star = Exact.flatten()[:,None]
# Domain bounds
lb = X_star.min(0)
ub = X_star.max(0)
xx1 = np.hstack((X[0:1,:].T, T[0:1,:].T))
uu1 = Exact[0:1,:].T
xx2 = np.hstack((X[:,0:1], T[:,0:1]))
uu2 = Exact[:,0:1]
xx3 = np.hstack((X[:, -1:], T[:, -1:]))
uu3 = Exact[:, -1:]

X_u_train = np.vstack([xx1, xx2, xx3])
X_f_train = lb + (ub-lb)*lhs(2, N_f)
X_f_train = np.vstack((X_f_train, X_u_train))
u_train = np.vstack([uu1, uu2, uu3])

idx = np.random.choice(X_u_train.shape[0], N_u, replace=False)

X_u_train = X_u_train[idx, :]
u_train = u_train[idx,:]

X_u_train_tf = tf.convert_to_tensor(X_u_train, dtype=tf.float32)
u_train_tf = tf.convert_to_tensor(u_train, dtype=tf.float32)
X_f_train_tf = tf.convert_to_tensor(X_f_train, dtype=tf.float32)
```

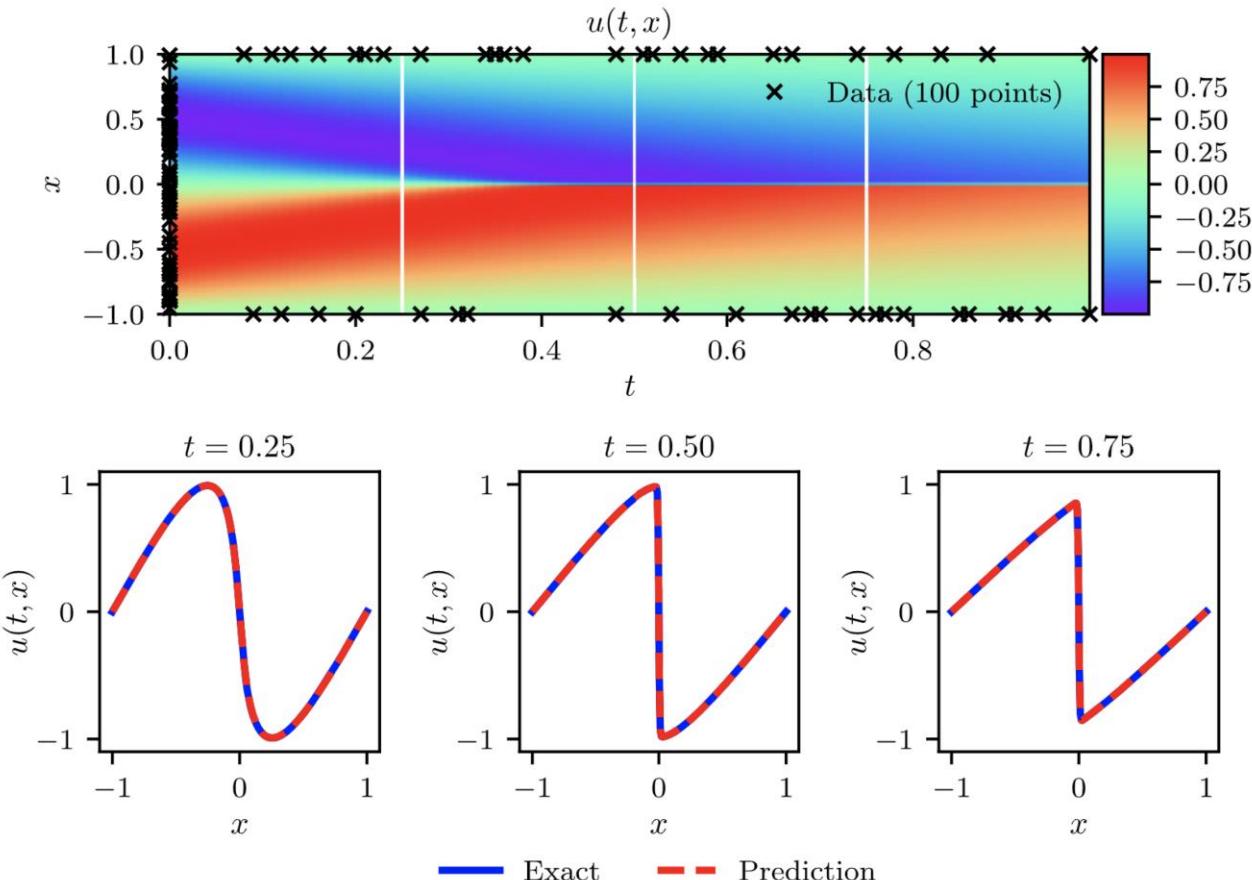
8. Optimizer and PINN Training

```
lr = 1e-3
optimizer = tf.optimizers.Adam(learning_rate=lr)

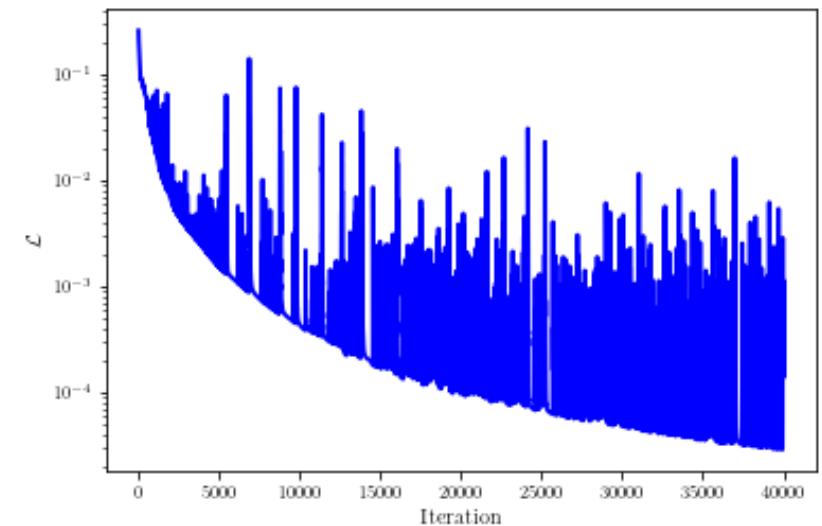
start_time = time.time()
n=0
loss = []
while n <= Nmax:
    loss_ = train_step(W, b, X_u_train_tf, u_train_tf, X_f_train_tf, optimizer, nu)
```

PINN for Burger's Equation: Results

8. Actual vs Predicted Solutions at Different Times



9. Loss Function



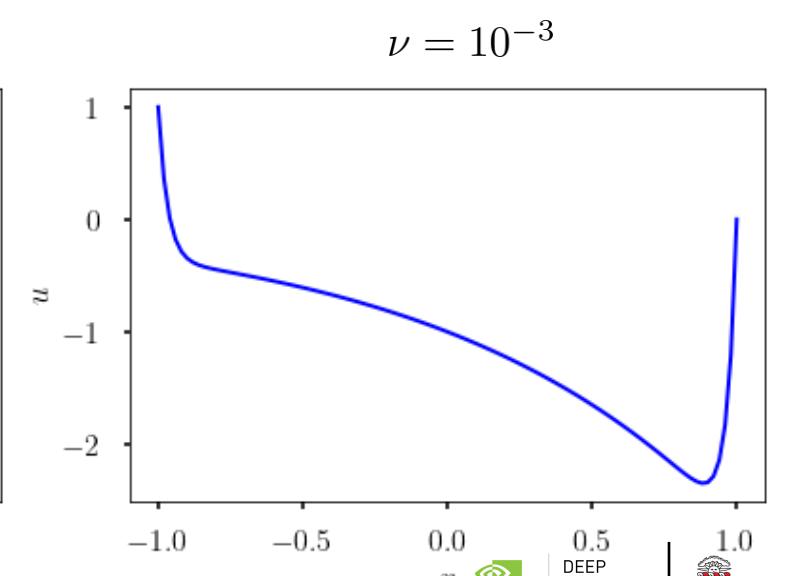
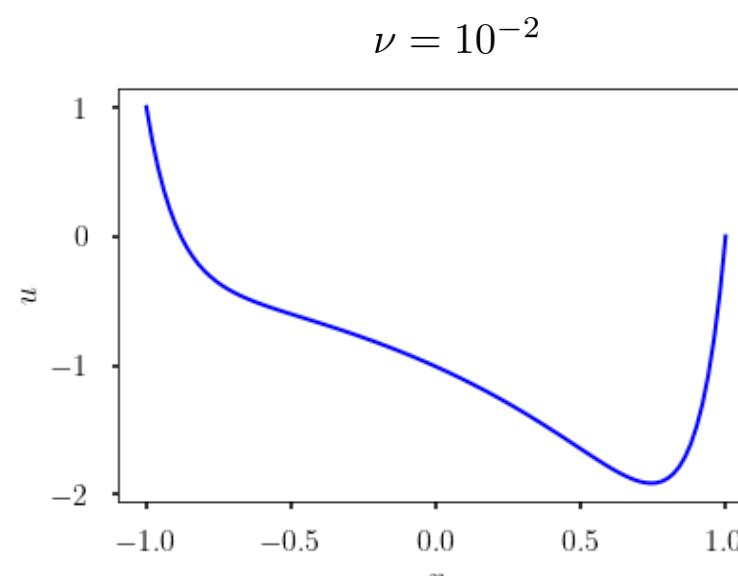
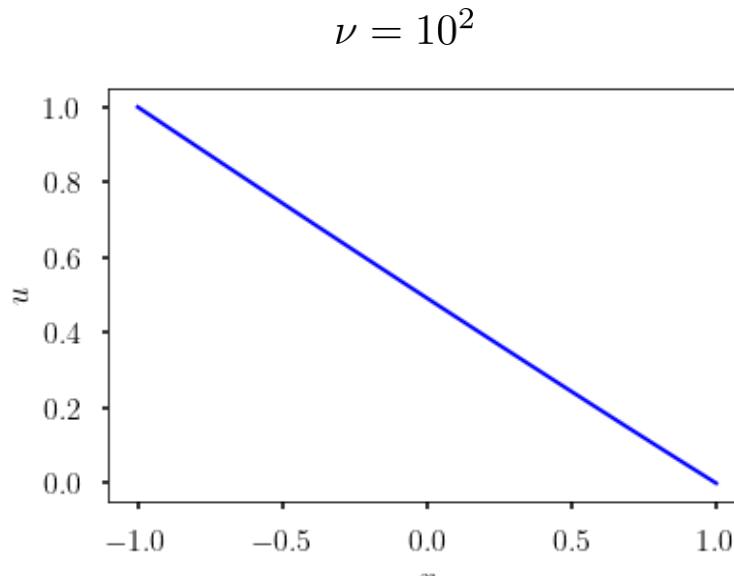
PINNs: Training For Boundary Value Problems

- Define an ODE for given boundary conditions

$$\nu \frac{d^2 u}{dx^2} - u = e^x,$$

where $x \in [-1, 1]$, $u[-1] = 1$, $u[1] = 0$, $\nu = 10^{-3}$.

- Boundary Layer vs Viscosity



Solution using Central Finite Difference Method

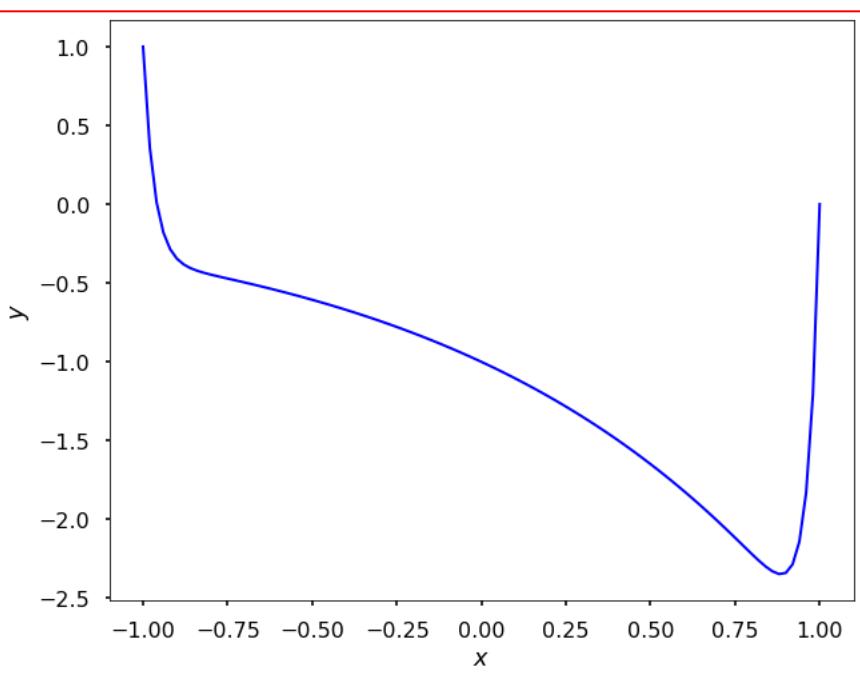


$$u_0 = 1$$

$$\frac{u_{i-1} - 2u_i + u_{i+1}}{h^2} - u_i = \exp(x_i)$$

$$u_m = 0, \quad \nu = 10^{-3}$$

Solution computed using Central Finite-Difference Scheme



```
1  ### Solution of Equation using Central Finite Difference Equation
2  import numpy as np
3  import matplotlib.pyplot as plt
4  plt.style.use('seaborn-poster')
5  %matplotlib inline
6  ### Number of Gridpoints
7  n = 100
8  h = (1+1) / n
9  x = np.linspace(-1,1,n+1)
10 # Difference Operator
11 A = np.zeros((n+1, n+1))
12
13 ## Coefficient For Boundary Condition
14 A[0, 0] = 1
15 A[n, n] = 1
16
17 ### Matrix for Interior Point
18 for i in range(1, n):
19     A[i, i-1] = 1
20     A[i, i] = -(2 + (10**3)*h**2)
21     A[i, i+1] = 1
22
23 # Get b
24 b = np.zeros(n+1)
25 b = (np.exp(x))*h*h*(10**3)
26
27 ##### Imposition Boundary Condition
28 b[0] = 1
29 b[-1] = 0
30 # solve the linear equations
31 y = np.linalg.solve(A, b)
32 ### Plot of Equation
33 plt.figure(figsize=(10,8))
34 plt.plot(x, y, "-", lw=2.0, color="b")
35 plt.xlabel('$x$')
36 plt.ylabel("$y$")
37 plt.show()
```

Approach 1: PINNs with Soft Constraints

- Loss = Boundary Losses + Residual Loses

$$\mathcal{L} = \lambda_{B_1} \mathcal{L}_{B_1} + \lambda_{B_2} \mathcal{L}_{B_2} + \lambda_f \mathcal{L}_f$$

PINN Architecture	
Hyper-parameters	Value
No. of Layers	6
Neurons per layer	4
No. of Iterations	3000
Learning rate	5×10^{-3}
No. of Residuals Points	200
$(\lambda_{B_1}, \lambda_{B_2}, \lambda_f)$	(1, 1, 1)

Approach 1: PINNs with Soft Constraints

- Code

1. Import Modules and Initial Setup

```
import sys
sys.path.insert(0, 'Utilities')
import os

from scipy.interpolate import griddata
from pyDOE import lhs
from plotting import newfig, savefig
from mpl_toolkits.mplot3d import Axes3D
import matplotlib.gridspec as gridspec
from mpl_toolkits.axes_grid1 import make_axes_locatable
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import time
import scipy.io

np.random.seed(seed=1234)
tf.random.set_seed(1234)
tf.config.experimental.enable_tensor_float_32_execution(False)
lb = -1
ub = 1
```

2. Routines for Initialization, Forward Pass and Residual computation

```
# Initialization of Network
def hyper_initial(size):
    in_dim = size[0]
    out_dim = size[1]
    std = np.sqrt(2.0/(in_dim + out_dim))
    return tf.Variable(tf.random.truncated_normal(shape=size, s

# Neural Network
def DNN(X, W, b):
    A = 2.0*(X - lb)/(ub - lb) - 1.0
    L = len(W)
    for i in range(L-1):
        A = tf.tanh(tf.add(tf.matmul(A, W[i]), b[i]))
    Y = tf.add(tf.matmul(A, W[-1]), b[-1])
    return Y

def train_vars(W, b):
    return W + b

def net_u(x,w, b):
    u = DNN(x, w, b)
    return u

#@tf.function(jit_compile=True)
@tf.function
def net_f(x,W, b, nu):
    with tf.GradientTape(persistent=True) as tape1:
        tape1.watch([x])
    with tf.GradientTape(persistent=True) as tape2:
        tape2.watch([x])
        u=net_u(x, W, b)
        u_x = tape2.gradient(u, x)
    del tape2
    u_xx = tape1.gradient(u_x, x)
    del tape1
    f = u_xx - (1/nu)*u-(1/nu)*tf.exp(x)
    return f
```

Approach 1: PINNs with Soft Constraints

- Code

3. Routines for Training and Predict

```
#@tf.function(jit_compile=True)
@tf.function
def train_step(W, b, X_u_train_tf, u_train_tf, X_f_train_tf, opt, nu):
    x_u = X_u_train_tf
    x_f = X_f_train_tf
    with tf.GradientTape() as tape:
        tape.watch([W,b])
        u_nn = net_u(x_u, W, b)
        f_nn = net_f(x_f,W, b, nu)
        loss = tf.reduce_mean(tf.square(u_nn - u_train_tf)) + tf.reduce_mean(tf.square(f_nn))
    grads = tape.gradient(loss, train_vars(W,b))
    opt.apply_gradients(zip(grads, train_vars(W,b)))
    return loss

def predict(X_star_tf, w, b):
    u_pred = net_u(X_star_tf, w, b)
    return u_pred
```

4. Driver

```
nu = 10**-3
noise = 0.0
N_f = 300
Nmax=3000

layers = [1, 4,4,4,4,4,4, 1]
L = len(layers)
W = [hyper_initial([layers[1-1], layers[1]]) for l in range(1, L)]
b = [tf.Variable(tf.zeros([1, layers[1]])) for l in range(1, L)]

x_0 = -1
x_1 = 1
u_0 = 1
u_1 = 0

X_u_train = np.vstack([x_0, x_1])
u_train = np.vstack([u_0, u_1])

X_f_train = lb + (ub-lb)*lhs(1, N_f)
X_f_star = np.linspace(-1,1,200)
X_f_star = X_f_star.reshape((-1,1))

X_u_train_tf = tf.convert_to_tensor(X_u_train, dtype=tf.float32)
u_train_tf = tf.convert_to_tensor(u_train, dtype=tf.float32)
X_f_train_tf = tf.convert_to_tensor(X_f_train, dtype=tf.float32)
X_star_tf = tf.convert_to_tensor(X_f_star, dtype=tf.float32)

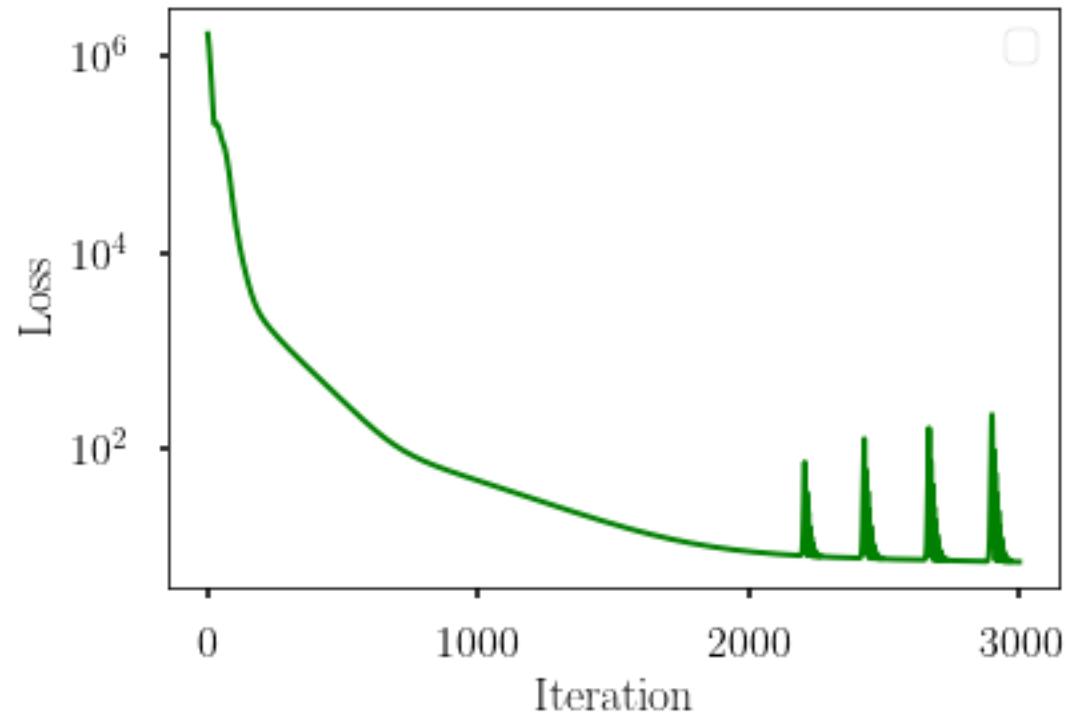
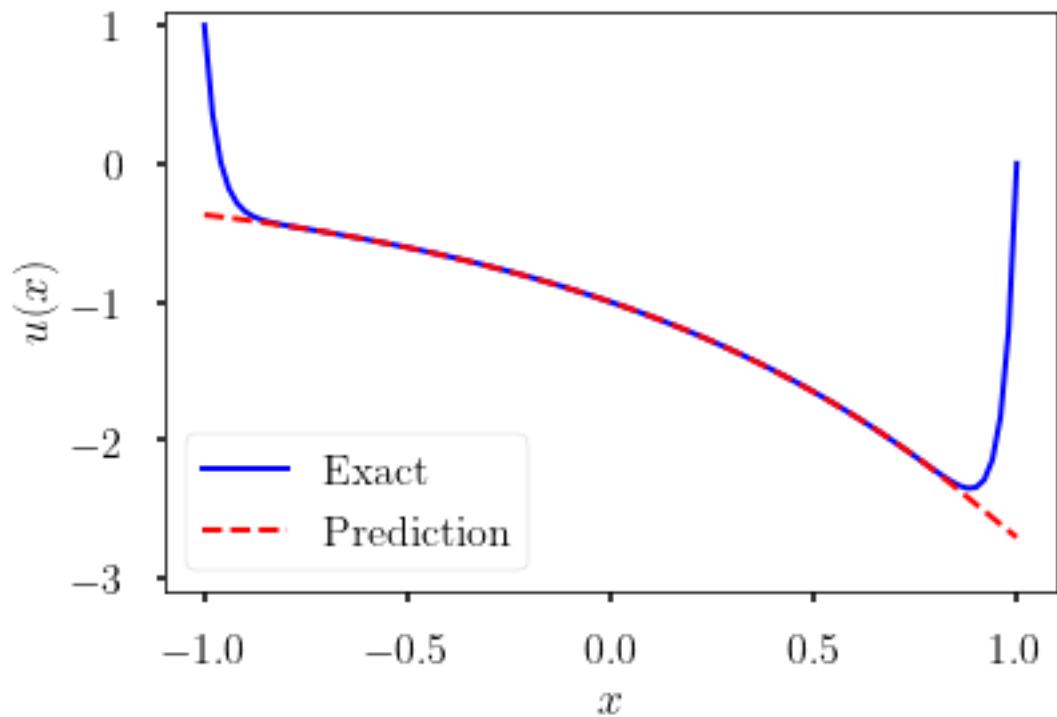
lr = 5e-3
optimizer = tf.optimizers.Adam(learning_rate=lr)

start_time = time.time()
n=0
loss = []
while n <= Nmax:
    loss_ = train_step(W, b, X_u_train_tf, u_train_tf, X_f_train_tf, optimizer, nu)
    loss.append(loss_)
    print(f"Iteration is: {n} and loss is: {loss_}")
    n+=1

elapsed = time.time() - start_time
print('Training time: %.4f' % (elapsed))
```

Approach 1: PINNs with Soft Constraints

- Results



Approach 2: Self-Adaptive PINNs

PDE is Defined as

$$\begin{aligned}\mathcal{N}_{\mathbf{x},t}[u(\mathbf{x}, t)] &= 0, \quad \mathbf{x} \in \Omega, t \in [0, T] \\ (\mathbf{x}, t) &= g(\mathbf{x}, t), \quad \mathbf{x} \in \partial\Omega, t \in [0, T] \\ u(\mathbf{x}, 0) &= h(\mathbf{x}), \quad \mathbf{x} \in \Omega\end{aligned}$$

Loss function

$$\mathcal{L}(\mathbf{w}) = \mathcal{L}_s(\mathbf{w}) + \mathcal{L}_r(\mathbf{w}) + \mathcal{L}_b(\mathbf{w}) + \mathcal{L}_0(\mathbf{w})$$

$$\mathcal{L}_s(\mathbf{w}) = \frac{1}{N_s} \sum_{i=1}^{N_s} |u(\mathbf{x}_s^i, t_s^i; \mathbf{w}) - y_s^i|^2$$

$$\mathcal{L}_r(\mathbf{w}) = \frac{1}{N_r} \sum_{i=1}^{N_r} r(\mathbf{x}_r^i, t_r^i; \mathbf{w})^2$$

$$\mathcal{L}_b(\mathbf{w}) = \frac{1}{N_b} \sum_{i=1}^{N_b} |u(\mathbf{x}_b^i, t_b^i; \mathbf{w}) - g_b^i|^2$$

$$\mathcal{L}_0(\mathbf{w}) = \frac{1}{N_0} \sum_{i=1}^{N_0} |u(\mathbf{x}_0^i, 0; \mathbf{w}) - h_0^i|^2$$

- McClenny L, Braga-Neto U. Self-adaptive physics-informed neural networks using a soft attention mechanism. arXiv preprint arXiv:2009.04544. 2020.

Self-adaptive PINN utilizes the following loss function

$$\mathcal{L}(w, \lambda_r, \lambda_b, \lambda_0) = \mathcal{L}_s(w) + \mathcal{L}_r(w, \lambda_r) + \mathcal{L}_b(w, \lambda_b) + \mathcal{L}_0(w, \lambda_0)$$

where $\lambda_r = (\lambda_r^1, \dots, \lambda_r^{N_r})$, $\lambda_b = (\lambda_b^1, \dots, \lambda_b^{N_b})$, and $\lambda_0 = (\lambda_0^1, \dots, \lambda_0^{N_0})$ are trainable selfadaptation weights for the initial, boundary, and collocation points, respectively, and

$$\begin{aligned}\mathcal{L}_r(w, \lambda_r) &= \frac{1}{N_r} \sum_{i=1}^{N_r} [\lambda_r^i \cdot (x_r^i, t_r^i; w)]^2 \\ \mathcal{L}_b(w, \lambda_b) &= \frac{1}{N_b} \sum_{i=1}^{N_b} [\lambda_b^i (u(x_b^i, t_b^i; w) - g_b^i)]^2 \\ \mathcal{L}_0(w, \lambda_0) &= \frac{1}{N_0} \sum_{i=1}^{N_0} [\lambda_0^i (u(x_0^i, 0; w) - h_0^i)]^2\end{aligned}$$

The key feature of self-adaptive PINNs is that the loss $\mathcal{L}(w, \lambda_r, \lambda_b, \lambda_0)$ is minimized with respect to the network weights w , as usual, but is maximized with respect to the self-adaptation weights $\lambda_r, \lambda_b, \lambda_0$; in other words, training seeks a saddle point

$$\min_w \max_{r, \lambda_b, \lambda_0} \mathcal{L}(w, \lambda_r, \lambda_b, \lambda_0).$$

This can be accomplished by a gradient descentascent procedure, with updates given by:

gradient descent

gradient ascent

$$w^{k+1} = w^k - \eta_k \nabla_w \mathcal{L}(w^k, \lambda_r^k, \lambda_b^k, \lambda_0^k)$$

$$\lambda_r^{k+1} = \lambda_r^k + \eta_k \nabla_{\lambda_r} \mathcal{L}(w^k, \lambda_r^k, \lambda_b^k, \lambda_0^k)$$

$$\lambda_b^{k+1} = \lambda_b^k + \eta_k \nabla_{\lambda_b} \mathcal{L}(w^k, \lambda_r^k, \lambda_b^k, \lambda_0^k)$$

$$\lambda_0^{k+1} = \lambda_0^k + \eta_k \nabla_{\lambda_0} \mathcal{L}(w^k, \lambda_r^k, \lambda_b^k, \lambda_0^k)$$

Approach 2: Implementation

- Code

1. Import Modules and Initial Setup

```
import sys
sys.path.insert(0, 'Utilities/')
import os

from scipy.interpolate import griddata
from pyDOE import lhs
from plotting import newfig, savefig
from mpl_toolkits.mplot3d import Axes3D
import matplotlib.gridspec as gridspec
from mpl_toolkits.axes_grid1 import make_axes_locatable
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import time
import scipy.io

np.random.seed(seed=1234)
tf.random.set_seed(1234)
tf.config.experimental.enable_tensor_float_32_execution(False)
lb = -1
ub = 1
```

Self Adaptive Weight

2. Routines for Initialization, Forward Pass and Residual computation

```
# Initialization of Network
def hyper_initial(size):
    in_dim = size[0]
    out_dim = size[1]
    std = np.sqrt(2.0/(in_dim + out_dim))
    return tf.Variable(tf.random.truncated_normal(shape=size, stddev = std))

# Neural Network
def DNN(X, W, b):
    A = 2.0*(X - lb)/(ub - lb) - 1.0
    L = len(W)
    for i in range(L-1):
        A = tf.tanh(tf.add(tf.matmul(A, W[i]), b[i]))
    Y = tf.add(tf.matmul(A, W[-1]), b[-1])
    return Y

def train_vars_nn(W, b):
    return W + b

def train_vars_total(W, b, lambda_r, lambda_b):
    return W + b + lambda_r + lambda_b

def train_vars_sa(lambda_r, lambda_b):
    return lambda_r + lambda_b

def net_u(x,w, b):
    u = DNN(x, w, b)
    return u

def loss_weight(N_r, N_b):
    alpha_b = tf.Variable(tf.reshape(tf.repeat(1000.0, N_b), (N_b, -1)))
    alpha_r = tf.Variable(tf.ones(shape=[N_r, 1]), dtype=tf.float32)
    return alpha_r, alpha_b
```

Approach 2: Implementation

3. Min-Max Training

```
#@tf.function(jit_compile=True)
@tf.function
def train_step(W, b, X_u_train_tf, u_train_tf, X_f_train_tf, opt, nu, lambda_r,
    x_u = X_u_train_tf
    x_f = X_f_train_tf
    with tf.GradientTape(persistent=True) as tape:
        tape.watch([W,b,lambda_r,lambda_b])
        u_nn = net_u(x_u, W, b)
        f_nn = net_f(x_f,W, b, nu)
        loss_r = tf.square(lambda_r*f_nn)
        loss_b = tf.square(lambda_b*(u_nn-u_train_tf))
        loss = tf.reduce_mean(loss_b) + tf.reduce_mean(loss_r)
        grads = tape.gradient(loss, train_vars_nn(W, b))
        grads_u = tape.gradient(loss, lambda_r)
        grads_b = tape.gradient(loss, lambda_b)
        opt.apply_gradients(zip(grads, train_vars_nn(W,b)))
        opt.apply_gradients(zip([-grads_u], [lambda_r]))
        opt.apply_gradients(zip([-grads_b], [lambda_b]))
    return loss

def predict(X_star_tf, w, b):
    u_pred = net_u(X_star_tf, w, b)
    return u_pred
```

Maximize Loss with
Self Adaptive Weight

4. Driver

```
nu = 10**-3
Nmax= 1500
N_f = 500
N_b = 2

layers = [1, 8,8,8,8,8, 1]
L = len(layers)
W = [hyper_initial([layers[l-1], layers[l]]) for l in range(1, L)]
b = [tf.Variable(tf.zeros([1, layers[l]])) for l in range(1, L)]

alpha_r, alpha_b = loss_weight(N_f, N_b)

x_0 = -1
x_1 = 1
u_0 = 1
u_1 = 0

X_u_train = np.vstack([x_0, x_1])
u_train = np.vstack([u_0, u_1])

X_f_train = lb + (ub-lb)*lhs(1, N_f)
X_f_star = np.linspace(-1,1,200)
X_f_star = X_f_star.reshape((-1,1))

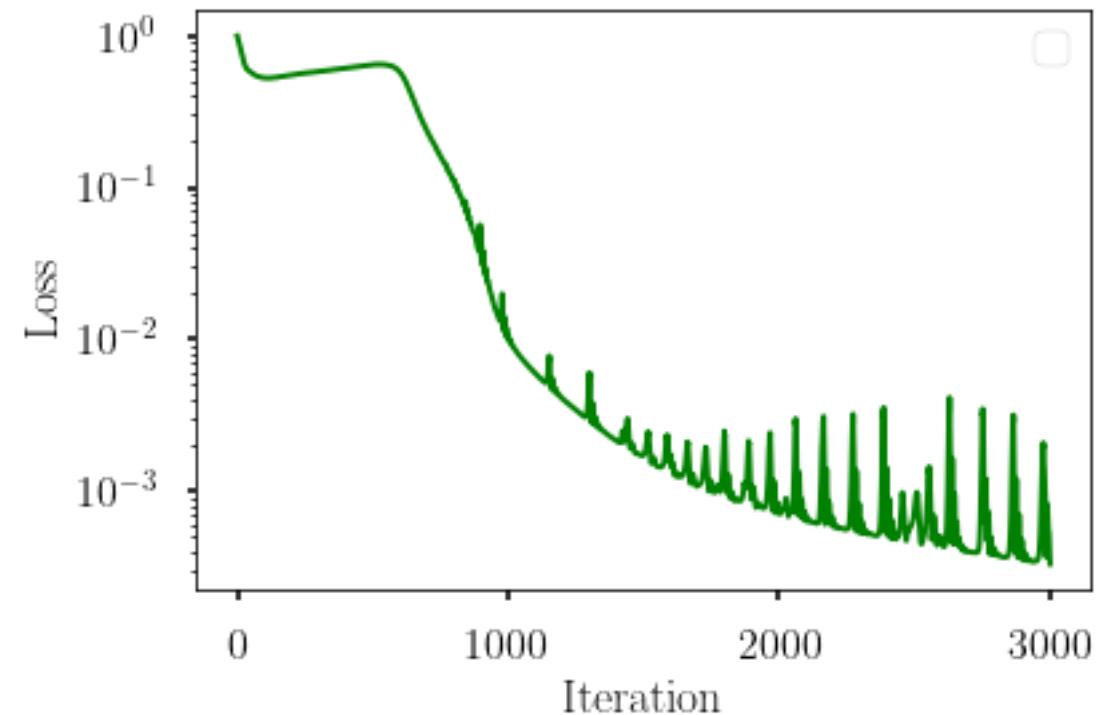
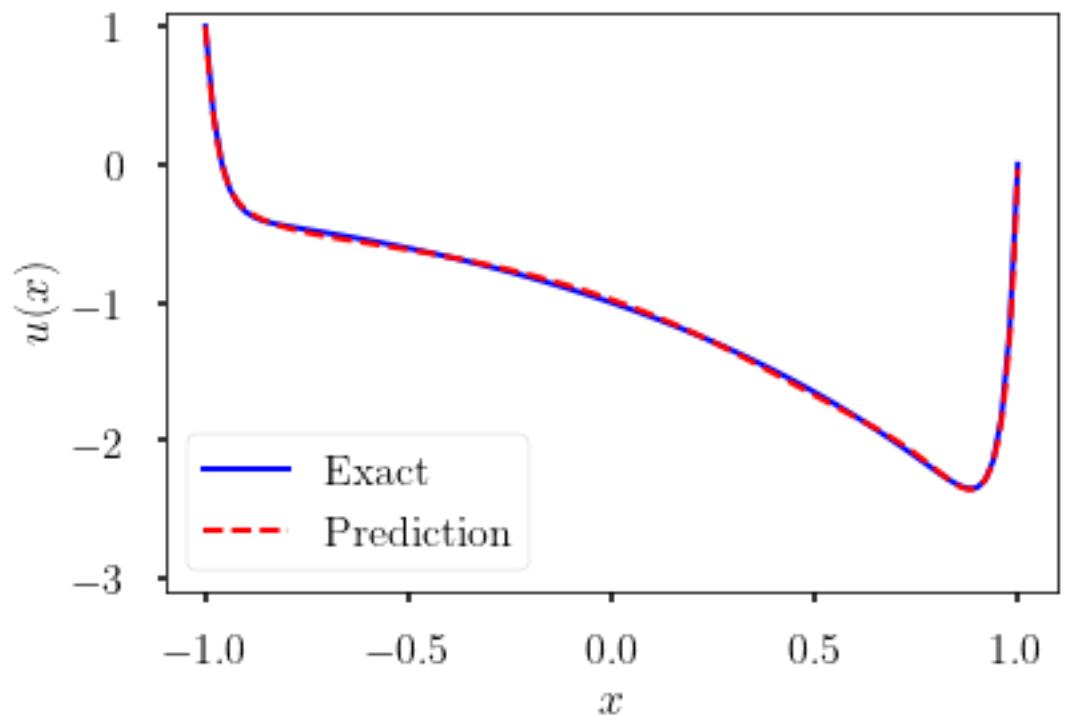
X_u_train_tf = tf.convert_to_tensor(X_u_train, dtype=tf.float32)
u_train_tf = tf.convert_to_tensor(u_train, dtype=tf.float32)
X_f_train_tf = tf.convert_to_tensor(X_f_train, dtype=tf.float32)
X_star_tf = tf.convert_to_tensor(X_f_star, dtype=tf.float32)

lr = 1e-3
optimizer = tf.optimizers.Adam(learning_rate=lr)

start_time = time.time()
n=0
loss = []
while n <= Nmax:
    loss_= train_step(W, b, X_u_train_tf, u_train_tf, X_f_train_tf, optimizer, nu, alpha_r, alpha_b)
    loss.append(loss_)
    print(f"Iteration is: {n} and loss is: {loss_}")
    n+=1

elapsed = time.time() - start_time
print('Training time: %.4f' % (elapsed))
```

Approach 2: Results



Self-Adaptive PINN Architecture	
Hyper-parameters	Value
No. of Layers	8
Neurons per layer	8
No. of Iterations	1500
Learning rate	1×10^{-3}
No. of Residuals Points	500

Approach 3: Dynamic Weights for PINNs

- Wang S, Yu X, Perdikaris P. When and why PINNs fail to train: A neural tangent kernel perspective. *Journal of Computational Physics*. 2022 Jan 15;449:110768.

- Loss = Boundary Losses + Residual Losses

$$\mathcal{L} = \lambda_B \mathcal{L}_B + \lambda_f \mathcal{L}_f$$

- Parameters update

$$\theta^{(k+1)} = \theta^{(k)} - \eta \nabla_{\theta} \mathcal{L}_f - \eta \alpha \nabla_{\theta} \mathcal{L}_B,$$

where $\alpha^{(k+1)} = (1 - \lambda)\alpha^{(k)} + \lambda \hat{\alpha}^{(k+1)}$, with $\hat{\alpha}^{(k+1)} = \frac{\overline{|\nabla_{\theta} \mathcal{L}_f|}}{\overline{|\nabla_{\theta} \mathcal{L}_B|}}$,

- $\overline{|\nabla_{\theta} \mathcal{L}_B|}$ denote the means of $|\nabla_{\theta} \mathcal{L}_B|$.
- $0 \leq \lambda \leq 1$.

Approach 3: Implementation

- Code

1. Import Modules and Initial Setup

```
1 import sys
2 sys.path.insert(0, 'Utilities/')
3 import os
4
5 from scipy.interpolate import griddata
6 from pyDOE import lhs
7 from plotting import newfig, savefig
8 from mpl_toolkits.mplot3d import Axes3D
9 import matplotlib.gridspec as gridspec
10 from mpl_toolkits.axes_grid1 import make_axes_locatable
11 import tensorflow as tf
12 import numpy as np
13 import matplotlib.pyplot as plt
14 import time
15 import scipy.io
```

2. Routines for Initialization, Forward Pass and Residual computation

```
1 np.random.seed(seed=1234)
2 tf.random.set_seed(1234)
3 tf.config.experimental.enable_tensor_float_32_execution(False)
4 #os.environ['TF_ENABLE_AUTO_MIXED_PRECISION'] = '1'
5
6 # Initialization of Network
7 def hyper_initial(size):
8     in_dim = size[0]
9     out_dim = size[1]
10    std = np.sqrt(2.0/(in_dim + out_dim))
11    return tf.Variable(tf.random.truncated_normal(shape=size, stddev = std))
12
13 # Neural Network
14 def DNN(X, W, b):
15     A = X
16     L = len(W)
17     for i in range(L-1):
18         A = tf.tanh(tf.add(tf.matmul(A, W[i]), b[i]))
19     Y = tf.add(tf.matmul(A, W[-1]), b[-1])
20     return Y
21
22 def train_vars(W, b):
23     return W + b
24
25 def net_u(x, t, w, b):
26     u = DNN(tf.concat([x,t],1), w, b)
27     return u
28
29
30 #@tf.function(jit_compile=True)
31 @tf.function
32 def net_f(x,t,W, b, nu):
33     with tf.GradientTape(persistent=True) as tape1:
34         tape1.watch([x, t])
35         with tf.GradientTape(persistent=True) as tape2:
36             tape2.watch([x, t])
37             u=net_u(x,t, W, b)
38             u_t = tape2.gradient(u, t)
39             u_x = tape2.gradient(u, x)
40             u_xx = tape1.gradient(u_x, x)
41             del tape1
42             f = u_t + u*u_x - nu*u_xx
43     return f
44
```

Approach 3: Implementation

3. Training and Dynamic Weights Implementation

```
@tf.function
def train_step(layers, W, b, X_u_train_tf, u_train_tf, X_f_train_tf, opt, nu, lambda_b, beta):
    x_u = X_u_train_tf[:,0:1]
    x_f = X_f_train_tf[:,0:1]
    adaptative_constant_bcs_list = []
    lambda_b_list = []
    with tf.GradientTape(persistent=True) as tape:
        tape.watch([W,b])
        u_nn = net_u(x_u,W, b)
        f_nn = net_f(x_f,W, b, nu)
        bc_loss = tf.reduce_mean(tf.square(u_nn - u_train_tf))
        phys_loss = tf.reduce_mean(tf.square(f_nn))
        loss = lambda_b[-1]*bc_loss + phys_loss
        loss = loss / (1 + lambda_b[-1])
        p_loss = phys_loss/(1 + lambda_b[-1])
        b_loss = lambda_b[-1]*bc_loss/(1 + lambda_b[-1])
        grad_loss = tape.gradient(loss, train_vars(W,b))
        opt.apply_gradients(zip(grad_loss, train_vars(W,b)))
        grads_bc = tape.gradient(bc_loss, train_vars(W,b))
        grads_phys = tape.gradient(phys_loss, train_vars(W,b))
        for i in range(len(layers) - 1):
            adaptative_constant_bcs_list.append(
                tf.reduce_mean(tf.abs(grads_phys[i])) / tf.reduce_mean(tf.abs(grads_bc[i])))
    lambda_b_new = tf.reduce_mean(tf.stack(adaptative_constant_bcs_list))
    lambda_b_new = (1-beta)*lambda_b[-1] + beta*lambda_b_new
    lambda_b_list.append(lambda_b_new)

    return loss, lambda_b_list, p_loss, b_loss
```

Dynamic weights

4. Driver

```
nu = 10**-3
noise = 0.0
N_f = 5000
Nmax=15000

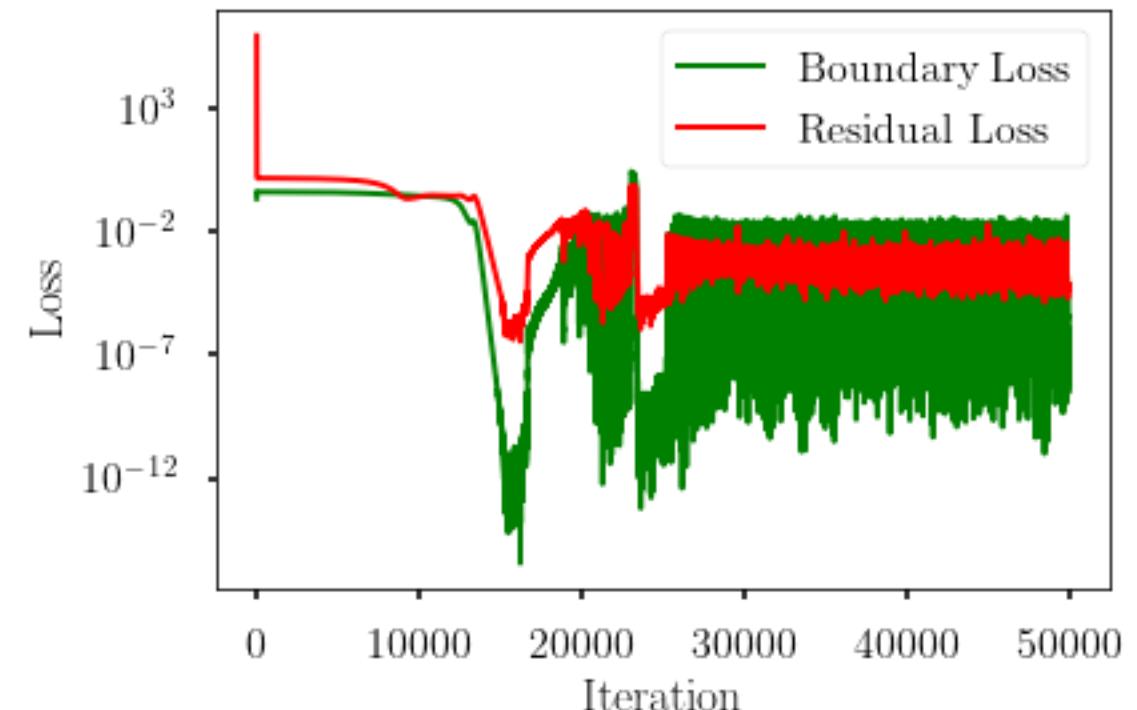
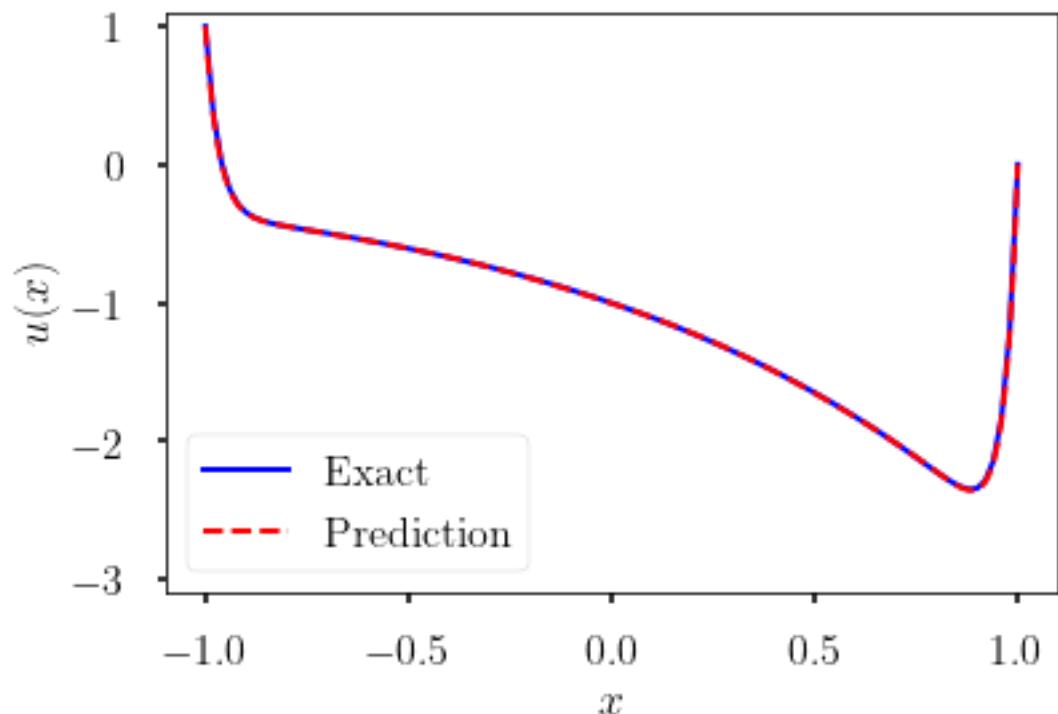
layers = [1, 8,8,8,8,8, 1]
L = len(layers)
W = [hyper_initial([layers[l-1], layers[l]]) for l in range(1, L)]
b = [tf.Variable(tf.zeros([1, layers[l]])) for l in range(1, L)]

x_0 = -1
x_1 = 1
u_0 = 1
u_1 = 0

X_u_train = np.vstack([x_0, x_1])
u_train = np.vstack([u_0, u_1])
X_f_train = lb + (ub-lb)*lhs(1, N_f)
X_f_star = np.linspace(-1,1,200)
X_f_star = X_f_star.reshape((-1,1))
X_u_train_tf = tf.convert_to_tensor(X_u_train, dtype=tf.float32)
u_train_tf = tf.convert_to_tensor(u_train, dtype=tf.float32)
X_f_train_tf = tf.convert_to_tensor(X_f_train, dtype=tf.float32)
X_star_tf = tf.convert_to_tensor(X_f_star, dtype=tf.float32)
lam_b = np.array([1.0])
lam_b_tf = tf.convert_to_tensor(lam_b, dtype=tf.float32)
lambda_b_list = [lam_b_tf]
lr = 1e-3
optimizer = tf.optimizers.Adam(learning_rate=lr)
start_time = time.time()
n=0
loss = []
bc_loss = []
phys_loss = []
beta = 0.1

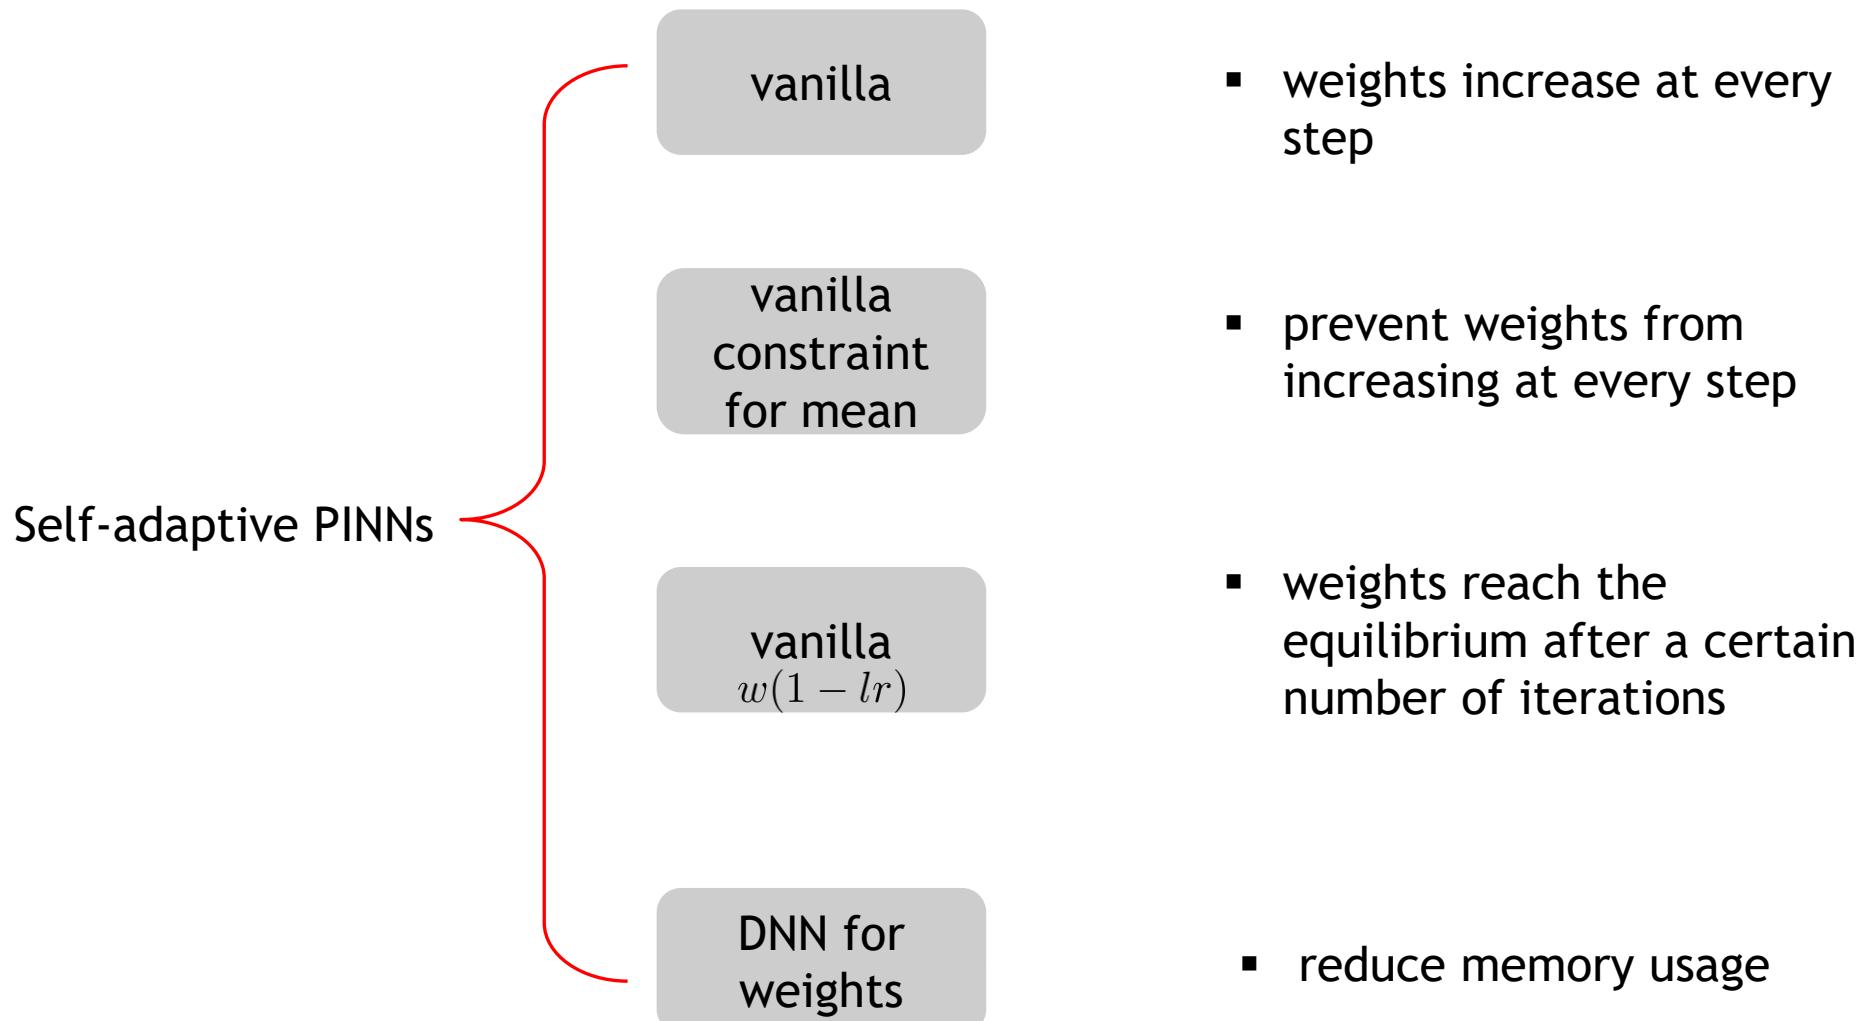
while n <= Nmax:
    loss_, lambda_b_out, phys_loss_, bc_loss_ = train_step(layers, W, b, X_u_train_tf,\n                                            u_train_tf, X_f_train_tf, optimizer, nu,lambda_b_list, beta)
    lambda_b_list = lambda_b_out
    loss.append(loss_)
    bc_loss.append(bc_loss_)
    phys_loss.append(phys_loss_)
```

Approach 3: Results

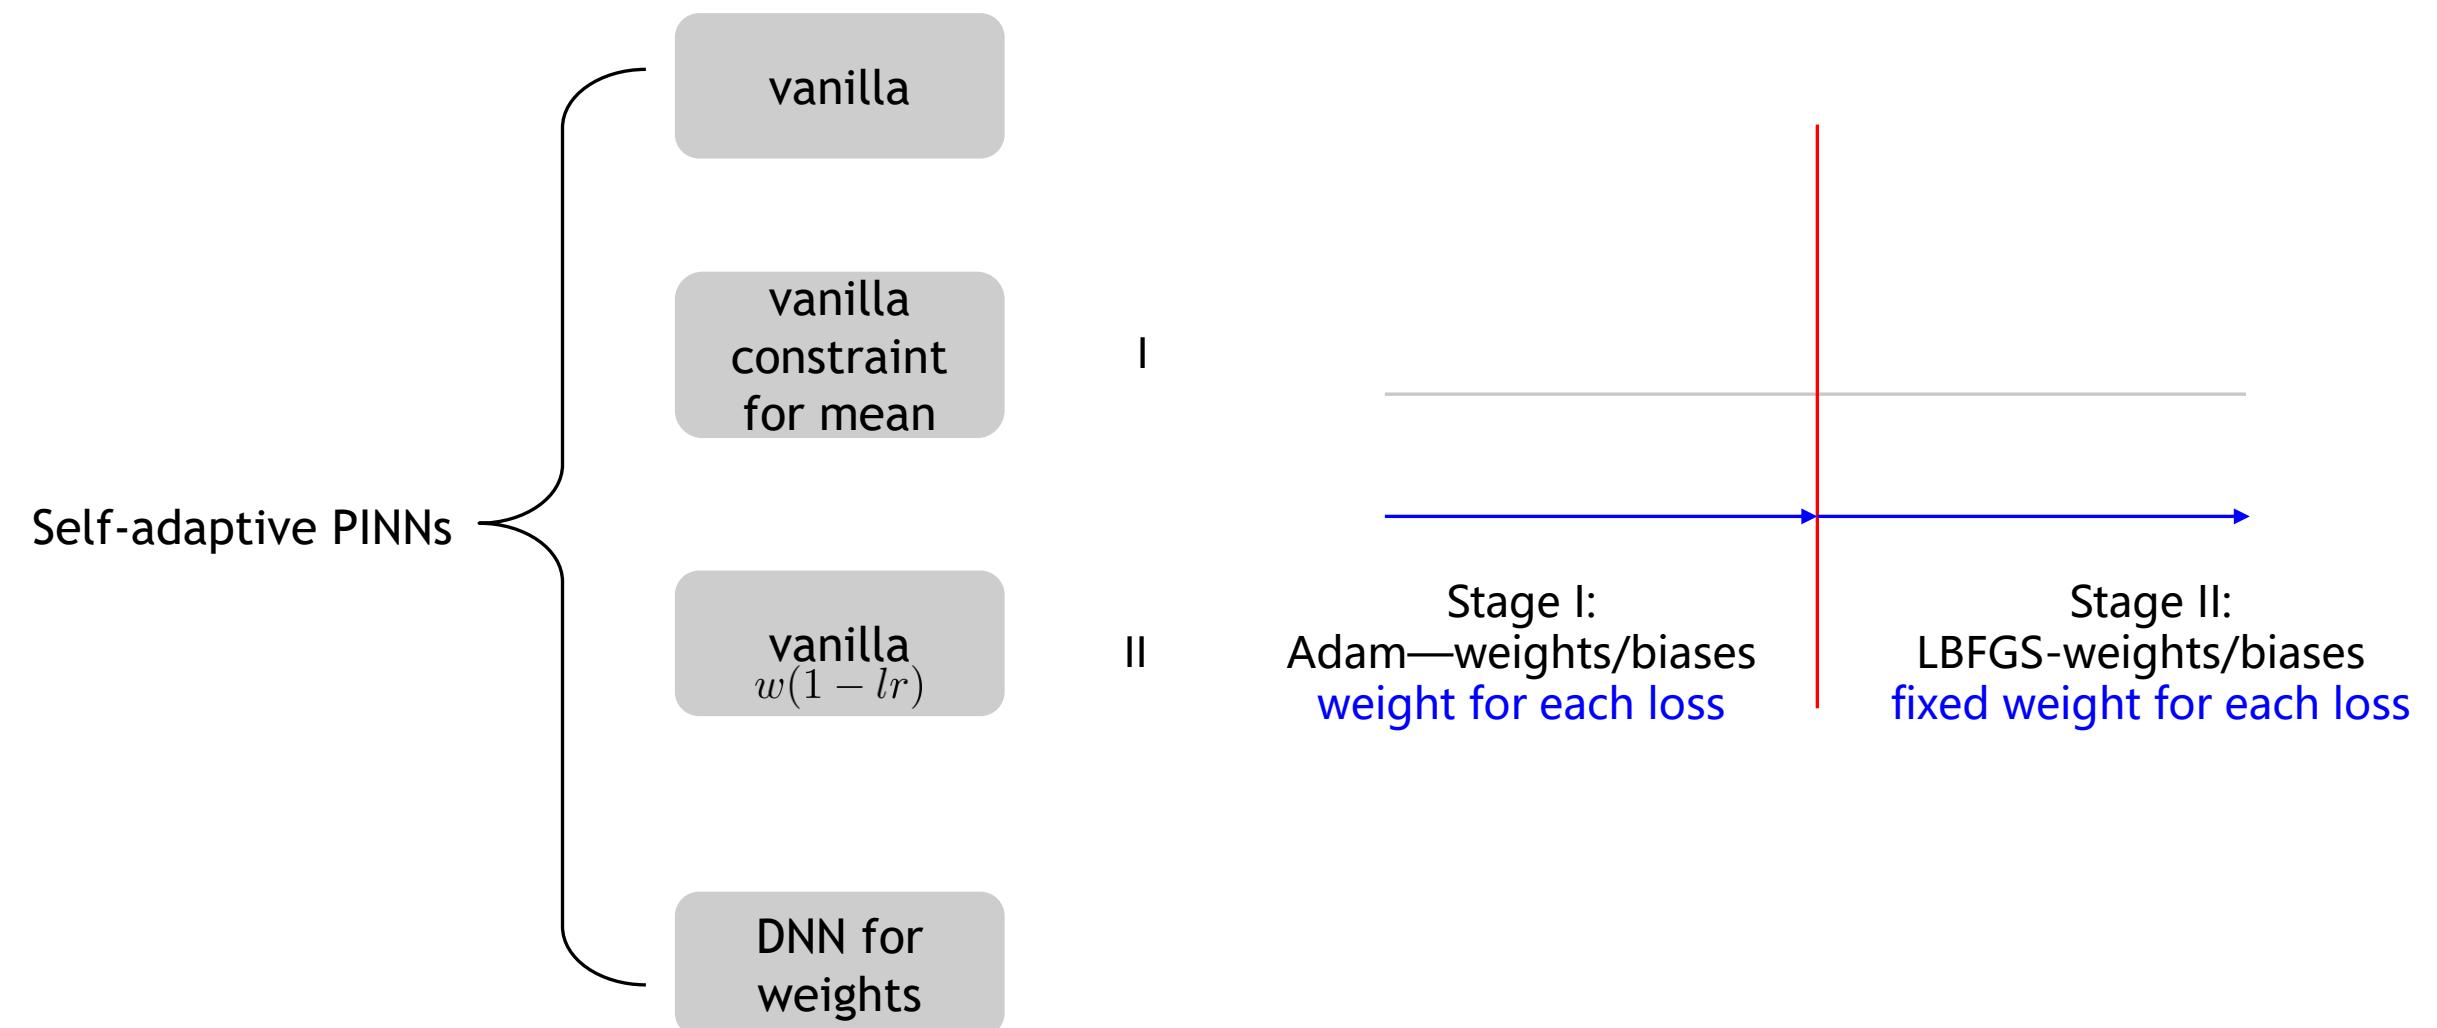


Adaptive weights PINN Architecture	
Hyper-parameters	Value
No. of Layers	6
Neurons per layer	8
No. of Iterations	5000
Learning rate	1×10^{-3}
No. of Residuals Points	500

Different Versions of Self-Adaptive Weights in PINNs



Training



Self-Adaptive Weights: Alan-Cahn Equation

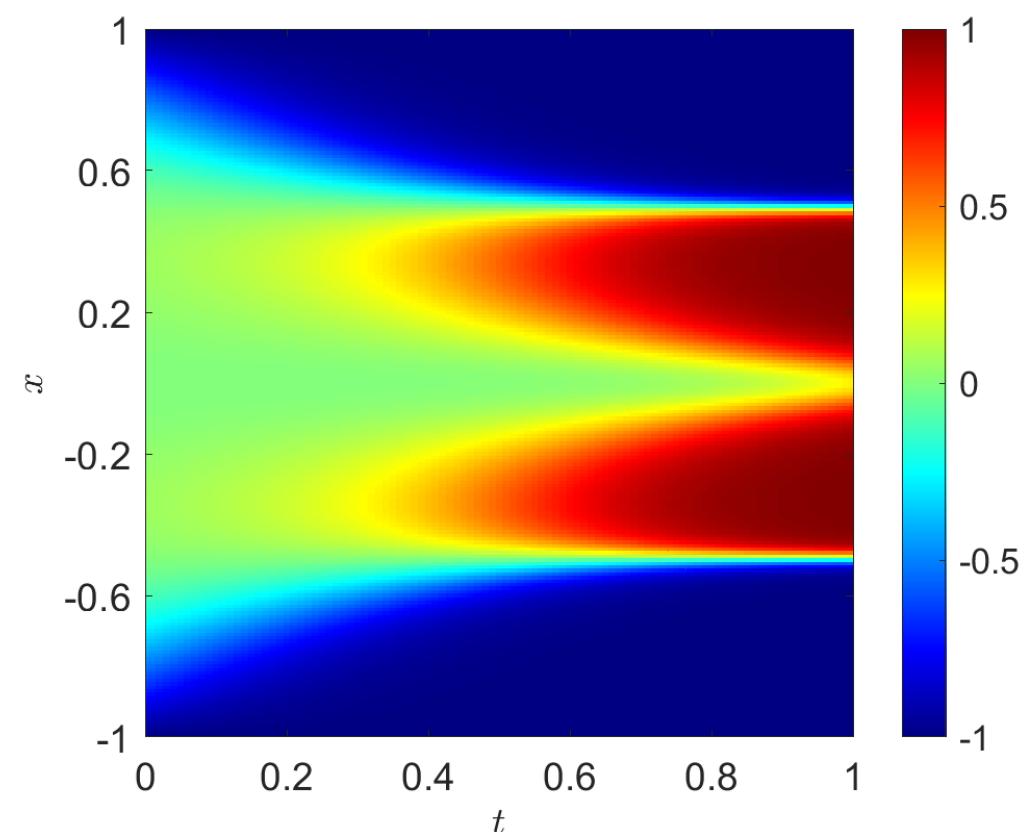
- Alan-Cahn Equation

$$\partial_t u - D \partial_x^2 u + 5(u^3 - u) = 0, D = 0.0001$$
$$t \in [0, 1], x \in [-1, 1]$$

- Initial and Boundary condition

$$u(x, 0) = x^2 \cos(\pi x) (IC)$$
$$u(-1, t) = u(1, t) = -1 (BCs)$$

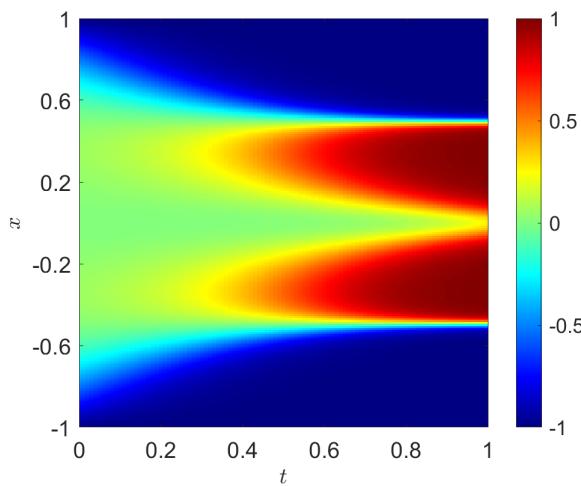
- DNN: 4 hidden layers with 128 neurons per layer
- 20,000 random points for residual
- 200 uniform points for initial conditions
- 100 uniform points for boundary conditions



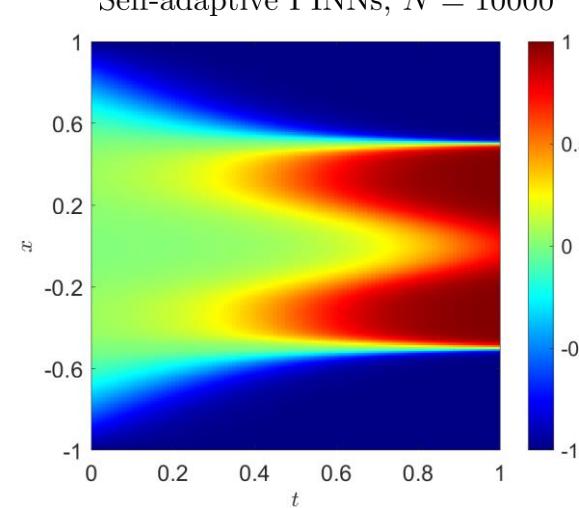
Reference Solution for Alan-Cahn Equation

Self-Adaptive Weights: Alan-Cahn Equation (vanilla)

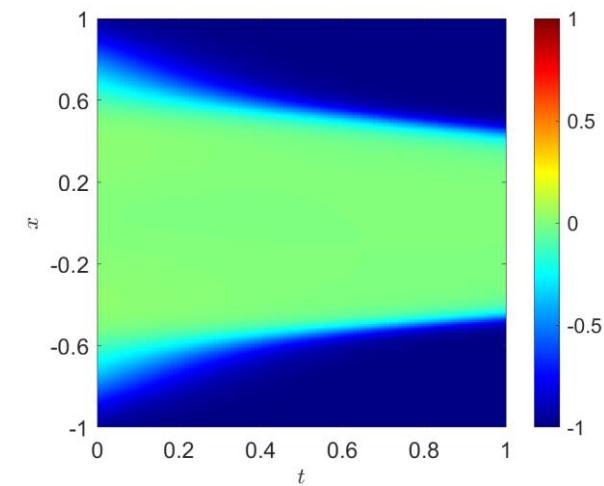
Reference solution



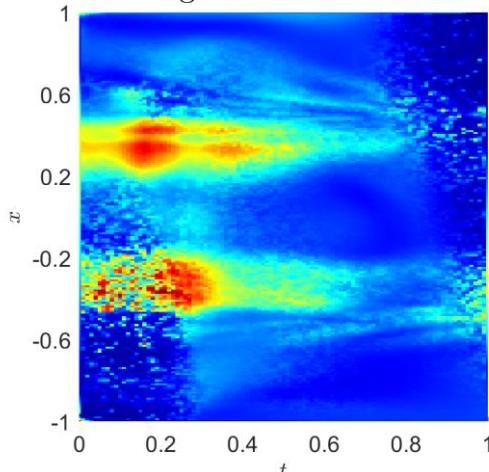
Self-adaptive PINNs, $N = 10000$



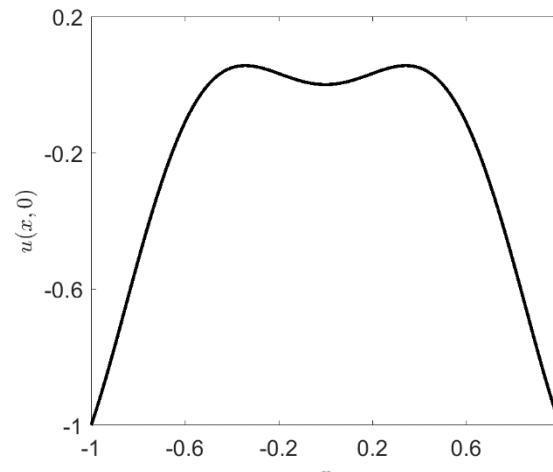
Vanilla PINNs



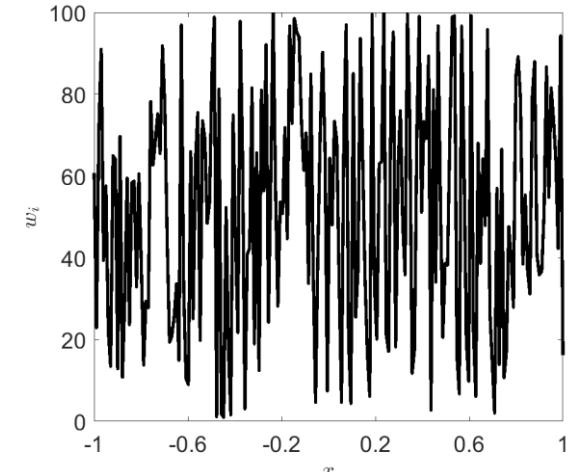
Weights for residuals



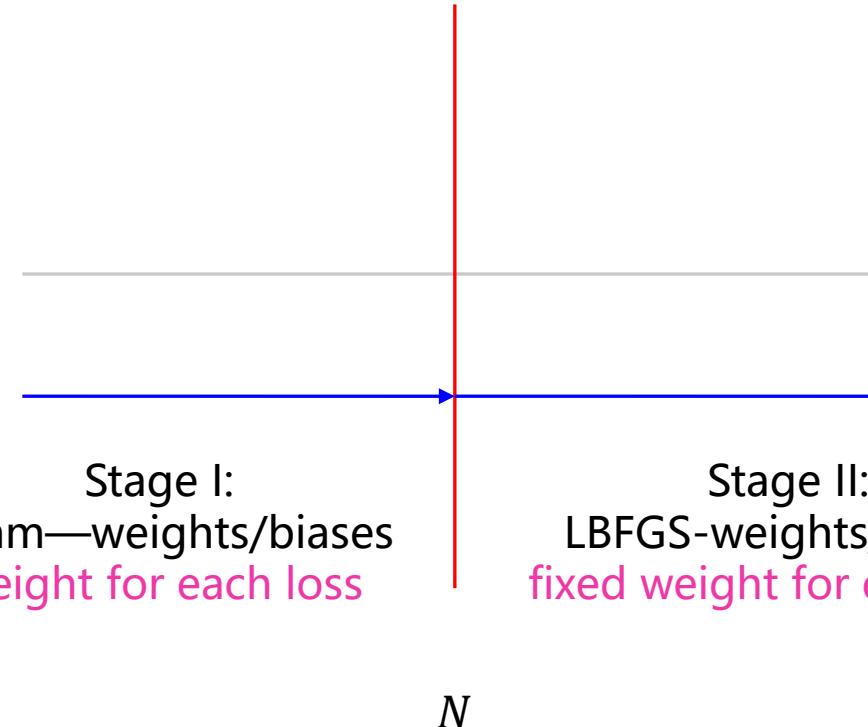
Initial condition



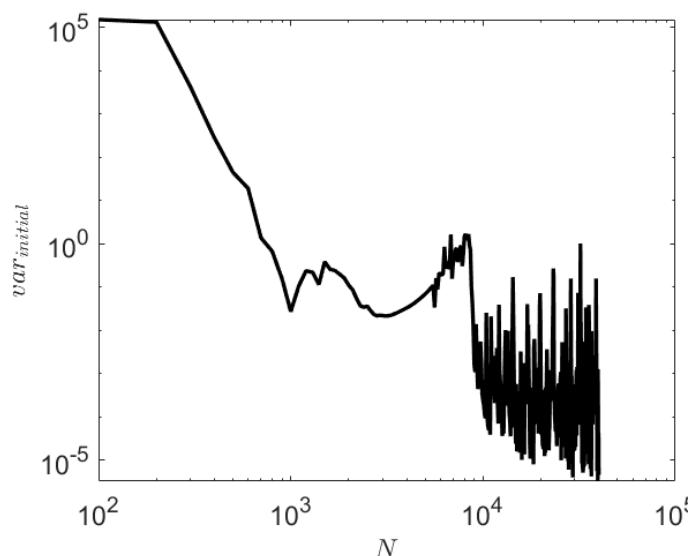
Weights for initial condition



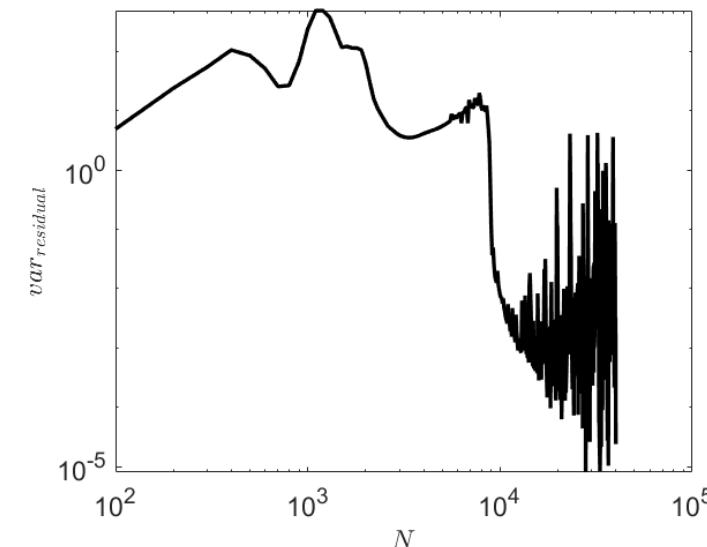
I. Variance for Losses



Variance of loss in initial condition

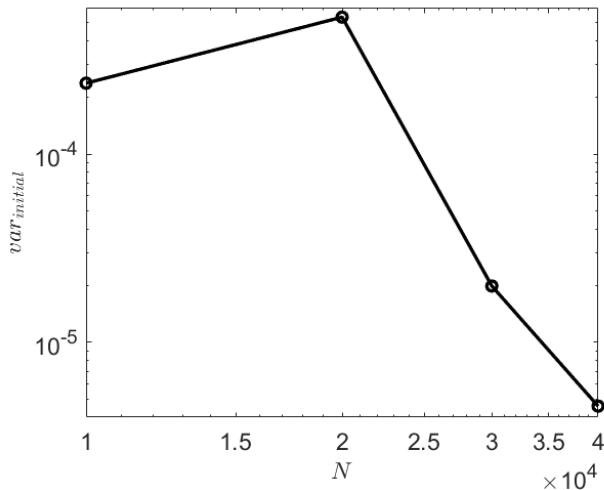


Variance of loss in residuals

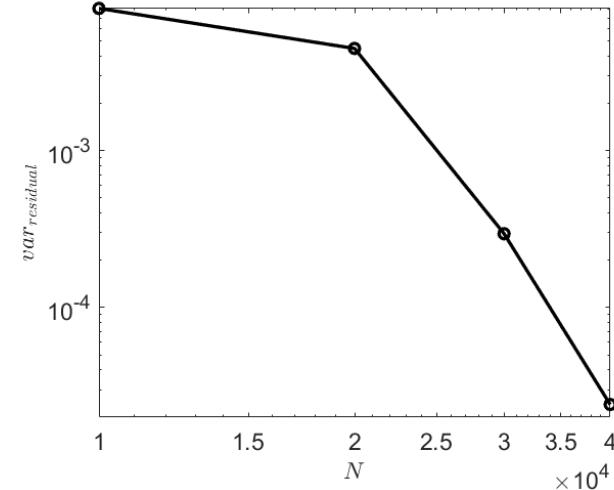


I. Variance for Losses

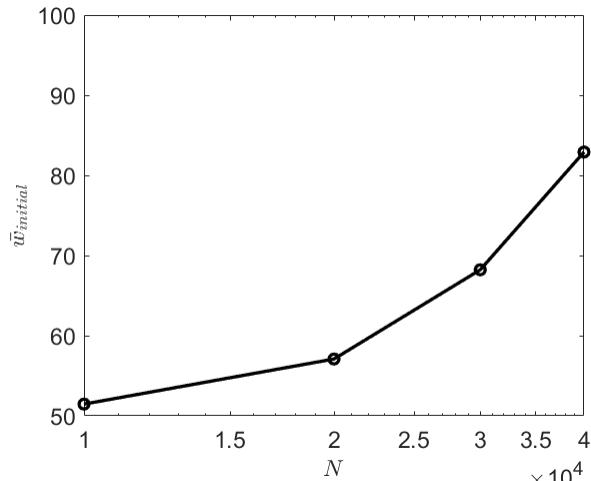
Variance of loss in initial condition



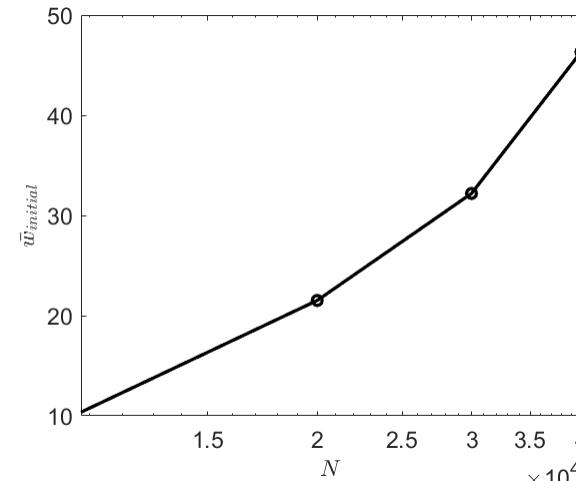
Variance of loss in residuals



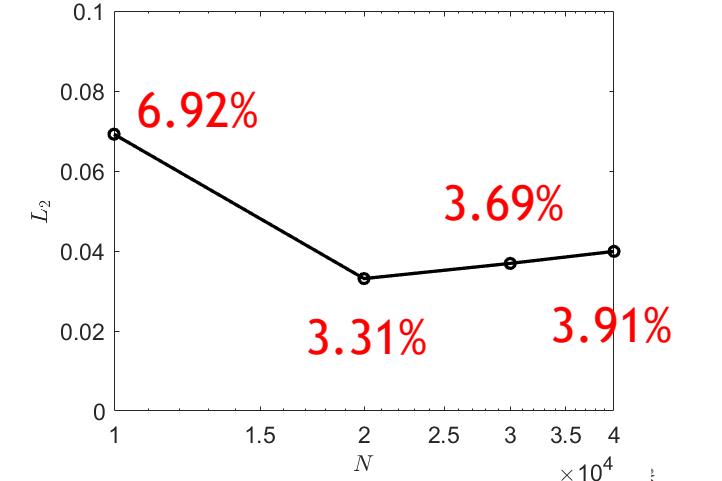
Average of self adaptive weight for initial conditions



Average of self adaptive weight for residual



L_2 Error between actual and predicted solution



II. Variance for Losses: $\mathbf{W}' = \mathbf{W}(1 - lr)$

Learning rates are important!

The key feature of self-adaptive PINNs is that the loss $\mathcal{L}(\mathbf{w}, \boldsymbol{\lambda}_r, \boldsymbol{\lambda}_b, \boldsymbol{\lambda}_0)$ is minimized with respect to the network weights \mathbf{w} , as usual, but is maximized with respect to the self-adaptation weights $\boldsymbol{\lambda}_r, \boldsymbol{\lambda}_b, \boldsymbol{\lambda}_0$; in other words, training seeks a saddle point

$$\min_{\mathbf{w}} \max_{\boldsymbol{\lambda}_r, \boldsymbol{\lambda}_b, \boldsymbol{\lambda}_0} \mathcal{L}(\mathbf{w}, \boldsymbol{\lambda}_r, \boldsymbol{\lambda}_b, \boldsymbol{\lambda}_0)$$

gradient ascent

This can be accomplished by a gradient descent/ascent procedure, with updates given by:

$$\mathbf{w}^{k+1} = \mathbf{w}^k - \eta_k \nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w}^k, \boldsymbol{\lambda}_r^k, \boldsymbol{\lambda}_b^k, \boldsymbol{\lambda}_0^k)$$

$$\boldsymbol{\lambda}_r^{k+1} = \boldsymbol{\lambda}_r^k + \eta_k \nabla_{\boldsymbol{\lambda}_r} \mathcal{L}(\mathbf{w}^k, \boldsymbol{\lambda}_r^k, \boldsymbol{\lambda}_b^k, \boldsymbol{\lambda}_0^k)$$

$$\boldsymbol{\lambda}_b^{k+1} = \boldsymbol{\lambda}_b^k + \eta_k \nabla_{\boldsymbol{\lambda}_b} \mathcal{L}(\mathbf{w}^k, \boldsymbol{\lambda}_r^k, \boldsymbol{\lambda}_b^k, \boldsymbol{\lambda}_0^k)$$

$$\boldsymbol{\lambda}_0^{k+1} = \boldsymbol{\lambda}_0^k + \eta_k \nabla_{\boldsymbol{\lambda}_0} \mathcal{L}(\mathbf{w}^k, \boldsymbol{\lambda}_r^k, \boldsymbol{\lambda}_b^k, \boldsymbol{\lambda}_0^k)$$

$$\begin{aligned} & \Rightarrow \boldsymbol{\lambda}^{k+1} = \boldsymbol{\lambda}^k + \eta_k \nabla_{\boldsymbol{\lambda}_k} L, \quad \nabla_{\boldsymbol{\lambda}_k} L > 0 \\ & \Rightarrow \boldsymbol{\lambda}^{k+1} = \boldsymbol{\lambda}^{k+1}(1 - \beta), \quad \beta = lr \end{aligned}$$

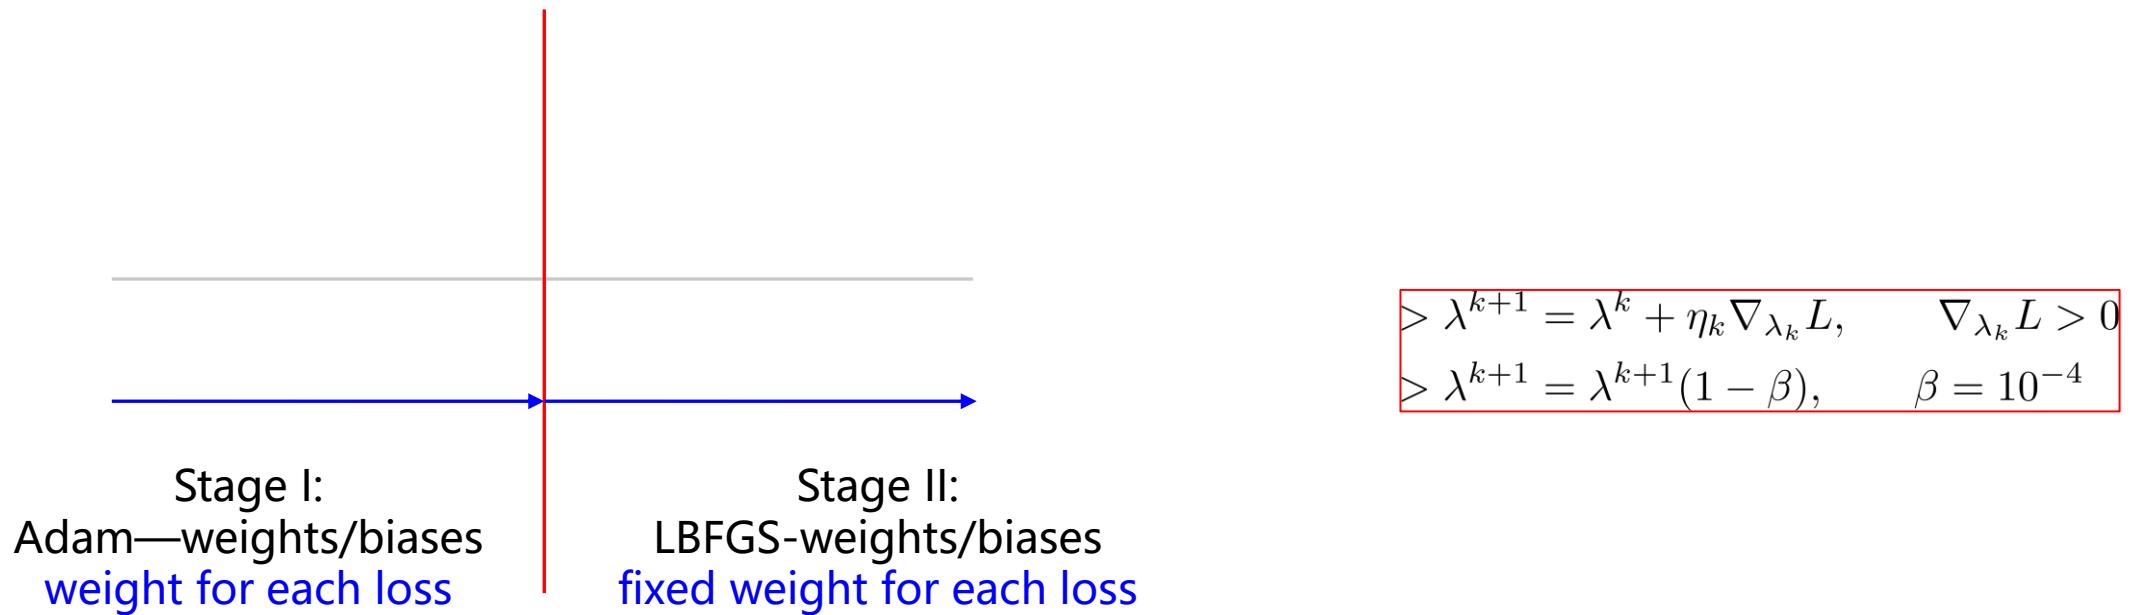
1. First increase λ , then decrease
2. Increment and decrement in λ is defined as

$$\Delta\lambda_+^{k+1} = \eta_k \nabla_{\lambda_k} L, \quad \Delta\lambda_-^{k+1} = -\beta\lambda^{k+1}$$

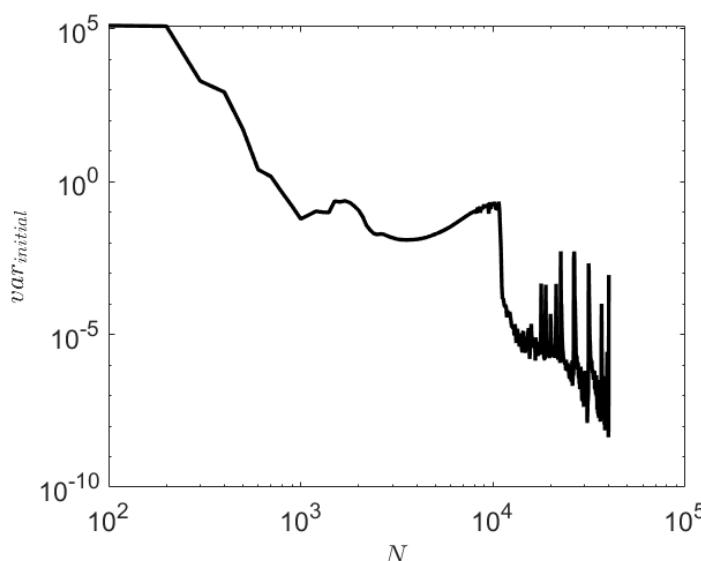
3. Combine λ

$$\Delta\lambda^{k+1} = \Delta\lambda_+^{k+1} + \Delta\lambda_-^{k+1} = \eta_k \nabla_{\lambda_k} L - \beta\lambda^{k+1} \quad \left\{ \begin{array}{l} \eta_k \nabla_{\lambda_k} L \uparrow, \lambda^{k+1} \uparrow \\ \eta_k \nabla_{\lambda_k} L \downarrow, \lambda^{k+1} \downarrow \end{array} \right.$$

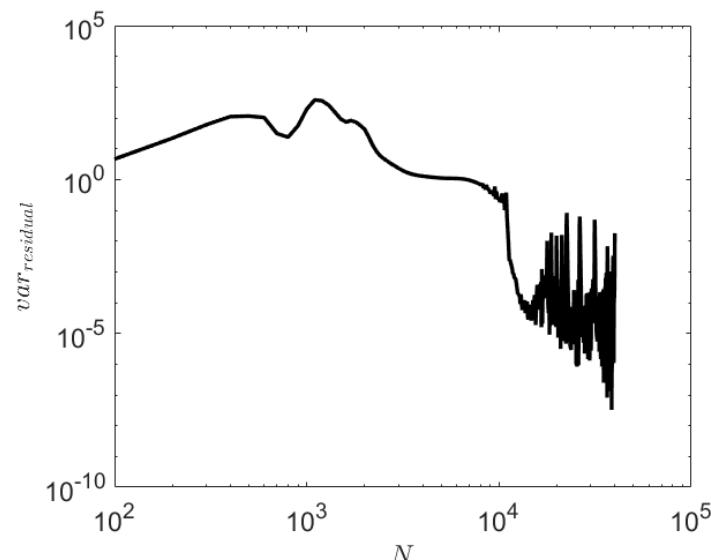
II. Variance for Losses: $W' = W(1 - lr)$



Variance of loss in initial condition

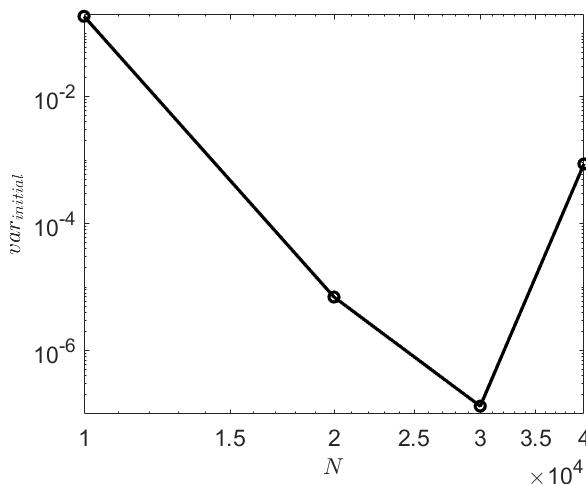


Variance of loss in residuals

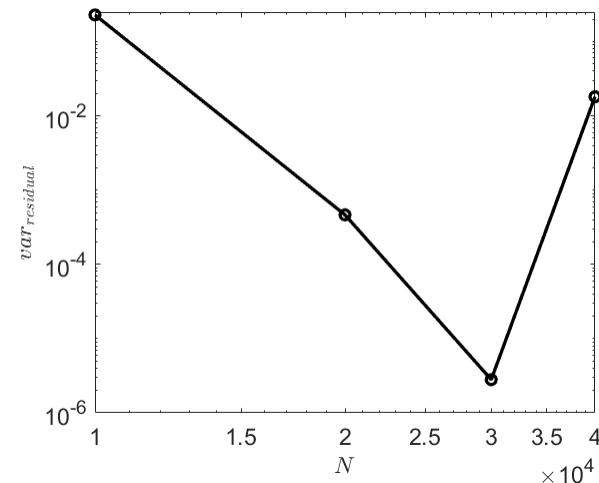


II. Variance for Losses: $W' = W(1 - lr)$

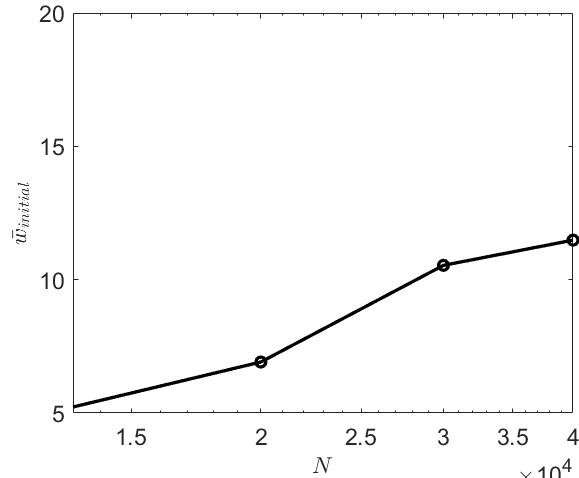
Variance of loss in initial condition



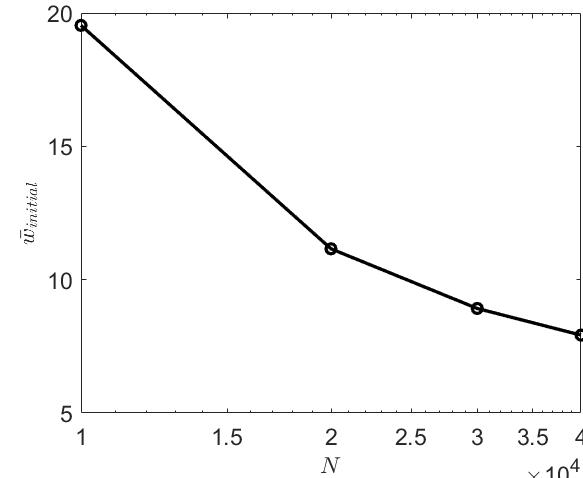
Variance of loss in residuals



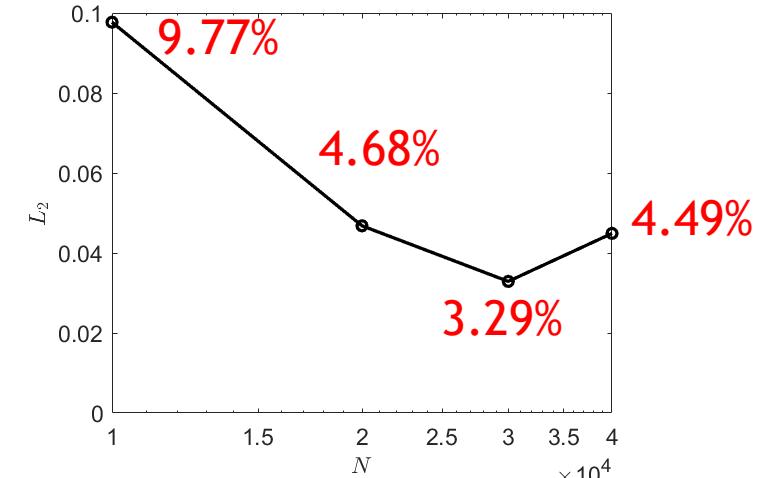
Average of self adaptive weight for initial conditions



Average of self adaptive weight for residual



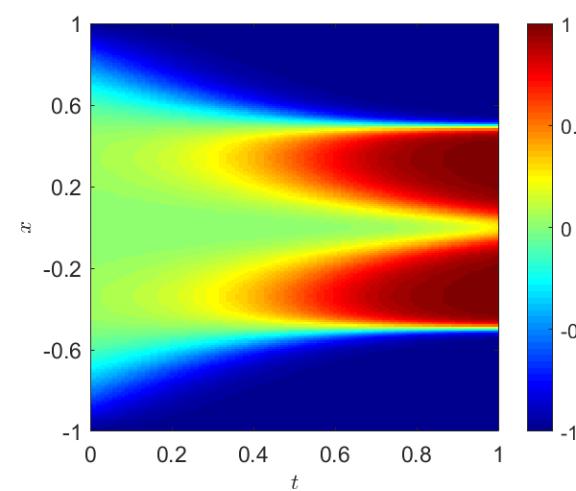
L_2 Error between actual and predicted solution



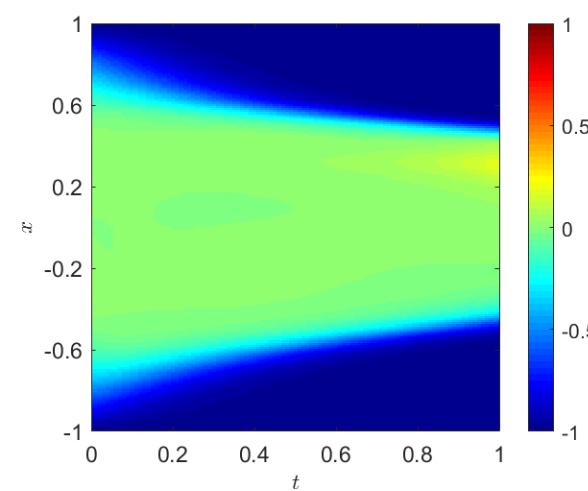
II. Variance for Losses: $W' = W(1 - lr)$ Effect of Learning Rate

$$\begin{aligned} &> \lambda^{k+1} = \lambda^k + \eta_k \nabla_{\lambda_k} L, \quad \nabla_{\lambda_k} L > 0 \\ &> \lambda^{k+1} = \lambda^{k+1}(1 - \beta), \quad \beta = 10^{-3} \end{aligned}$$

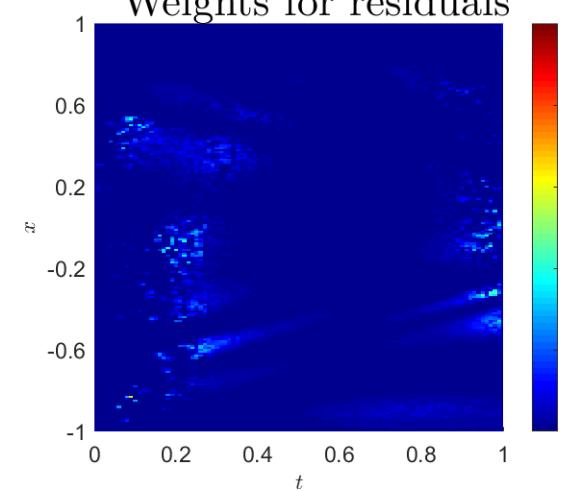
Reference solution



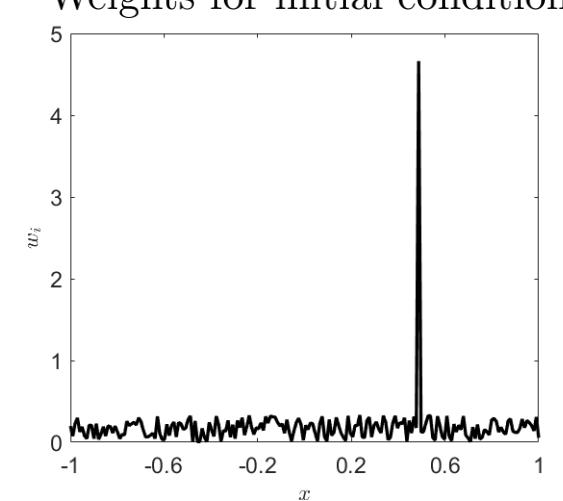
Self-adaptive PINNs ($N=10000$)



Weights for residuals



Weights for initial condition



Hard Constraints: Dirichlet BC and IC

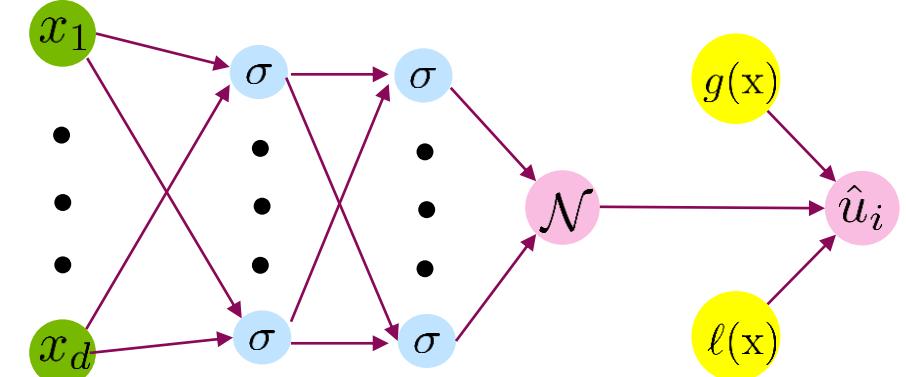
Dirichlet BC or IC : $u_i(\mathbf{x}) = g_0(\mathbf{x}), \quad \mathbf{x} \in \Gamma_D,$

Trial solution: $\hat{u}_i(\mathbf{x}; \boldsymbol{\theta}_u) = g(\mathbf{x}) + \ell(\mathbf{x})\mathcal{N}(\mathbf{x}; \boldsymbol{\theta}_u),$

$g(x)$ is a continuous extension of $g_0(x)$ from Γ_D to Ω where $\mathcal{N}(\mathbf{x}; \boldsymbol{\theta}_u)$ is the network output, and ℓ is a function satisfying the following two conditions:

$$\begin{cases} \ell(\mathbf{x}) = 0, & \mathbf{x} \in \Gamma_D \\ \ell(\mathbf{x}) > 0, & \mathbf{x} \in \Omega - \Gamma_D \end{cases}$$

For example: $u(0) = 0, u(1) = 1 : g(x) = x, \ell(x) = x(1 - x)$



[1] P. L. Lagaris, L. H. Tsoukalas, S. Safarkhani, and I. E. Lagaris, Systematic construction of neural forms for solving partial differential equations inside rectangular domains, subject to initial, boundary and interface conditions, *Int. J. Artif. Intell.*, 29 (2020), 2050009.

[2] L. Lu, R. Pestourie, W. Yao, Z. Wang, F. Verdugo, & S. G. Johnson. Physics-informed neural networks with hard constraints for inverse design. *SIAM Journal on Scientific Computing*, 43(6), B1105–B1132, 2021.

Hard Constraints: Periodic BC and IC

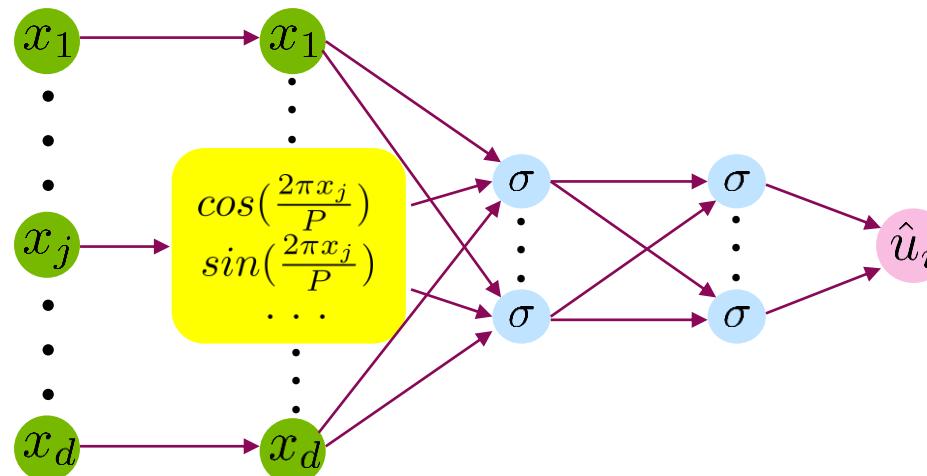
$u_i(x)$ is a periodic function with respect to x_j of the period P .

$\Rightarrow u_i(x)$ can be decomposed into weighted summation of the Fourier series:

$$\{1, \cos\left(\frac{2\pi x_j}{P}\right), \sin\left(\frac{2\pi x_j}{P}\right), \cos\left(\frac{4\pi x_j}{P}\right), \sin\left(\frac{4\pi x_j}{P}\right), \dots\}$$

\Rightarrow replace x_j with Fourier basis functions to impose periodicity.

$$u_i(x) = \mathcal{N} \left(x_1, \dots, x_{j-1}, \left[1, \cos\left(\frac{2\pi x_j}{P}\right), \sin\left(\frac{2\pi x_j}{P}\right), \cos\left(\frac{4\pi x_j}{P}\right), \sin\left(\frac{4\pi x_j}{P}\right), \dots \right], x_{j+1}, \dots, x_d \right)$$



can use as few terms as $\{\cos\left(\frac{2\pi x_j}{P}\right), \sin\left(\frac{2\pi x_j}{P}\right)\}$

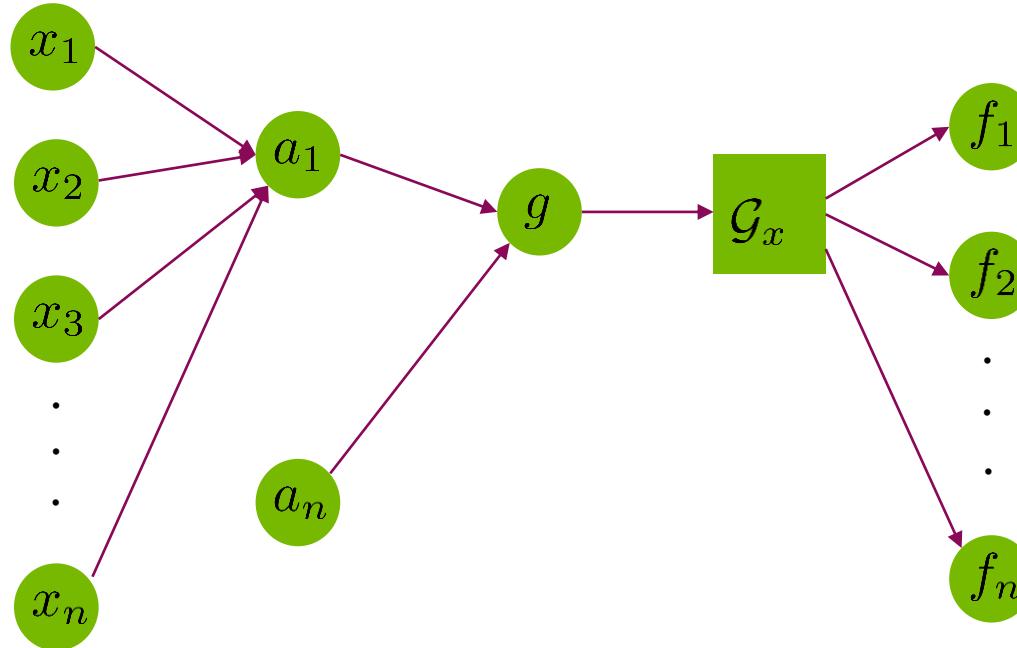
Linearly Constrained Neural Networks

For example: divergence-free neural networks

$$\nabla \cdot \mathbf{f} = \frac{\partial f_1}{\partial x} + \frac{\partial f_2}{\partial y} = 0$$

Divergence of the curl is 0 : $\nabla \cdot (\nabla \times g) \equiv 0$, so choose

$$\mathbf{f} = \nabla \times g$$



Linearly Constrained Neural Networks: Implementation

- Problem definition

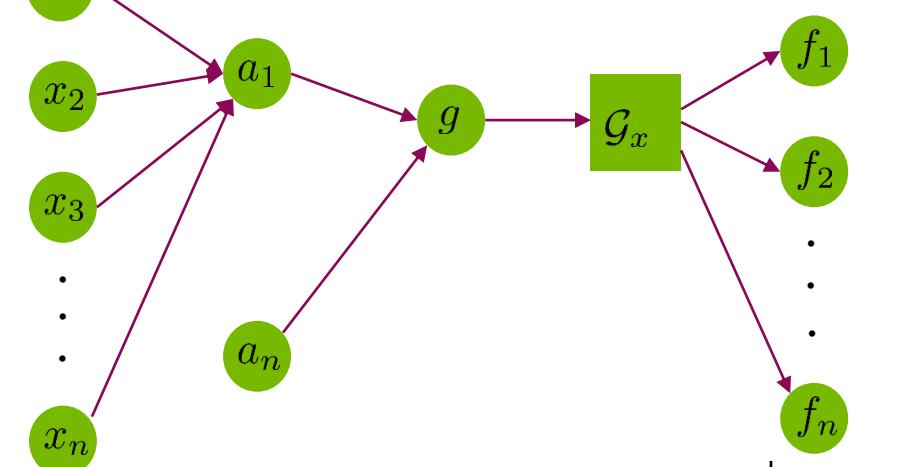
$$f_1(x_1, x_2) = \exp(-ax_1x_2)(ax_1 \sin(x_1x_2) - x_1 \cos(x_1x_2)),$$

$$f_2(x_1, x_2) = \exp(-ax_1x_2)(x_2 \cos(x_1x_2) - ax_2 \sin(x_1x_2)),$$

where a is a constant.

- This vector field satisfies the constraint $\frac{\partial f_1}{\partial x_1} + \frac{\partial f_2}{\partial x_2} = 0$.
- A neural network based model satisfying these constraints is given by

$$\mathbf{f} = \begin{bmatrix} \frac{\partial}{\partial x_2} \\ -\frac{\partial}{\partial x_1} \end{bmatrix} g$$



Linearly Constrained Neural Networks: Code

1. Import Modules and Initial Setup

```
import torch
from matplotlib import pyplot as plt
from torch.utils import data
import numpy as np
import torch.autograd as ag
import torch.nn as nn
torch.manual_seed(1234)
epochs = 400
n_data = 200
### Input Data
def vector_field(x, y, a=0.01):
    v1 = torch.exp(-a*x*y)*(a*x*torch.sin(x*y) - x*torch.cos(x*y))
    v2 = torch.exp(-a*x*y)*(y*torch.cos(x*y) - a*y*torch.sin(x*y))
    return (v1, v2)

## ----- set up models----- ##
# set network size
n_in = 2
n_h1 = 100
n_h2 = 50
n_o = 1

# two outputs for the unconstrained network
n_o_uc = 2
```

2. Model class for divergence free system

```
# define model class
class DivFree2D(torch.nn.Module):
    def __init__(self, base_net):
        super(DivFree2D, self).__init__()
        self.base_net = base_net

    def forward(self, x):
        x.requires_grad = True
        y = self.base_net(x)
        dydx = ag.grad(outputs=y, inputs=x, create_graph=True, grad_outputs=torch.ones(y.size()),
                        retain_graph=True, only_inputs=True)[0]
        return y, dydx[:,1].unsqueeze(1), -dydx[:,0].unsqueeze(1)

model = DivFree2D(nn.Sequential(nn.Linear(n_in,n_h1),nn.Tanh(),nn.Linear(n_h1,n_h2),
                                nn.Tanh(),nn.Linear(n_h2,n_o)))
```

3. Unconstrained Model for Vanilla PINN

```
model_uc = torch.nn.Sequential(
    torch.nn.Linear(n_in, n_h1),
    torch.nn.Tanh(),
    torch.nn.Linear(n_h1, n_h2),
    torch.nn.Tanh(),
    torch.nn.Linear(n_h2, n_o_uc),
)
```

Linearly Constrained Neural Networks: Code

4. Training and Testing data Generation

```
# pregenerate validation data
x_val = 4.0 * torch.rand(2000, 2)
x1_val = x_val[:, 0].unsqueeze(1)
x2_val = x_val[:, 1].unsqueeze(1)

(v1, v2) = vector_field(x1_val, x2_val)
y1_val = v1 + 0.1 * torch.randn(x1_val.size())
y2_val = v2 + 0.1 * torch.randn(x1_val.size())
y_val = torch.cat((y1_val, y2_val), 1)

# Get the true function values on a grid
xv, yv = torch.meshgrid([torch.arange(0.0, 20.0) * 4.0 / 20.0, torch.arange(0.0, 20.0) * 4.0 / 20.0])
(v1, v2) = vector_field(xv, yv)

# generate training data
x_train = 4.0 * torch.rand(n_data, 2)
x1_train = x_train[:, 0].unsqueeze(1)
x2_train = x_train[:, 1].unsqueeze(1)

(v1_t, v2_t) = vector_field(x1_train, x2_train)
y1_train = v1_t + 0.1 * torch.randn(x1_train.size())
y2_train = v2_t + 0.1 * torch.randn(x1_train.size())
```

Linearly Constrained Neural Networks: Code

5. Data Class for Automatic Batching

```
# define Dataset class
class Dataset(data.Dataset):
    'Characterizes a dataset for PyTorch'

    def __init__(self, x1, x2, y1, y2):
        'Initialization'
        self.x1 = x1
        self.x2 = x2
        self.y1 = y1
        self.y2 = y2

    def __len__(self):
        'Denotes the total number of samples'
        return len(self.x1)

    def __getitem__(self, index):
        'Generates one sample of data'
        # Select sample
        x1 = self.x1[index]
        x2 = self.x2[index]
        y1 = self.y1[index]
        y2 = self.y2[index]

        return x1, x2, y1, y2

training_set = Dataset(x1_train, x2_train, y1_train, y2_train)

# data loader Parameters
DL_params = {'batch_size': 100,
             'shuffle': True,
             'num_workers': 0,
             'pin_memory': False}
training_generator = data.DataLoader(training_set, **DL_params)
```

6. Loss and LR scheduler

```
# ----- Set up and train the constrained model -----
criterion = torch.nn.MSELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.01)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer, patience=10,
    min_lr=1e-10,
    factor=0.5, cooldown=15)
```

7. Train and Eval routine for Constrained Model

```
def train(epoch):
    model.train()
    total_loss = 0
    n_batches = 0
    for x1_train, x2_train, y1_train, y2_train in training_generator:
        optimizer.zero_grad()
        x_train = torch.cat((x1_train, x2_train), 1)
        (yhat, v1hat, v2hat) = model(x_train)
        loss = (criterion(y1_train, v1hat) + criterion(y2_train, v2hat)) / 2
        loss.backward()
        optimizer.step()
        total_loss += loss
        n_batches += 1
    return total_loss / n_batches

def eval(epoch):
    model.eval()
    # with torch.no_grad():
    (yhat, v1hat, v2hat) = model(x_val)
    loss = (criterion(y1_val, v1hat) + criterion(y2_val, v2hat)) / 2
    return loss.cpu()
```

Linearly Constrained Neural Networks: Code

8. Training of Constrained Model

```
train_loss = np.empty([epochs, 1])
val_loss = np.empty([epochs, 1])

print('Training Constrained NN')

for epoch in range(epochs):
    train_loss[epoch] = train(epoch).detach().numpy()
    v_loss = eval(epoch)
    scheduler.step(v_loss)
    val_loss[epoch] = v_loss.detach().numpy()
    print('Constrained NN: epoch: ', epoch, 'training loss ', train_loss[epoch], \
          'validation loss', val_loss[epoch])

# work out the rms error for this one
x_pred = torch.cat((xv.reshape(20 * 20, 1), yv.reshape(20 * 20, 1)), 1)
(f_pred, v1_pred, v2_pred) = model(x_pred)
error_new = torch.cat((v1.reshape(400, 1) - v1_pred.detach(), v2.reshape(400, 1) - v2_pred.detach()), 0)
rms_error = torch.sqrt(sum(error_new * error_new) / 800)

# ----- Set up and train the unconstrained model -----
optimizer_uc = torch.optim.Adam(model_uc.parameters(), lr=0.01)
scheduler_uc = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer_uc, patience=10,
                                                       min_lr=1e-10,
                                                       factor=0.5,
                                                       cooldown=15)
```

9. Training of Unconstrained Model

```
def train_uc(epoch):
    model_uc.train()
    total_loss = 0
    n_batches = 0
    for x1_train, x2_train, y1_train, y2_train in training_generator:
        optimizer_uc.zero_grad()
        x_train = torch.cat((x1_train, x2_train), 1)
        vhat = model_uc(x_train)
        y_train = torch.cat((y1_train, y2_train), 1)
        loss = criterion(y_train, vhat)
        loss.backward()
        optimizer_uc.step()
        total_loss += loss.cpu()
        n_batches += 1
    return total_loss / n_batches

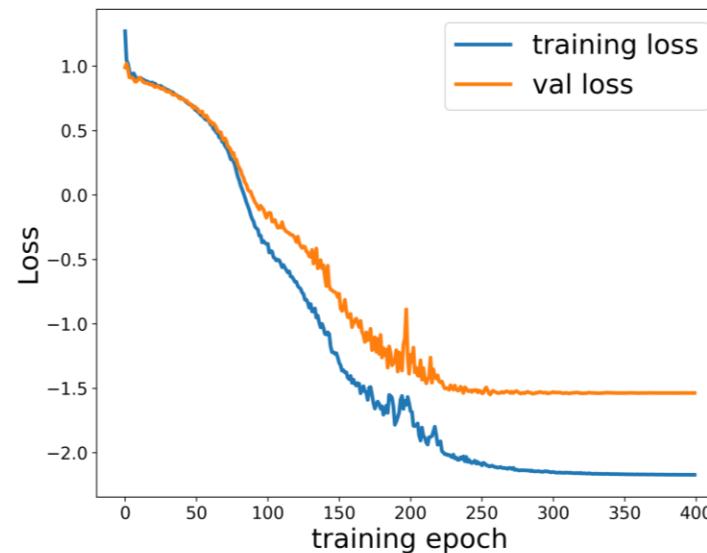
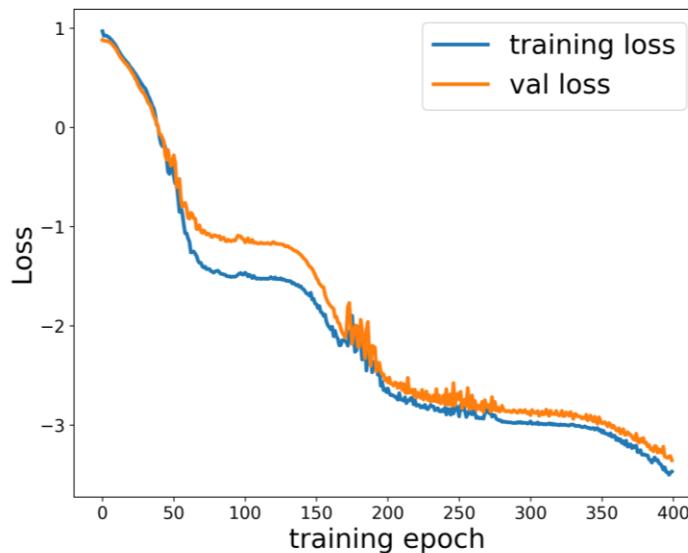
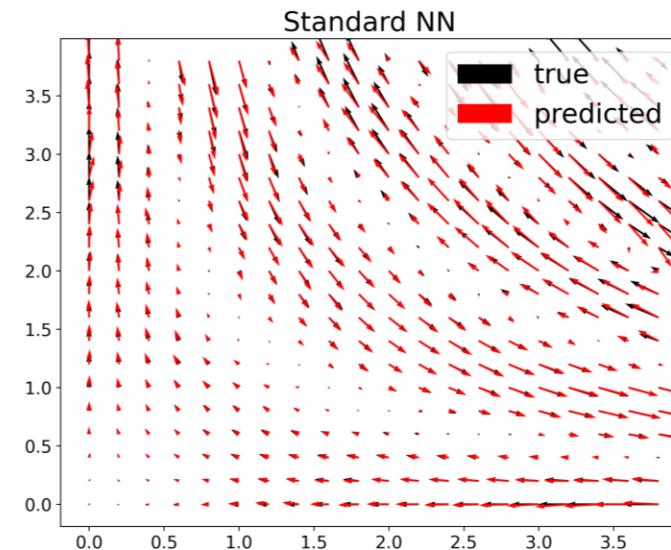
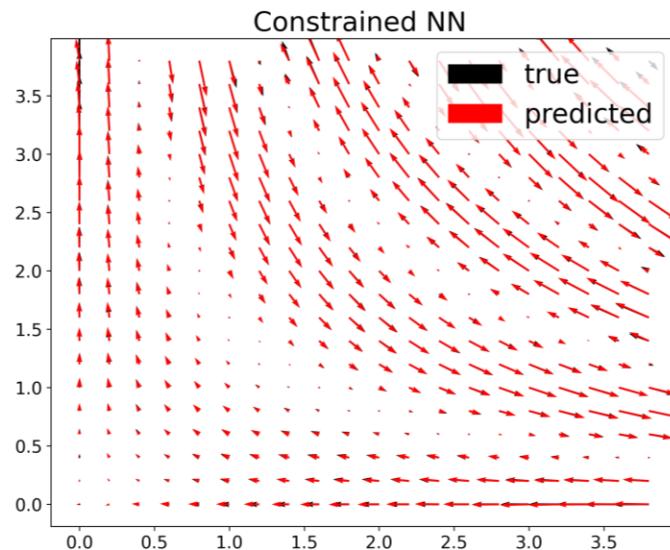
def eval_uc(epoch):
    model_uc.eval()
    with torch.no_grad():
        (vhat) = model_uc(x_val)
        loss = criterion(y_val, vhat)
    return loss.cpu()

train_loss_uc = np.empty([epochs, 1])
val_loss_uc = np.empty([epochs, 1])

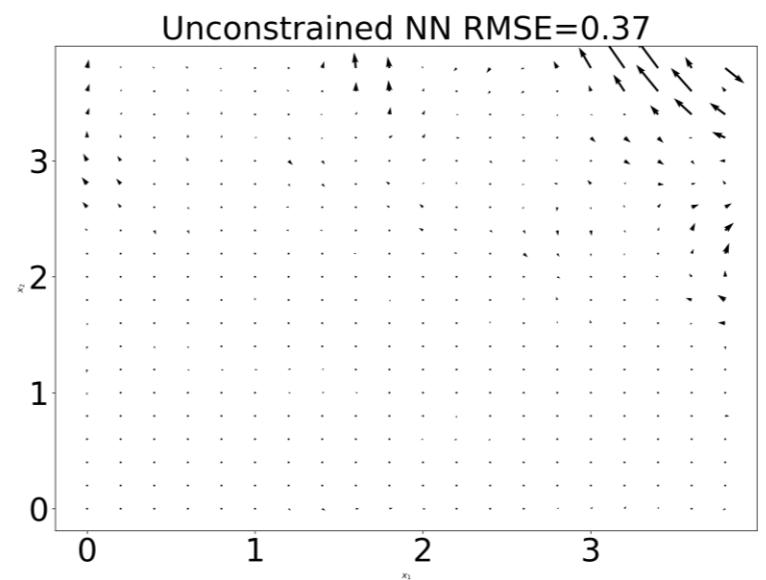
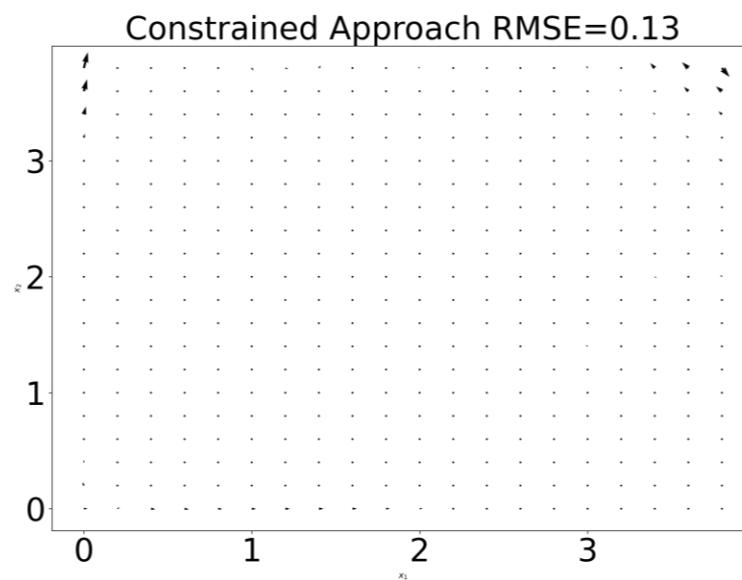
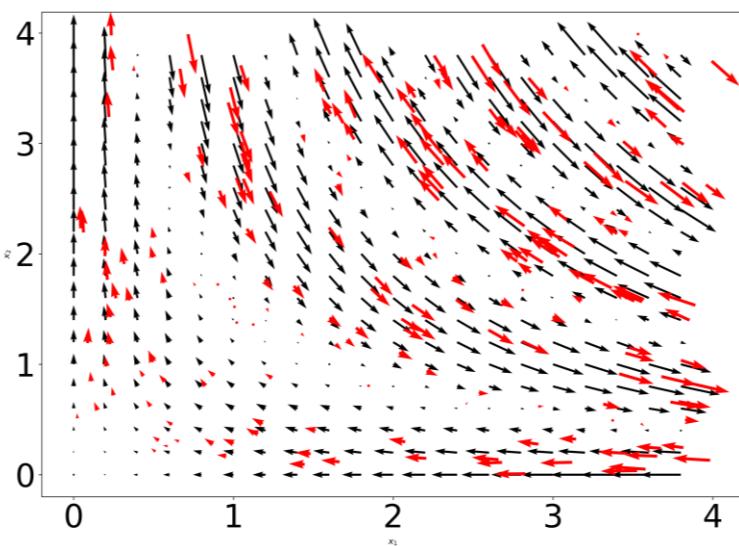
print('Training standard NN')

for epoch in range(epochs):
    train_loss_uc[epoch] = train_uc(epoch).detach().numpy()
    v_loss = eval_uc(epoch)
    scheduler_uc.step(v_loss)
    val_loss_uc[epoch] = v_loss.detach().numpy()
    print('Standard NN: epoch: ', epoch, 'training loss ', \
          train_loss_uc[epoch], 'validation loss', val_loss_uc[epoch])
```

Linearly Constrained Neural Networks: Results



Linearly Constrained Neural Networks: Results



PINNs Family: Weighted Residual Methods

- Approximation: $u \approx \tilde{u} = u_{NN}$

- Strong-Form Residuals:

$$r(\tilde{u}) = \mathcal{L}^q \tilde{u} - f \quad : \text{PDE residual}$$

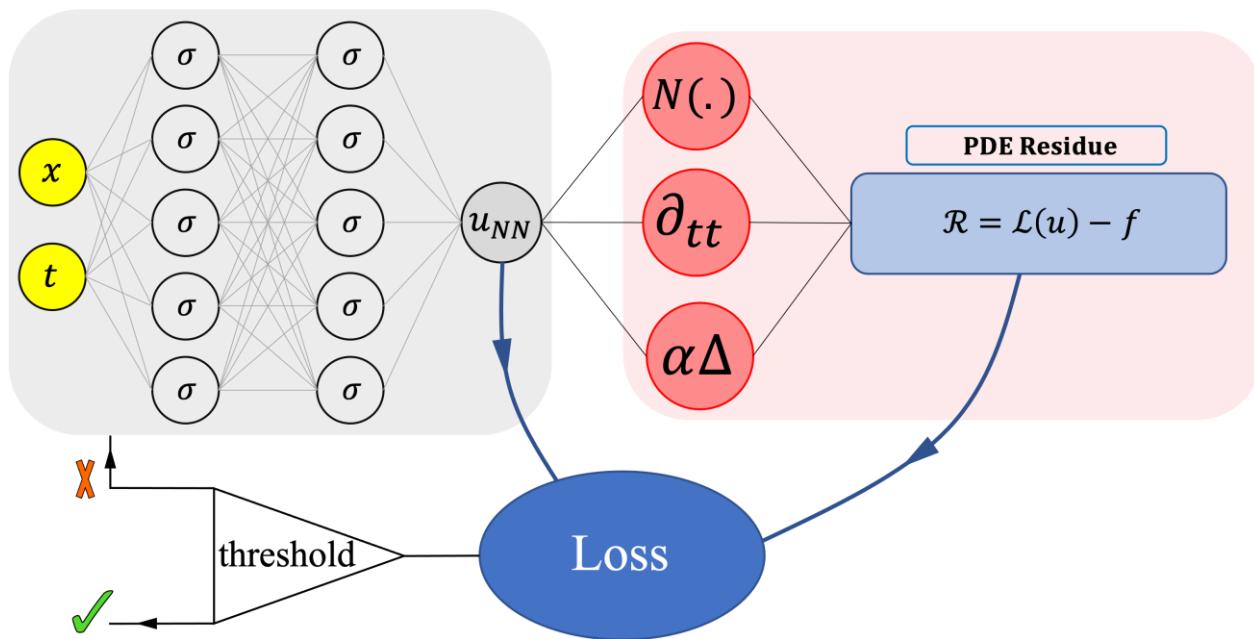
$$r_b(\tilde{u}) = \tilde{u} - h \quad : \text{boundary residual}$$

$$r_0(\tilde{u}) = \tilde{u} - g \quad : \text{initial residual}$$

Loss Function:

$$L^s(\tilde{u}) = \frac{1}{N_r} \sum_{i=1}^{N_r} |r(\mathbf{x}_r^i, t_r^i)|^2 + \tau_b \frac{1}{N_b} \sum_{i=1}^{N_b} |r_b(\mathbf{x}_b^i, t_b^i)|^2 + \tau_0 \frac{1}{N_0} \sum_{i=1}^{N_0} |r_0(\mathbf{x}_0^i)|^2$$

minimization problem: $\min L^s(\tilde{u})$



PINNs Family: Weighted Residual Methods

- Weighted/Variational Residuals:
projections onto space of test function V

$$\mathcal{R}_j(\tilde{u}) = \int_{\Omega \times (0, T]} r(\tilde{u}) v_j dx dt$$

$$\mathcal{R}_{b,j}(\tilde{u}) = \int_{\partial\Omega \times (0, T]} r_b(\tilde{u}) v_j dx dt$$

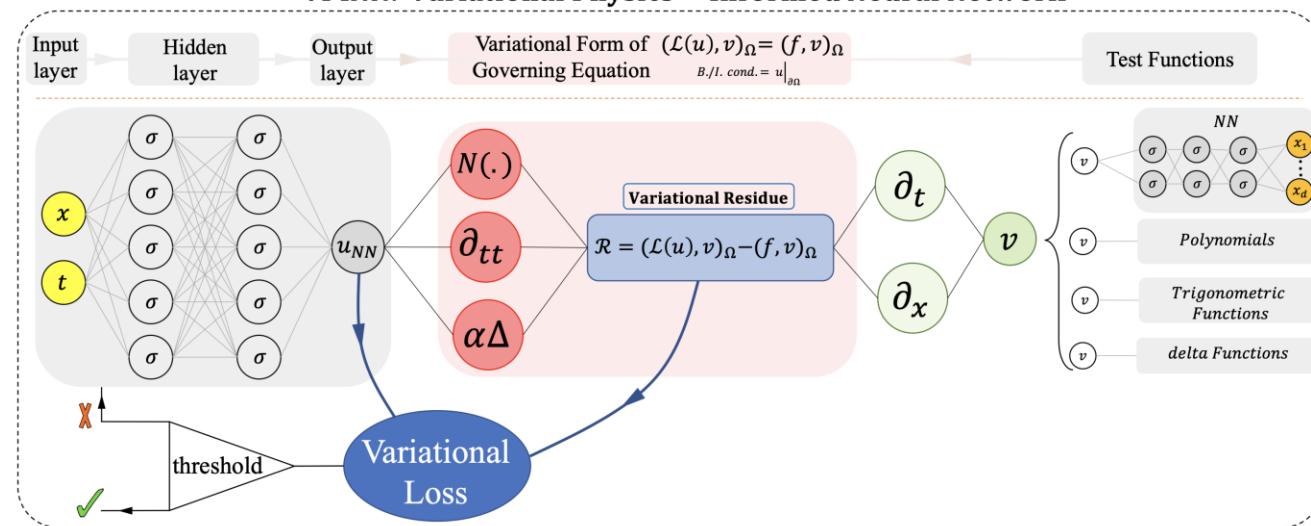
$$\mathcal{R}_{0,j}(\tilde{u}) = \int_{\Omega} r_0(\tilde{u}) v_j dx$$

Loss Function:

$$L^v(\tilde{u}, v) = w \sum_{j=1}^{N_r} \mathcal{R}_j^2(\tilde{u}) + w_b \sum_{j=1}^{N_b} \mathcal{R}_{b,j}^2(\tilde{u}) + w_0 \sum_{j=1}^{N_0} \mathcal{R}_{0,j}^2(\tilde{u})$$

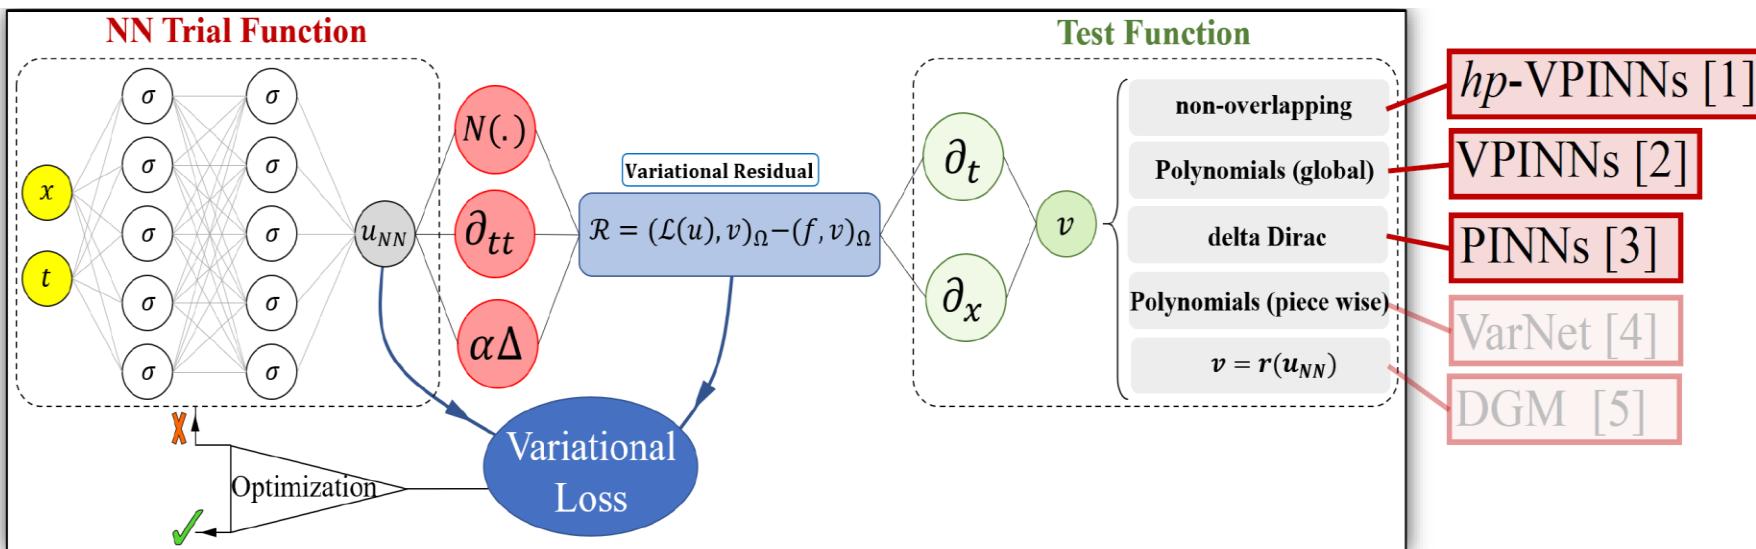
minimization problem: $\min L^v(\tilde{u}, v)$

VPINN: Variational Physics – Informed Neural Network



Kharazmi E, Zhang Z, Karniadakis GE. Variational physics-informed neural networks for solving partial differential equations. arXiv preprint arXiv:1912.00873. 2019 Nov 27.

hp-VPINNs: A General Framework for Solving PDEs



	Local NN	Global NN
Local v	Conservative VPINNs [6]	<i>hp</i> -VPINNs [5]
Global v	-	VPINNs[3], VarNet[4], D3M [7]

Global:= single domain
Local:= domain decomposition

[1] Kharazmi et al., CMAME (2020).

[5] Sirignano et al., JCP (2018)

[2] Kharazmi et al., *arXiv* (2019).

[6] Jagtap et al., CMAME (2020)

[3] Raissi et al., JCP (2019)

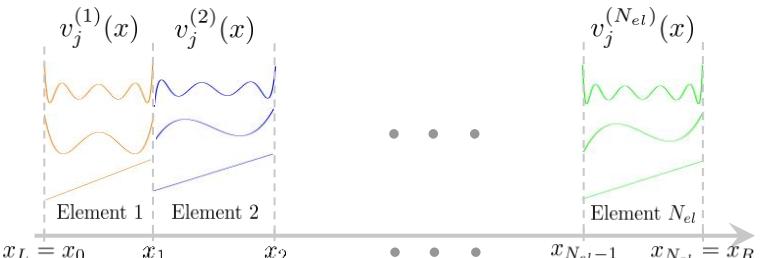
[7] Li et al., arXiv (2019)

[4] Khodayi-Mehr et al., arXiv (2019).

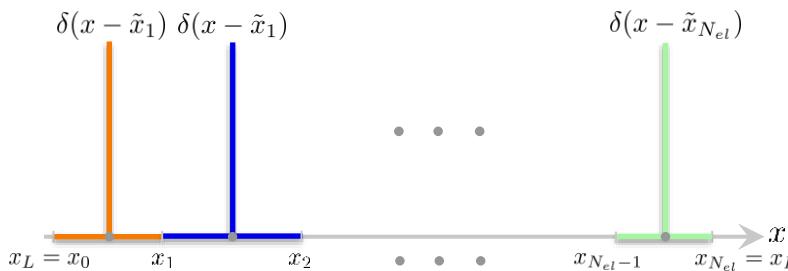
hp-VPINNs: A General Framework for Solving PDEs

h-refinement & *p*-refinement

$$\Omega = \bigcup_e \Omega^{(e)}$$
$$v_j^{(e)} = \begin{cases} \bar{v}_j \neq 0, & x \in \Omega^{(e)} \\ 0, & \text{otherwise} \end{cases}$$

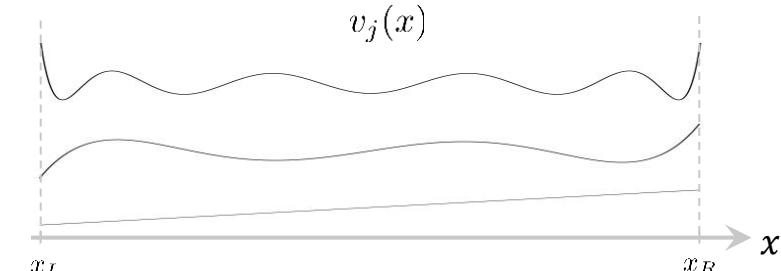


PINNs



Delta Dirac test functions

hp-VPINNs



p-refinement

VPINNs

- nonlinear approximation via neural networks
- *h*-refinement via domain decomposition
- *p*-refinement via projection onto high-order polynomial space

hp-VPINNs: A General Framework for Solving PDEs

Loss Function

elemental/local variational loss

hp-VPINNs:
$$L^v = \sum_{e=1}^{N_{el}} \frac{1}{K^{(e)}} \sum_{k=1}^{K^{(e)}} \left| \mathcal{R}_k^{(e)} \right|^2 + \tau_b \frac{1}{N_b} \sum_{i=1}^{N_b} \left| r_b(\mathbf{x}_b^i, t_b^i) \right|^2 + \tau_0 \frac{1}{N_0} \sum_{i=1}^{N_0} \left| r_0(\mathbf{x}_0^i) \right|^2$$

global variational loss

VPINNs:
$$L^v = \frac{1}{K} \sum_{j=1}^K \left| \mathcal{R}_j \right|^2 + \tau_b \frac{1}{N_b} \sum_{i=1}^{N_b} \left| r_b(\mathbf{x}_b^i, t_b^i) \right|^2 + \tau_0 \frac{1}{N_0} \sum_{i=1}^{N_0} \left| r_0(\mathbf{x}_0^i) \right|^2$$

strong-form loss

PINNs:
$$L^s = \frac{1}{N_r} \sum_{i=1}^{N_r} \left| r(\mathbf{x}_r^i, t_r^i) \right|^2 + \tau_b \frac{1}{N_b} \sum_{i=1}^{N_b} \left| r_b(\mathbf{x}_b^i, t_b^i) \right|^2 + \tau_0 \frac{1}{N_0} \sum_{i=1}^{N_0} \left| r_0(\mathbf{x}_0^i) \right|^2$$

Variational Neural Networks (VNNs): Function Approximation

Variational residual:

$$\mathcal{R}_k^{(e)} = \int_{\Omega_e} (u_{NN}(\boldsymbol{x}) - u(\boldsymbol{x})) v_k^{(e)}(\boldsymbol{x}) d\Omega_e$$

Variational loss:

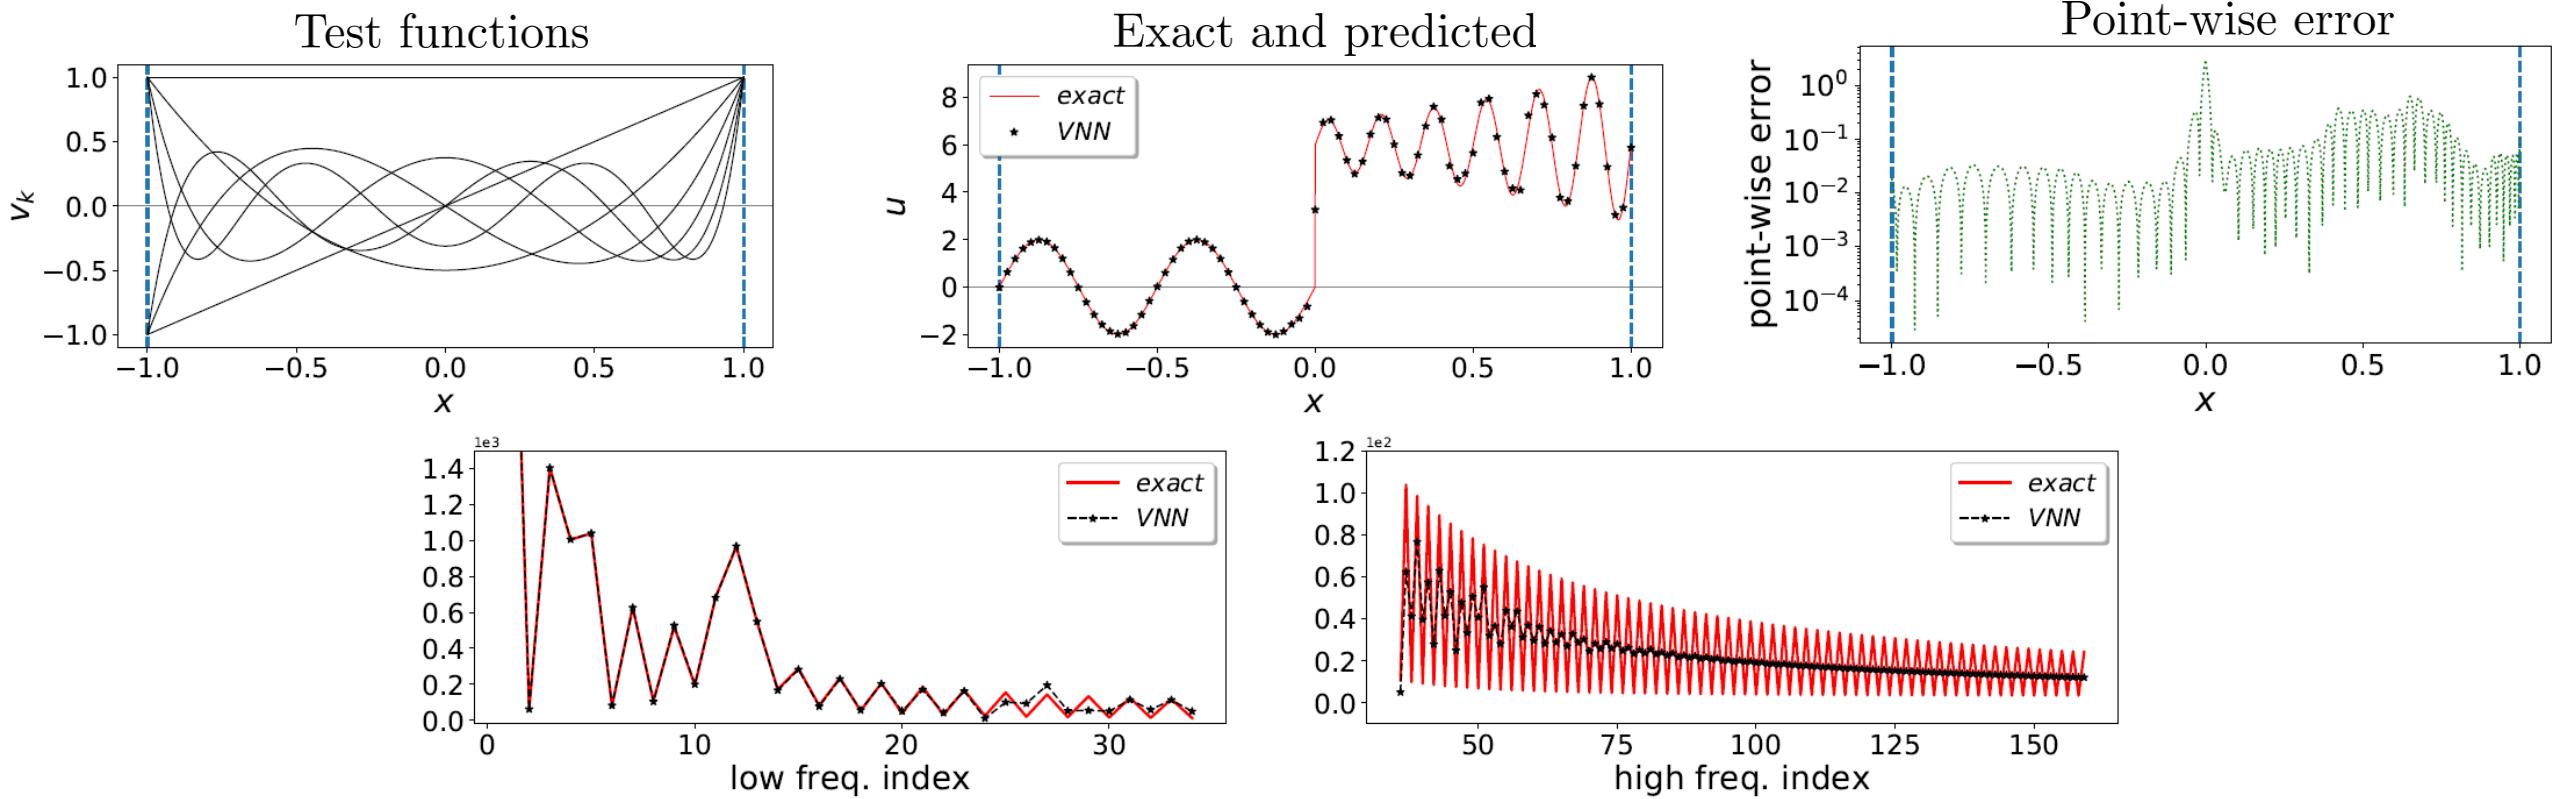
$$L^{\mathfrak{v}} = \sum_{e=1}^{N_{el}} \frac{1}{K^{(e)}} \sum_{k=1}^{K^{(e)}} \left| \mathcal{R}_k^{(e)} \right|^2$$

Example (discontinuous function approximation): the target function is

$$u^{exact} = \begin{cases} 2 \sin(4\pi x) & x \in [-1, 0], \\ 6 + \exp^{1.2x} \sin(12\pi x) & x \in (0, 1] \end{cases}$$

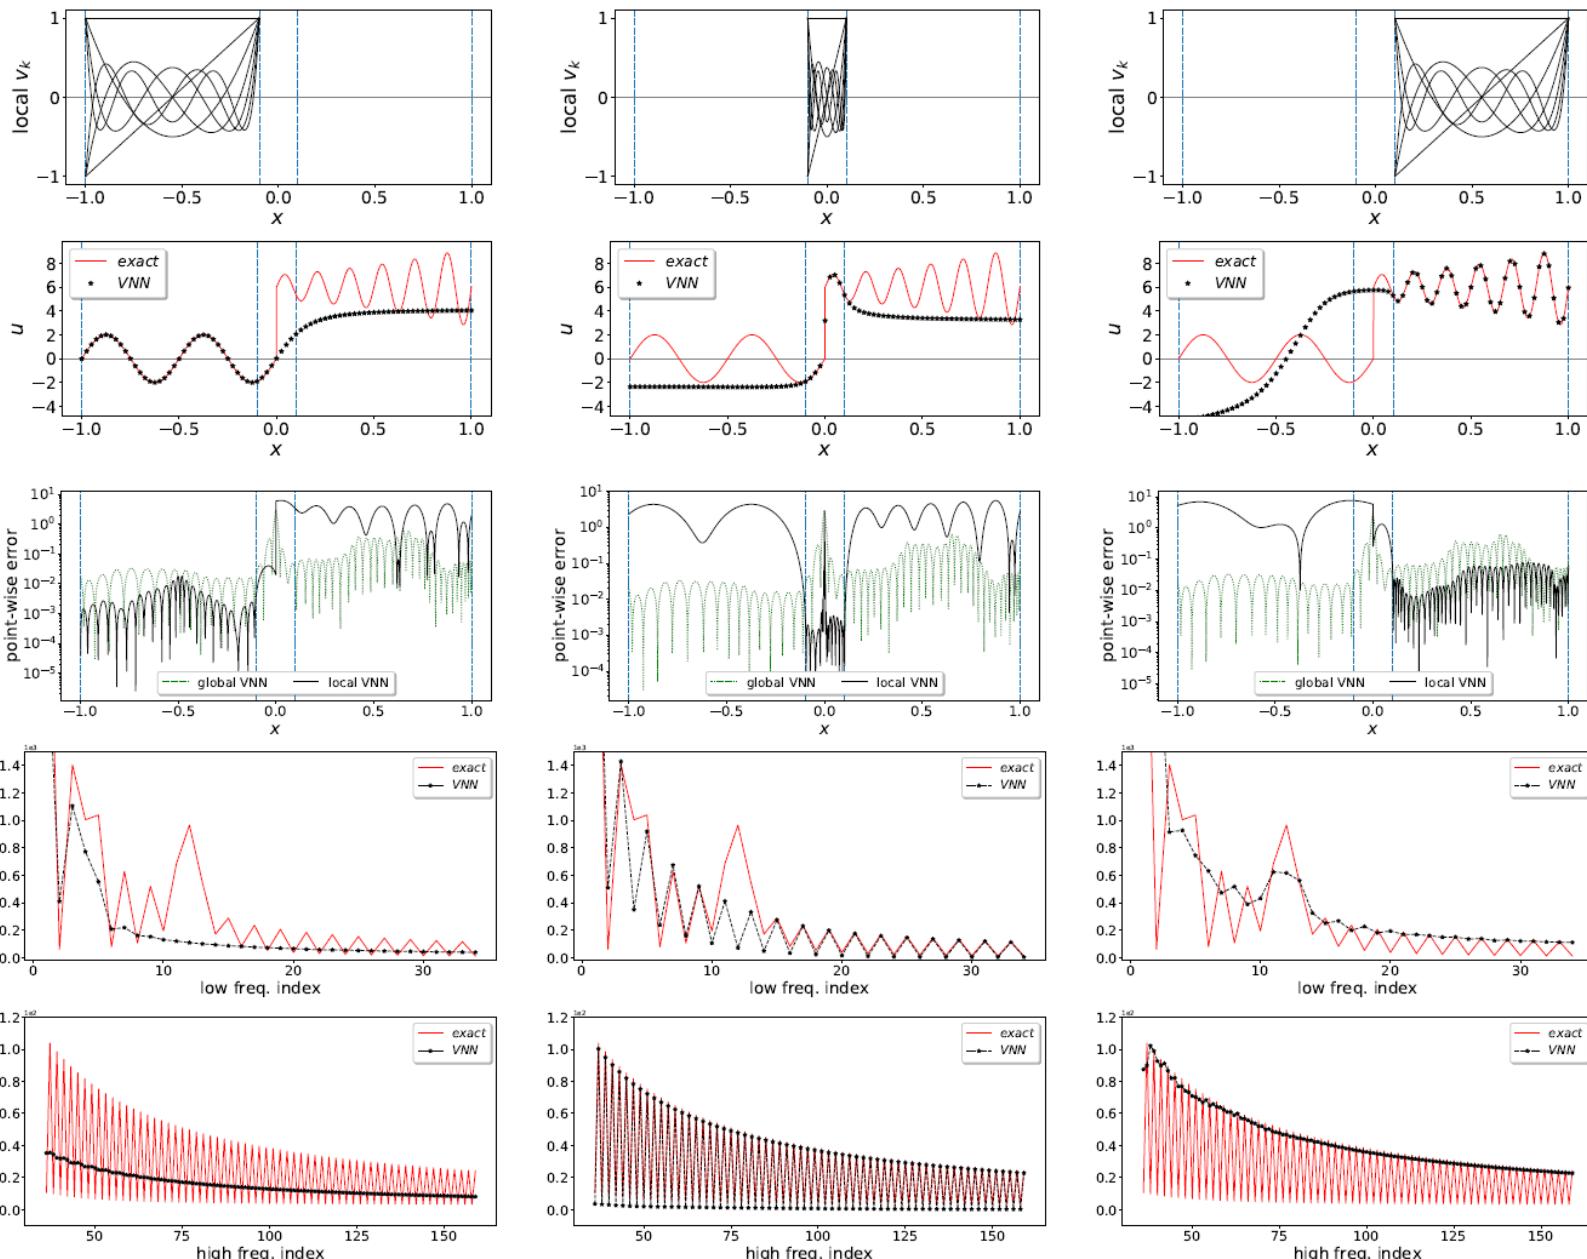
Variational Neural Networks (VNNs): Function Approximation

Global test functions (single element):



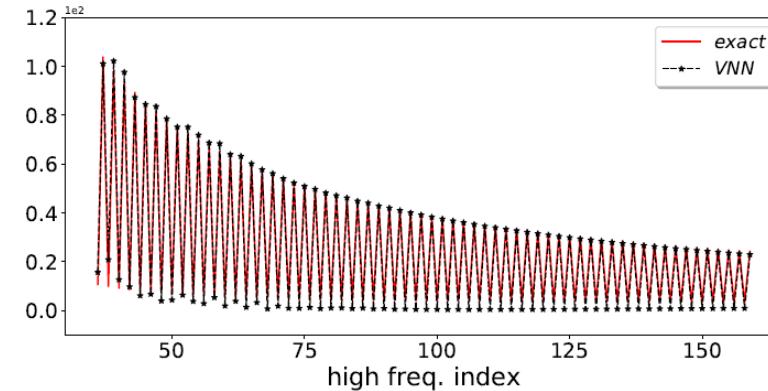
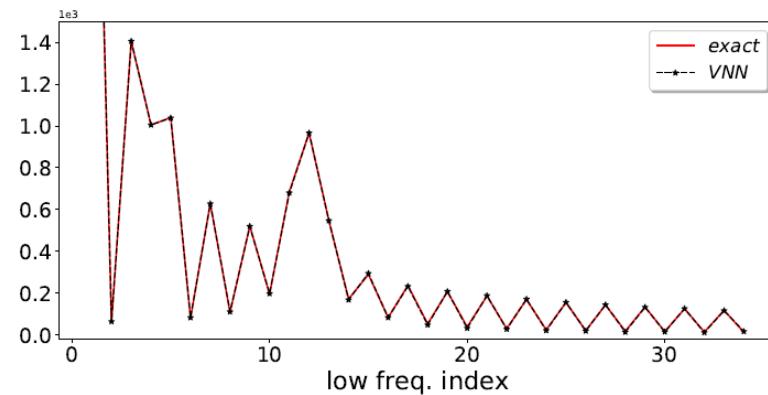
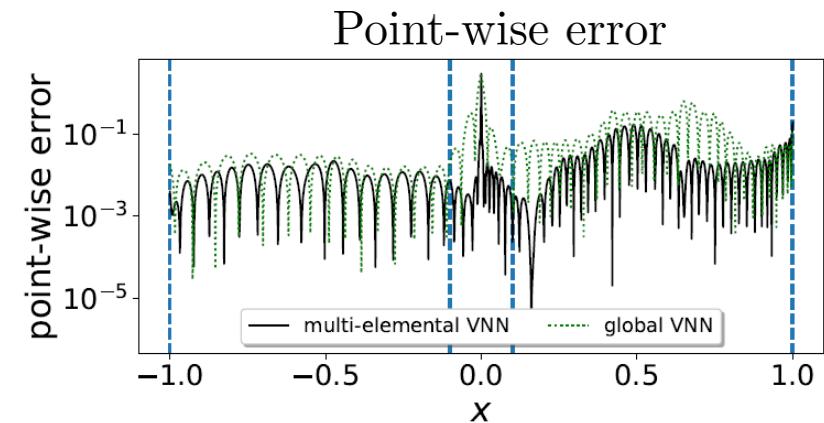
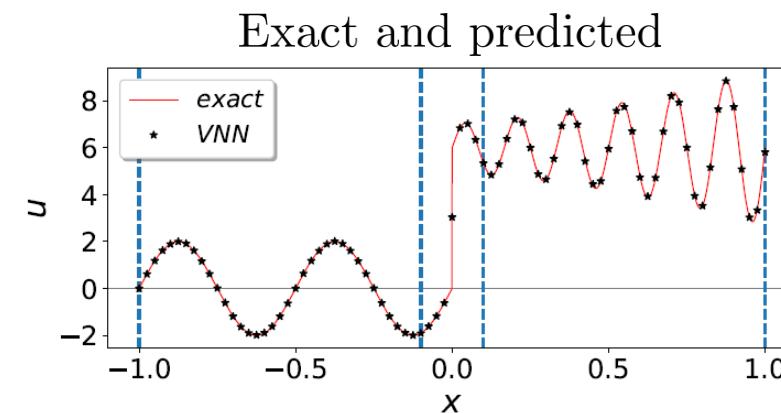
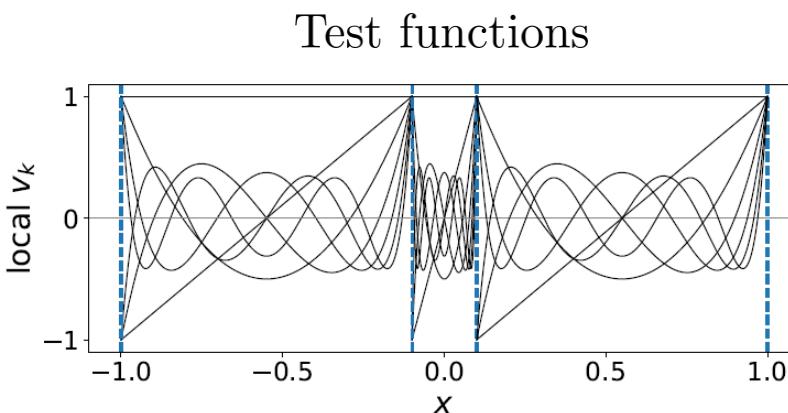
Variational Neural Networks (VNNs): Function Approximation

Local (elemental)
test functions:



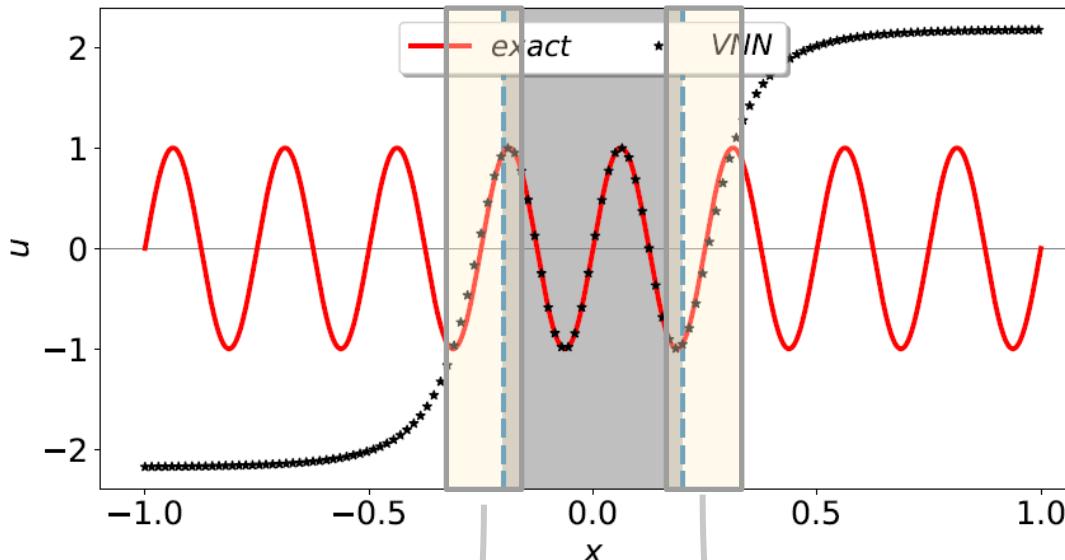
Variational Neural Networks (VNNs): Function Approximation

Multi-elemental test functions:

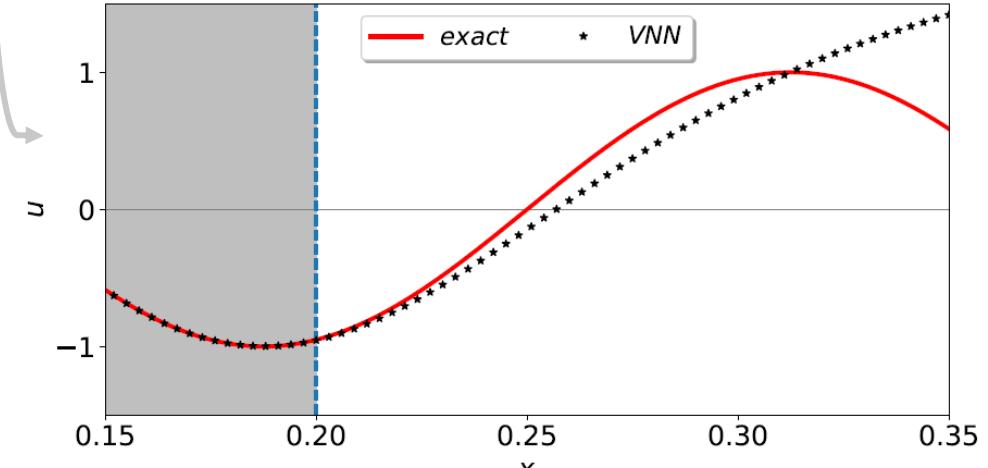
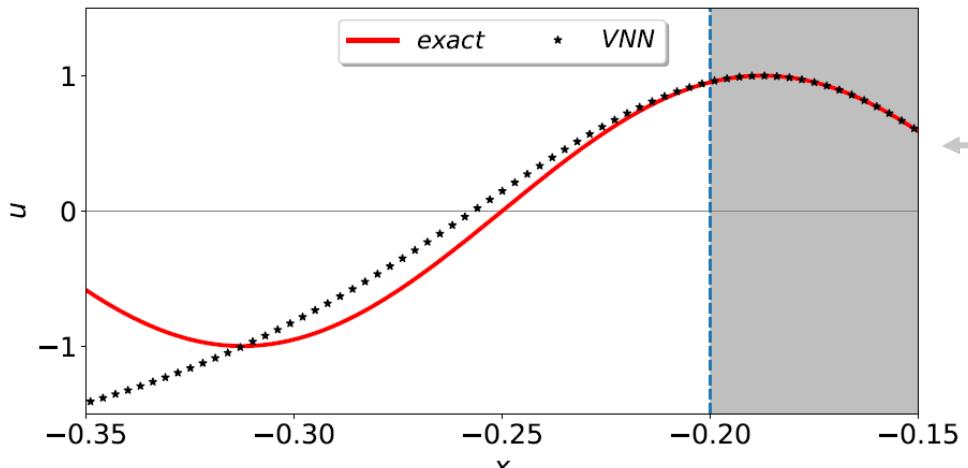


Variational Neural Networks (VNNs): Function Approximation

Localized learning & learning out-of-the-box



- The test functions have **non-zero** values only in the gray box ($-0.2 < x < 0.2$).
- The test functions have **zero** values in the white region.
- It is expected that VNN only learns the function in the grey box.
- But VNN learns the function slightly out side of the box, a notion of generalization to unseen domain. This is indicated as the yellow region.



hp-VPINNs: Poisson's Equation (1D)

- Governing Equation:

$$-\frac{d^2 u(x)}{dx^2} = f(x)$$

- Boundary conditions:

$$u(-1) = g$$

$$u(1) = h$$

- Variational residual:

$$\mathcal{R}_k = \sum_{e=1}^{N_{el}} \mathcal{R}_k^{(e)} = \sum_{e=1}^{N_{el}} \int_{x_{e-1}}^{x_e} \left(-\frac{d^2 u_{NN}(x)}{dx^2} - f(x) \right) v_k^{(e)}(x) dx,$$

hp-VPINNs: Poisson's Equation (1D)

- Reducing the regularity and simplifying the network by integration-by-parts:

$$^{(1)}\mathcal{R}_k^{(e)} = - \int_{x_{e-1}}^{x_e} \frac{d^2 u_{NN}(x)}{dx^2} v_k^{(e)}(x) dx - F_k^{(e)},$$

$$^{(2)}\mathcal{R}_k^{(e)} = \int_{x_{e-1}}^{x_e} \frac{du_{NN}(x)}{dx} \frac{dv_k^{(e)}(x)}{dx} dx - \frac{du_{NN}(x)}{dx} v_k^{(e)}(x) \Big|_{x_{e-1}}^{x_e} - F_k^{(e)},$$

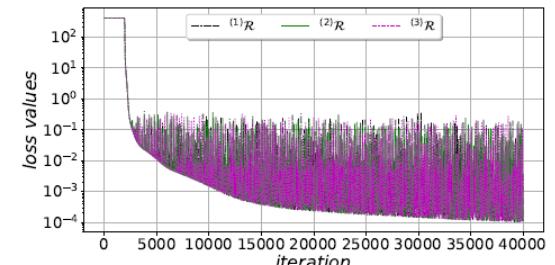
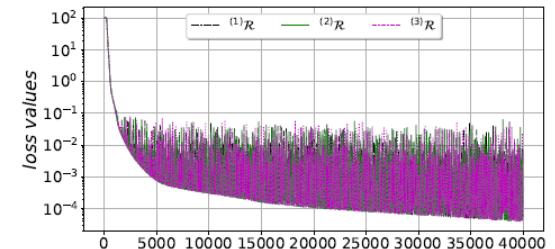
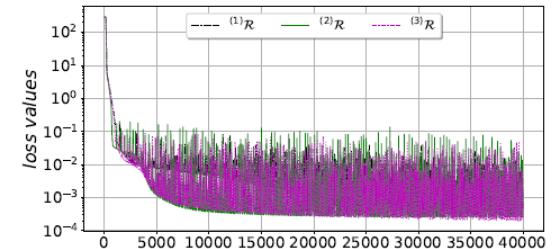
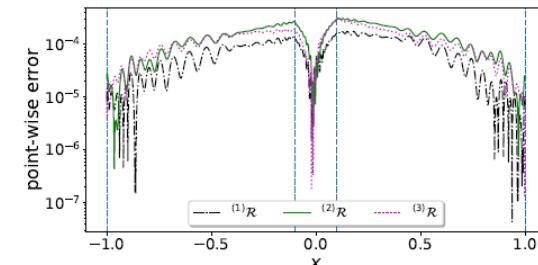
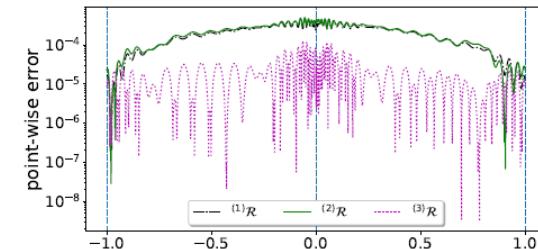
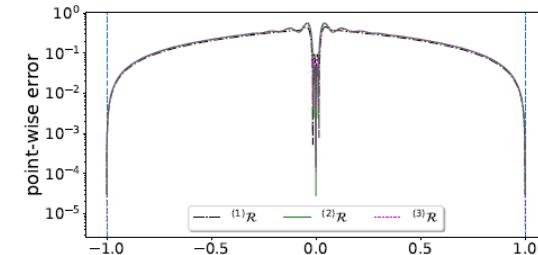
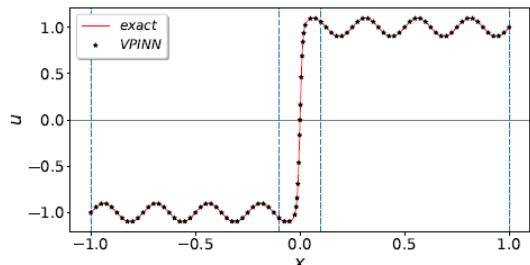
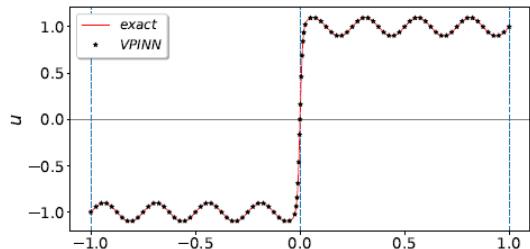
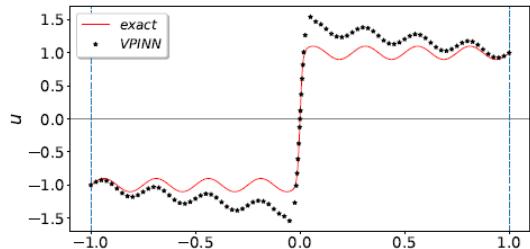
$$^{(3)}\mathcal{R}_k^{(e)} = - \int_{x_{e-1}}^{x_e} u_{NN}(x) \frac{d^2 v_k^{(e)}(x)}{dx^2} dx - \frac{du_{NN}(x)}{dx} v_k^{(e)}(x) \Big|_{x_{e-1}}^{x_e} + u_{NN}(x) \frac{dv_k^{(e)}(x)}{dx} \Big|_{x_{e-1}}^{x_e} - F_k^{(e)}$$

- Variational loss:

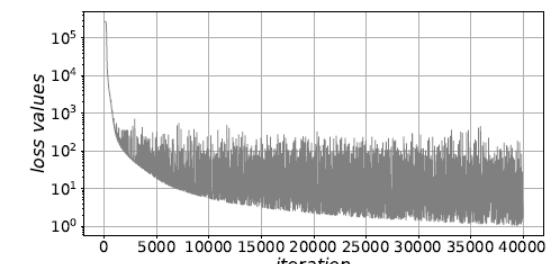
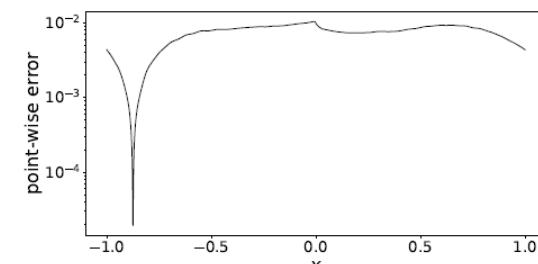
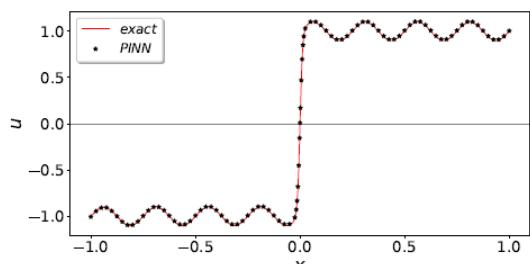
$$L^{\mathfrak{v}(i)} = \sum_{e=1}^{N_{el}} \frac{1}{K^{(e)}} \sum_{k=1}^{K^{(e)}} \left| {}^{(i)}\mathcal{R}_k^{(e)} \right|^2 + \frac{\tau_b}{2} \left(|u_{NN}(-1) - g|^2 + |u_{NN}(1) - h|^2 \right), \quad i = 1, 2, 3$$

hp-VPINNs: Poisson's Equation (1D)

VPINN



PINN



Example: Steep Solution

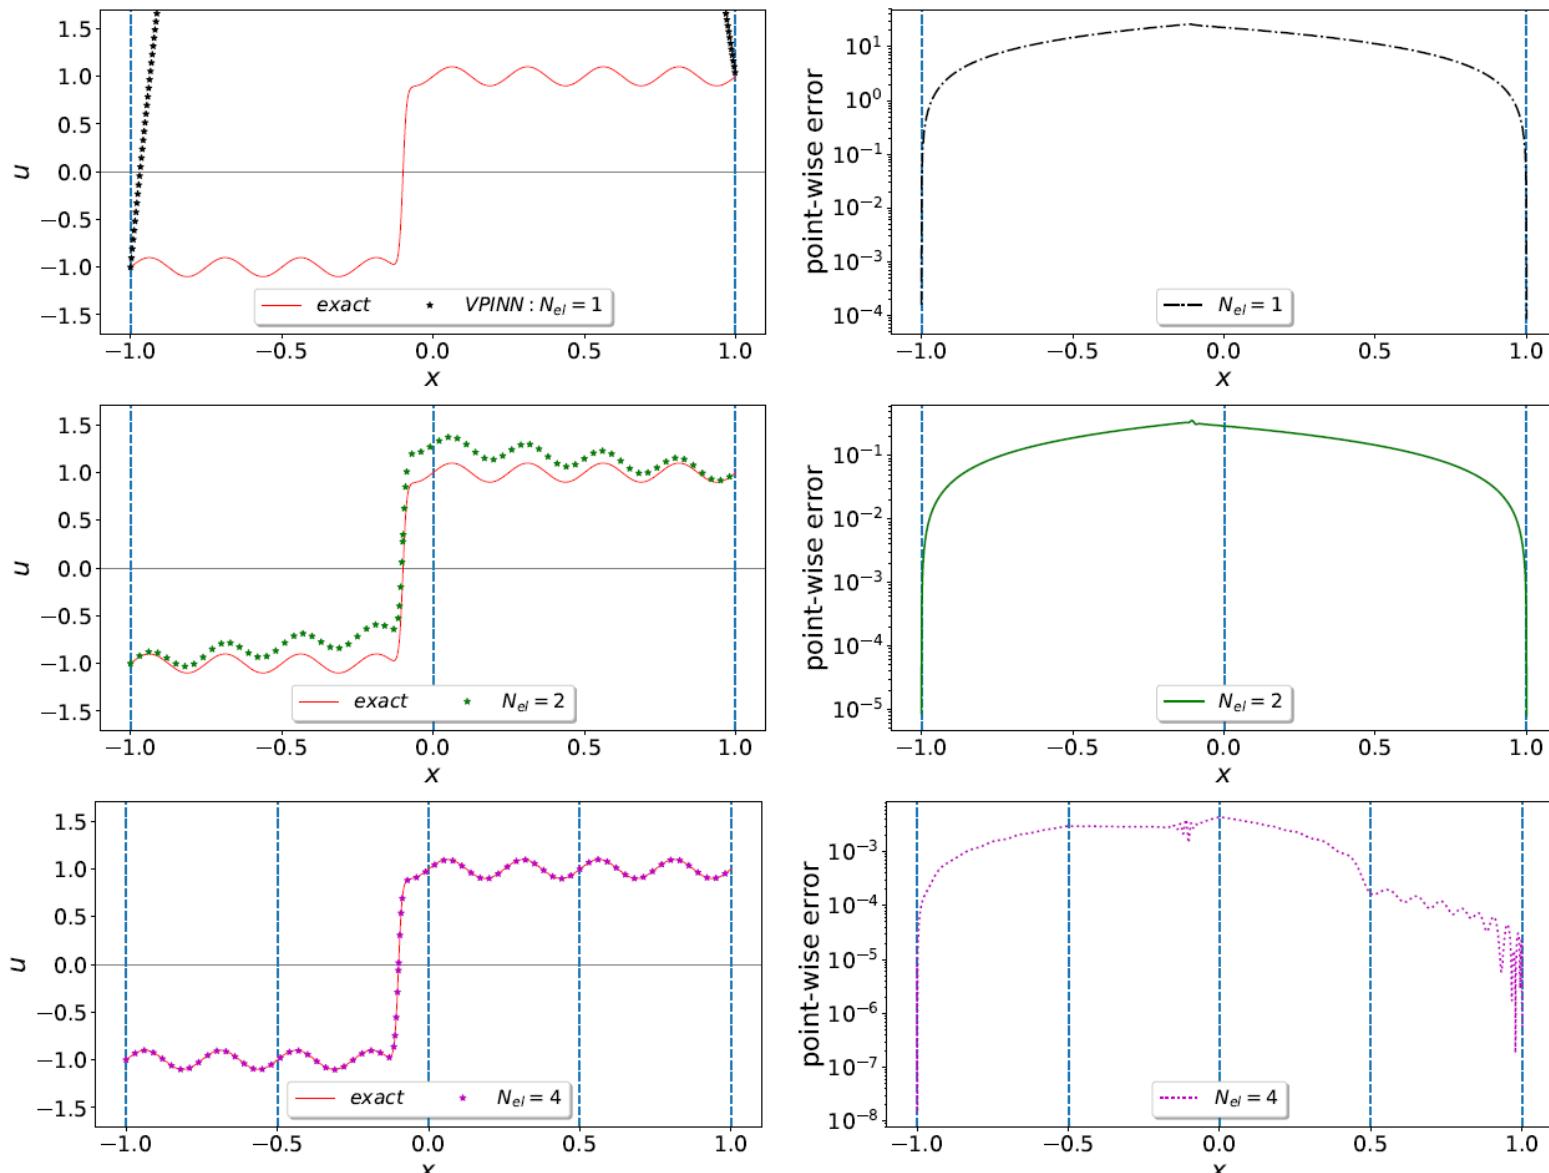
$$u^{exact}(x) = 0.1 \sin(8\pi x) + \tanh(80x)$$

hp-VPINNs: Poisson's Equation (1-D)

Example: Singularity Capturing

$$u^{\text{exact}}(x) = 0.1 \sin(8\pi x) + \tanh(80(x + 0.1))$$

asymmetric steep part



Kharazmi E, Zhang Z, Karniadakis GE. *hp*-VPINNs: Variational physics-informed neural networks with domain decomposition. Computer Methods in Applied Mechanics and Engineering. 2021 Feb 1;374:113547.

***hp*-VPINNs: Poisson's Equation (2-D)**

$$\nabla^2 u(x, y) = f(x, y), \quad (x, y) \in \Omega = [-1, 1] \times [-1, 1]$$

Discrete finite dimensional space, given by

$$\tilde{V} = \text{span} \{ v_{k_1 k_2}(x, y) = \phi_{k_1}(x)\phi_{k_2}(y), k_m = 1, 2, \dots, K_m, m = 1, 2 \}$$

$$\phi_{k_1}^{(e_x)}(x_{e_x-1}) = \phi_{k_1}^{(e_x)}(x_{e_x}) = 0, \quad k_1 = 1, 2, \dots, K_1,$$

$$\phi_{k_2}^{(e_y)}(y_{e_y-1}) = \phi_{k_2}^{(e_y)}(y_{e_y}) = 0, \quad k_2 = 1, 2, \dots, K_2.$$

The variational residual becomes

$$\mathcal{R}_{k_1 k_2}^{(e_x e_y)} = \int_{x_{e_x-1}}^{x_{e_x}} \int_{y_{e_y-1}}^{y_{e_y}} (\nabla^2 u_{NN}(x, y) - f(x, y)) \phi_{k_1}^{(e_x)}(x) \phi_{k_2}^{(e_y)}(y) dx dy$$

The variational loss

$$L^{v(i)} = \sum_{e_r=1}^{N_{el_r}} \sum_{e_y=1}^{N_{el_y}} \frac{1}{K_1 K_2} \sum_{k=1}^{K_1 K_2} \left| {}^{(i)}\mathcal{R}_k^{(e_x e_y)} \right|^2 + \tau_b \frac{1}{N_b} \sum_{i=1}^{N_b} \left| r_b(x_b^i, y_b^i) \right|^2, \quad i = 1, 2, 3$$

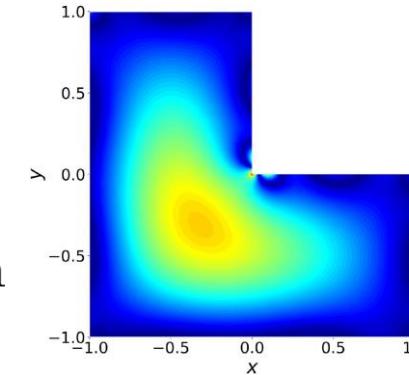
hp-VPINNs: Poisson's Equation (2-D)

2-D PDE: $\begin{cases} \nabla^2 u(x, y) = f(x, y) \\ \text{Homogeneous Dirichlet B.C.} \end{cases}$

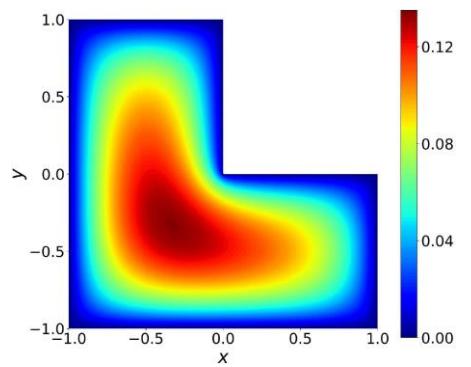
The homogeneous case $f(x, y) = 0$ gives the Laplace equation

hp-VPINN:⁽¹⁾ \mathcal{R} formulation

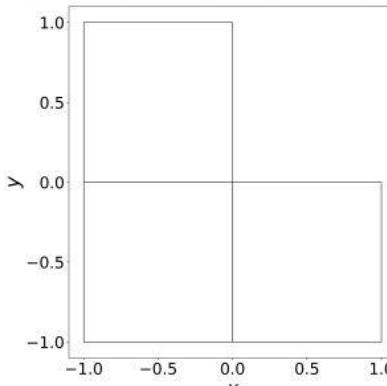
PINNs point-wise error



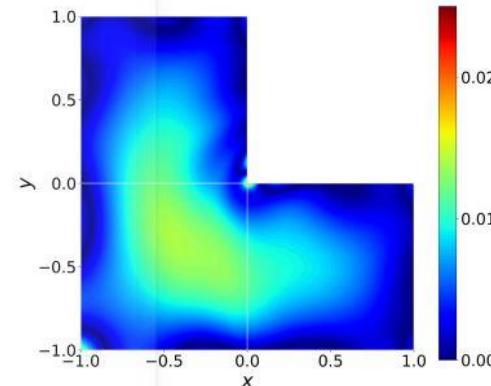
reference solution



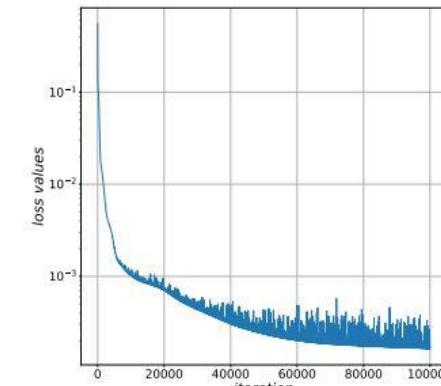
C domain decomposition



D point-wise error



E loss

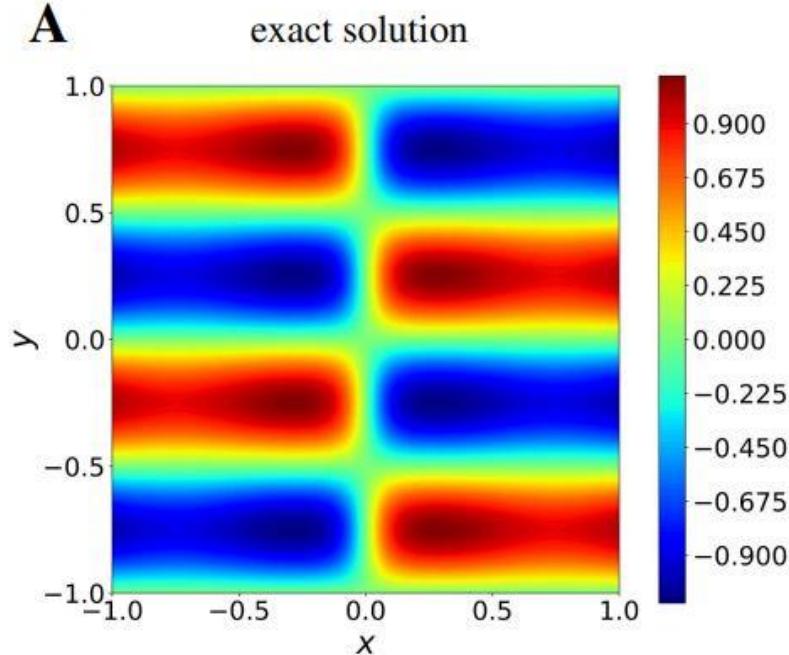


hp-VPINNs: Poisson's Equation

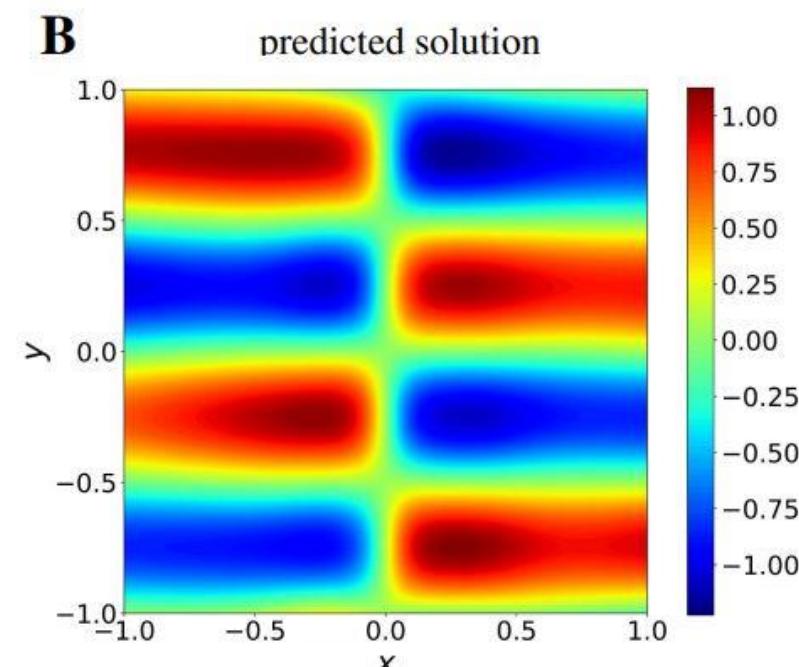
2-D PDE: $\begin{cases} \nabla^2 u(x, y) = f(x, y) \\ \text{Homogeneous Dirichlet B.C.} \end{cases}, \quad u^{\text{exact}}(x, y) = \sin(2\pi y) \times (0.1 \sin(2\pi x) + \tanh(10x))$

PINN

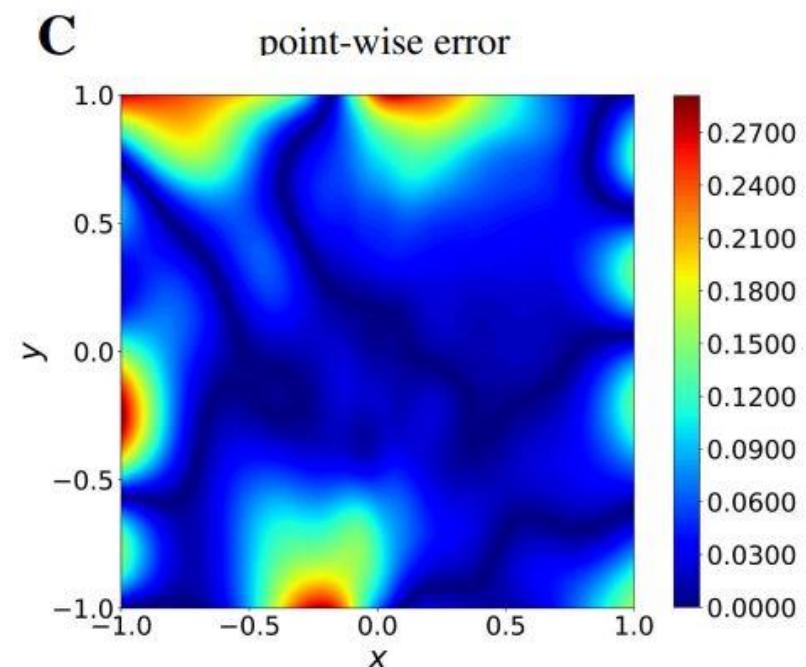
A



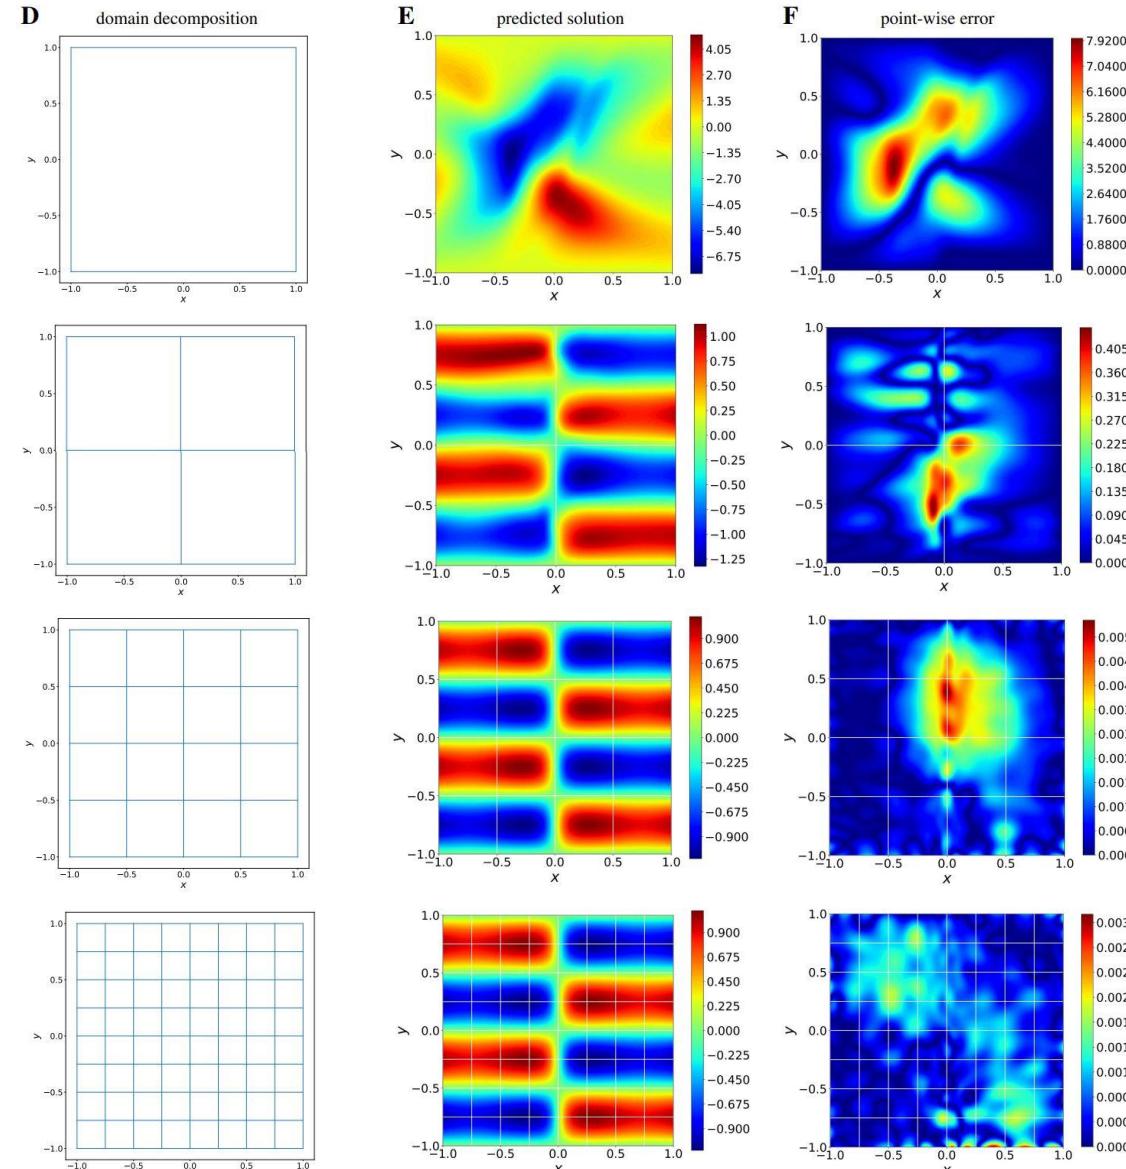
B



C



hp-VPINNs: Poisson's Equation



Kharazmi E, Zhang Z, Karniadakis GE. *hp*-VPINNs: Variational physics-informed neural networks with domain decomposition. *Computer Methods in Applied Mechanics and Engineering*. 2021 Feb 1;374:113547.

hp-VPINNs: Advection-Diffusion Equation

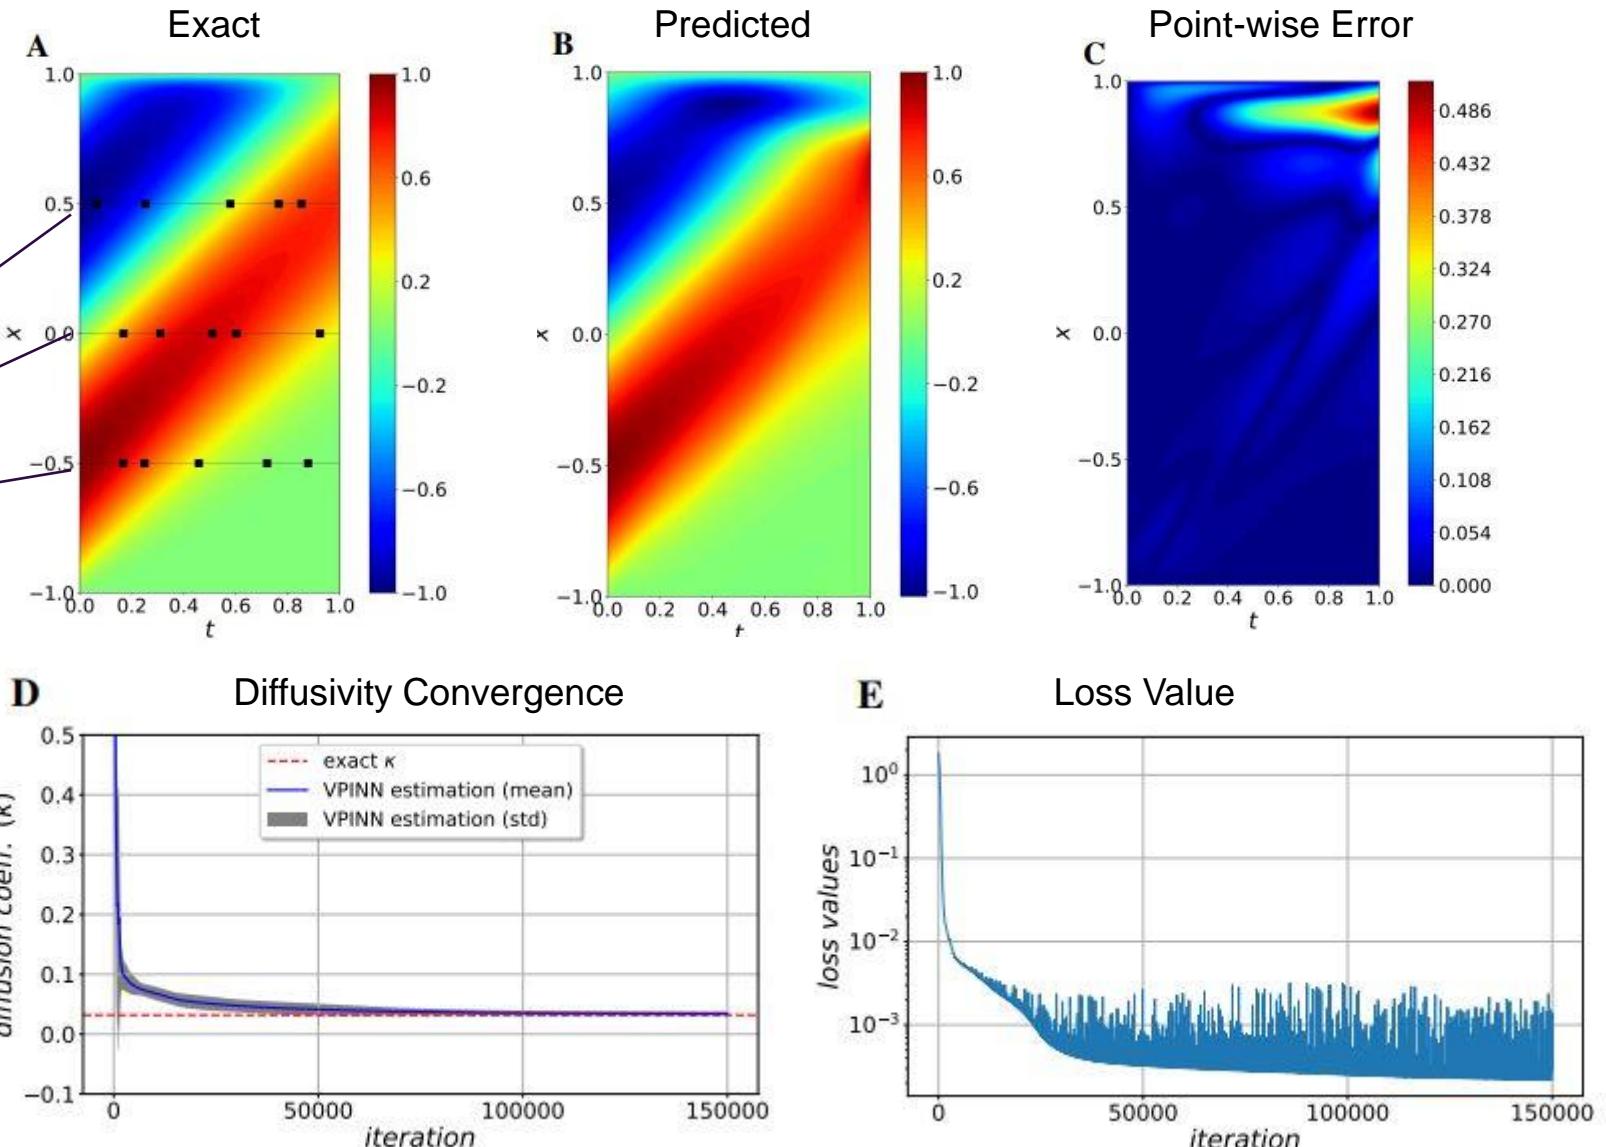
$$\Omega = [0, 1] \times [-1, 1]$$

$$\frac{\partial u}{\partial t} + v \frac{\partial u}{\partial x} = \kappa \frac{\partial^2 u}{\partial x^2},$$

$$u(-1, t) = u(1, t) = 0,$$

$$u(x, 0) = -\sin(\pi x),$$

Tree locations of sparse measurements in time





DEEP
LEARNING
INSTITUTE



Deep Learning for Science and Engineering Teaching Kit

Thank You

